
A Monte-Carlo HJ Reachability Sampling Scheme

Anonymous Author(s)

Affiliation

Address

email

Abstract

Backward reachable tubes (BRTs), computed via viscous Hamilton-Jacobi (HJ) partial differential equations, provide principled safety certificates for learned controllers and planning algorithms in trustworthy machine learning. However, classical grid-based HJ solvers require $O(M^n)$ memory footprint for M grid points per n state dimension. This renders them impractical for high-dimensional systems. We address this bottleneck with a local PDE linearization that enables a frozen-coefficient sampling scheme for the viscous HJ PDE: a generalized Cole-Hopf-type transformation reduces the nonlinear HJ equation to a sequence of linear heat equations whose solutions admit Gaussian heat-kernel representations. The value function and its spatial gradient are then recovered via roll-outs of Monte Carlo expectations on Gaussian densities, yielding a storage and grid-free algorithm that scales as $N \cdot n$ for N samples. This decoupling of memory from dimensionality enables reachability analysis on problems where grid-based methods are simply impossible. We prove a finite-sample concentration bound $O(N^{-1/2})$ error and conditional linear convergence for the introduced Monte-Carlo Picard iterative scheme. Numerical validation on pursuit-evasion games demonstrates relative L^2_{rel} errors of $0.03 - 0.20$, with $14 - 26$ second wall-clock times per 2D slice on a CPU. Crucially, the method scales with validation on up to (but not limited to) $n = 45$ -dimensional multi-agent games.

1 Introduction

A central challenge in trustworthy machine learning is certifying that a learned controller, neural policy, or planning algorithm will not violate safety constraints at deployment time. This concern appears directly in safe reinforcement learning [Li et al., 2021], model-based RL with learned dynamics [Berkenkamp et al., 2017], and any pipeline where a learned model must operate in a continuous, possibly high-dimensional state space under adversarial or uncertain inputs. Backward reachable (reach-avoid) tubes (BRTs/BRATs)—the set of all states from which a system is guaranteed to reach (avoid) a target region regardless of disturbances that may affect its behavioral evolution during deployment—provide a principled certificates for solving such problems. Computing a BRT amounts to solving a Hamilton-Jacobi-Isaacs (HJI) partial differential equation (PDE) [Mitchell et al., 2005, Lygeros, 2004], typically under two competing inputs.




Figure 1: Rockets BRT slices: Grid-based levelsetpy from [Molu, 2024b] vs. Monte Carlo (ours). The bottom row shows a heat map of the pointwise errors at various relative vehicle orientations θ ; errors are largest near the boundary of the usable part where $|Dv^\delta|$ is largest.

Central Contribution: Monte-Carlo Trajectory Optimization of Quasi-Linear BRTs

A *backward reachable tube* (BRT) is the set of initial states from which a controlled system is guaranteed to reach a target set within a given time horizon, for all adversarial inputs. Given system dynamics $\dot{x} = f(t; x, u, w)$ with controller u and disturbance w , the BRT is the zero sublevel set of a value function v^δ that solves the viscous HJI PDE (HJI-RCBRT-Visc). Classical grid-based solvers [Mitchell et al., 2005, Molu, 2025] discretize the state space on a grid with M points per dimension; the resulting memory cost is $O(M^n)$, which grows exponentially in the state dimension n . For $n = 6$ with $M = 100$, this already requires 10^{12} grid cells. The proposed method replaces the grid with N Monte Carlo samples of state space trajectories, reducing spatial memory to $O(N \cdot n)$.

31

32 The practical bottleneck for realizing the solutions to the BRT of complex systems is computational:
 33 classical HJ solvers scale as $O(M^n)$ in memory, where M is the number of grid points per dimension
 34 and n is the state dimension. While GPU-accelerated implementations [Molu, 2024b, 2025] reduce
 35 wall-clock time, they retain the same exponential memory footprint. Self-supervised, PINN-based
 36 approaches such as DeepReach [Bansal and Tomlin, 2021] train a neural network to minimize the
 37 HJ PDE residual directly via a physics-informed loss, without requiring a reference grid solution.
 38 However, the method requires a differentiable dynamics model and scales only to moderate dimen-
 39 sions (reported at 9D), while accuracy degrades as dimension increases. This paper introduces a
 40 fundamentally different approach: replacing grid storage with a sampling representation, so that the
 41 memory cost is $O(N \cdot n)$ and the algorithm is grid-free.

42 We consider first-order nonlinear scalar HJ PDEs of the form

$$v_t + \mathbf{H}(t; \mathbf{x}, Dv) = 0 \text{ in } \Omega \times (0, T], \quad v(0; \mathbf{x}) = \mathbf{g}(0; \mathbf{x}) \text{ on } \partial\Omega \times \{t = 0\} \quad (\text{HJ-IVP})$$

43 where the state $\mathbf{x}$ belongs to the open set $\Omega \subseteq \mathbb{R}^n$; v_t denotes the partial time derivative of $v(t; \mathbf{x})$;
 44 $\mathbf{H}(t; \mathbf{x}, Dv)$ is the Hamiltonian, continuously defined on $\mathbb{R}^n$, with Dv representing the spatial
 45 gradient; and $\mathbf{g}(\mathbf{x})$ being bounded and uniformly continuous (BUC) on $\mathbb{R}^n$.

46 Fixing a viscosity parameter $\delta > 0$, Crandall and Lions [Crandall and Lions, 1984] introduced the
 47 *vanishing viscosity* solution, v^δ , of (HJ-IVP) i.e. ,

$$v_t^\delta + \mathbf{H}(t; \mathbf{x}, Dv^\delta) = \frac{\delta}{2} \Delta v^\delta \text{ in } \Omega \times (0, T]; \quad v^\delta(0; \mathbf{x}) = \mathbf{g}(0; \mathbf{x}) \text{ on } \Omega \times \{t = 0\} \quad (\text{HJ-Visc})$$

48 such that (HJ-Visc) satisfies the uniform convergence bound

$$\sup_{t \in (0, T]} \sup_{\mathbf{x} \in \mathbb{R}^n} |v(t; \mathbf{x}) - v^\delta(t; \mathbf{x})| \leq k\sqrt{\delta} \quad (1)$$

49 for a constant $k > 0$. *Our key insight is that (HJ-Visc) can be reduced, via a generalized Cole-Hopf-*
 50 *type transformation, to a sequence of linear heat equations admitting explicit Gaussian heat-kernel*
 51 *solutions.* This yields a locally linearized PDE. The scheme is *not* an exact Cole-Hopf reduction for
 52 general Hamiltonians; it is an iterative (Picard) approximation in which the nonlinear coefficient is
 53 frozen at each step and updated after the linear solve. The exact reduction holds only for the quadratic
 54 case $\mathbf{H} = \frac{1}{2}|p|^2$, where the frozen coefficient is constant and the residual vanishes identically.

55 **Connections to ML.** Beyond reachability, the proposed scheme connects to several active ML
 56 research directions. (i) *Safe RL and policy verification*: certifying that a learned policy satisfies
 57 safety constraints requires computing BRTs for the closed-loop system; the memory bottleneck of
 58 grid-based solvers is the primary obstacle [Li et al., 2021, Berkenkamp et al., 2017]. (ii) *Model-based*
 59 *RL*: planning algorithms that learn a dynamics model and then verify safety via HJ reachability are
 60 limited to low-dimensional state spaces unless the memory cost is broken [Bansal and Tomlin, 2021].
 61 (iii) *Diffusion-based methods*: the Gaussian heat-kernel representation derived here is structurally
 62 similar to score-matching objectives and denoising score functions [Chaudhari et al., 2018, Heaton
 63 et al., 2024], opening potential bridges to score-based generative models for safety certificates.

64 **Contributions.** Our contributions are as follows. (i) A generalized Cole-Hopf-type transformation
 65 that decouples the nonlinear viscous HJ equation into linear heat equations with Gaussian heat-kernel
 66 solutions, enabling a frozen-coefficient quasi-linearization. (ii) A sampling-based algorithm (Algo-
 67 rithm 1) with $O(N \cdot n)$ memory complexity that scales with sample count rather than discretization
 68 dimension, offering an alternative to grid-based $O(M^n)$ approaches that become impractical in high
 69 dimensions. (iii) Theoretical guarantees: (i) a finite-sample concentration bound of $O(N^{-1/2})$ for the
 70 Monte Carlo value estimator, independent of state dimension, under explicit boundedness conditions
 71 (Theorem 2.5), and (ii) conditional linear convergence of the Picard iteration with explicit contraction
 72 constant (Theorem 2.8). (iv) Validation on two-player pursuit-evasion games achieving L_{rel}^2 errors
 73 of 0.03–0.20 (below the $O(\sqrt{\delta}) \approx 0.28$ viscosity bound) in 12–15 iterations. Demonstration on
 74 45-dimensional multi-agent problems where grid-based discretization becomes computationally
 75 prohibitive.

76 The rest of this work is structured as follows. Readers not familiar with existing literature can consult
 77 the background in Appendix B. Section 2 derives the quasi-linearization, states the algorithm, and
 78 establishes the theoretical guarantees. Section 3 presents numerical results. Section 4 concludes. A
 79 thorough background is provided in Appendix A; all proofs are in Appendix C; additional numerical
 80 experiments appear in Appendix D. Some of the bottlenecks we encountered during the process of
 81 making this paper are described in Appendix E.

82 2 Transformation of the HJ PDEs

83 We now transform the nonlinear viscous HJ PDE (HJ-Visc) into a linearized form and extract its
 84 solution via Gaussian heat-kernel expectations. All proofs are collected in Appendix C.

HJ Sign Convention

Throughout this paper we adopt the backward-reachability viscosity-solution convention:

$$v_t + \mathbf{H}(t; \mathbf{x}, D\mathbf{v}) = 0,$$

where the Hamiltonian is defined via the zero-sum differential game formulation (Appendix A.5):

$$\mathbf{H}(t; \mathbf{x}, p) = \max_{\mathbf{u}} \min_{\mathbf{w}} \langle p, f(t; \mathbf{x}, \mathbf{u}, \mathbf{w}) \rangle.$$

Negative signs appearing in the Hamiltonian expressions in Sections B and Appendix A arise from the backward-time transformation and the min-over-max structure of reachability propagation. This convention is standard in reachability literature (Mitchell et al. 2005, Tomlin et al. 2003) and ensures consistency with viscosity-solution theory.

85

2.1 Exact Reduction and Quasi-Linearization: Two Cases

86 Before stating the main results, we draw an explicit distinction that is essential for the interpretation
87 of the method.

88 **Exact case** ($\mathbf{H} = \frac{1}{2}|p|^2$). When the Hamiltonian is purely quadratic in the co-state, setting $\mathbf{c} =$
89 $1/\delta$ (a constant) and $\omega^\delta := \exp(-\mathbf{v}^\delta/\delta)$ is an *exact* Cole-Hopf transformation: ω^δ satisfies the
90 homogeneous heat equation $\omega_t^\delta = (\delta/2)\Delta\omega^\delta$ with no residual. In this case what follows below is
91 exact.

92 **General case (quasi-linearization)**. For a general Hamiltonian $\mathbf{H}(t; \mathbf{x}, p)$, define the spatially-
93 varying coefficient,
94

$$\mathbf{c}(t; \mathbf{x}) = \frac{2}{\delta} \cdot \mathbf{H}^\delta / |D\mathbf{v}^\delta|^2, \quad (2)$$

95 where $\mathbf{H}^\delta := \mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta)$. The transformation $\omega^\delta := \exp(-\mathbf{c}\mathbf{v}^\delta)$ applied to (HJ-Visc) yields a
96 heat equation with a non-zero residual term $\mathbf{R}(t; \mathbf{x})$ that involves spatial and temporal derivatives
97 of $\mathbf{c}$ (see Appendix C for the derivation). Our proposed algorithm treats this as a *Picard quasi-*
98 *linearization*: at iteration k , the coefficient $\mathbf{c}^{(k)}$ is *frozen* at the current iterate, the resulting linear
99 heat equation is solved for $\omega^{(k+1)}$, the value $\mathbf{v}^{(k+1)}$ is recovered, the gradient $D\mathbf{v}^{(k+1)}$ is updated,
100 and $\mathbf{c}^{(k+1)}$ is recomputed. When $\mathbf{c}^{(k)}$ stabilizes, $\mathbf{R} \rightarrow 0$ effectively. There is no joint solve of $\mathbf{c}$ with
101 the unknown iterate; the apparent circularity in (2) is resolved by the frozen-coefficient interpretation.
102 This distinction is stated precisely in Remark 2.4 below and is the basis of Algorithm 1.

2.2 HJ Linearization

103 Throughout this section we drop the arguments of $\mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta)$ and write $\mathbf{H}^\delta$ for brevity. For the
104 terminal cost $\mathbf{g}(0; \mathbf{x})$ we write $\mathbf{g}(\mathbf{x})$. With $\mathbf{c}^{(k)}$ frozen, the transformation,
105

$$\omega^\delta := \exp(-\mathbf{c}\mathbf{v}^\delta), \quad (\text{Linear-Trans})$$

106 reduces (HJ-Visc) to a heat equation under the frozen-coefficient approximation, as made precise in
107 Proposition 2.1. *For general nonlinear Hamiltonians, this transformation induces a residual term*
108 *and should be interpreted as a quasi-linearization, not an exact linearization.* The approximation is
109 exact only when the Hamiltonian is quadratic in the co-state.

110 **Proposition 2.1** (Heat-equation reduction under frozen coefficient). *With $\mathbf{c}^{(k)}$ frozen as a spatially-*
111 *varying coefficient (exact when $\mathbf{H} = \frac{1}{2}|p|^2$ and $\mathbf{c} = 1/\delta$; approximate otherwise via Picard*
112 *quasi-linearization as in §2.1), the transformed variable $\omega^\delta := \exp(-\mathbf{c}\mathbf{v}^\delta)$ satisfies,*

$$\omega_t^\delta - \frac{\delta}{2}\Delta\omega^\delta = 0 \text{ in } \Omega \times (0, T], \quad \omega^\delta = \exp(-\mathbf{c}\mathbf{g}(\mathbf{y})) \text{ on } \partial\Omega \times \{t = 0\}. \quad (\text{Linear-HJ})$$

113 The solution to (Linear-HJ) admits the following unique, explicit Green's-convolution representation,

$$\omega^\delta(t; \mathbf{x}) = \frac{1}{(\sqrt{2\pi\delta t})^n} \int_{\Omega} \exp\left(-\frac{|\mathbf{x} - \mathbf{y}|^2}{2\delta t}\right) \exp(-\mathbf{c}\mathbf{g}(\mathbf{y})) d\mathbf{y}, \quad \mathbf{x} \in \Omega, t > 0, \quad (3)$$

114 which can equivalently be written as a Gaussian expectation,

$$\omega^\delta(t; \mathbf{x}) = \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-\mathbf{c}\mathbf{g}(\mathbf{y}))]. \quad (4)$$

115 **Lemma 2.2** (Smoothed HJ Solution). *Under the frozen-coefficient interpretation of (2) (or exactly*
 116 *when $\mathbf{H} = \frac{1}{2}|p|^2$), the solution to the smoothed Hamilton-Jacobi equation (HJ-Visc) is,*

$$v^\delta(t; \mathbf{x}) = -\frac{1}{c^{(k)}} \log \left\{ \frac{1}{(\sqrt{2\pi\delta t})^n} \int_{\Omega} \exp \left(-\frac{|\mathbf{x} - \mathbf{y}|^2}{2\delta t} \right) \exp \left(-c^{(k)} g(\mathbf{y}) \right) d\mathbf{y} \right\}, \quad (5)$$

117 *in $\mathbb{R}^n \times (0, T]$, where $c^{(k)}$ is the frozen coefficient at iteration k . This can equivalently be written as,*

$$v^\delta(t; \mathbf{x}) = -\frac{1}{c^{(k)}} \cdot \log \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} \left[\exp \left(-c^{(k)} g(\mathbf{y}) \right) \right]. \quad (6)$$

118 **Lemma 2.3** (Spatial gradients of the HJ solution). *Under the same frozen-coefficient assumption as*
 119 *Lemma 2.2, the spatial gradient of the value function $v^\delta(t; \mathbf{x})$ admits the form,*

$$Dv^\delta = \frac{1}{t \cdot \delta \cdot c^{(k)}} \left(\mathbf{x} - \frac{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\mathbf{y} \cdot \exp(-c^{(k)} g(\mathbf{y}))]}{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-c^{(k)} g(\mathbf{y}))]} \right), \quad (7)$$

120 *where $\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)$, i.e., with mean $\mathbf{x}$ and variance $\delta t I_n$.*

121 **Remark 2.4** (Exact reduction vs. quasi-linearization). Lemmas 2.2 and 2.3 are exact when $\mathbf{H} = \frac{1}{2}|p|^2$,
 122 in which case $c = 1/\delta$ is a constant and the residual $\mathbf{R}(t; \mathbf{x})$ in the linearized equation vanishes
 123 identically (see Appendix C). For general Hamiltonians, $c^{(k)}$ depends on the unknown Dv^δ ; we
 124 resolve this by the Picard scheme of Algorithm 1: $c^{(k)}$ is formed from the previous iterate, frozen,
 125 and updated after each linear solve. The formulas in both lemmas apply at each iteration with $c^{(k)}$ in
 126 place of c .

127 2.3 Relations to Backward Reachable Sets/Tubes

128 Since we now have the derivatives that constitute the viscous HJ equation, it is easy for us to rewrite
 129 (HJI-RCBRT) as

$$v_t^\delta(t; \mathbf{x}) + \min\{0, \mathbf{H}^\delta - \frac{\delta}{2} \Delta v^\delta\} = 0, \quad v^\delta(0; \mathbf{x}) = g(0; \mathbf{x}). \quad (\text{HJI-RCBRT-Visc})$$

130 so that the solution to (HJI-RCBRT-Visc) is from integrating the quantities given in Lemma 2.3
 131 is tantamount to sampling from the kernels of the given Gaussian densities. The invariant set $\mathcal{L}_0$
 132 obtained at T is

$$\mathcal{L}_0(T) = \{\mathbf{x} \in \mathbb{R}^n \mid v^\delta(0; \mathbf{x}) \leq 0\}, \quad (8)$$

133 and the RCBRT is as given in (HJI-RCBRT). Similarly, the viscous version of the *reach-avoid BRT*
 134 i.e. (HJI-RCBRAT) can be represented as

$$\min\{v_t^\delta(t; \mathbf{x}) + \mathbf{H}^\delta - \frac{\delta}{2} \Delta v^\delta, g(t; \mathbf{x}) - \ell(t; \mathbf{x})\} \leq 0, \quad v^\delta(0; \mathbf{x}) = g(0; \mathbf{x}). \quad (\text{HJI-RCBRAT-Visc})$$

135 2.4 Weighted Importance Sampling of Reachable Sets

136 While we may not recover the exact value function system, the mean and variance of its derivatives can
 137 be evaluated with the kernel expectations of Lemma 2.2, and 2.3. Previously, self-supervised physics-
 138 informed neural network (PINN) approaches such as DeepReach [Bansal and Tomlin, 2021] train
 139 neural networks to minimize the HJ PDE residual directly without requiring external supervision or a
 140 reference grid solution. However, neural network-based approximation methods encounter inherent
 141 scaling limitations; DeepReach scales only to moderate dimensions (reported to 9D/10D in [Bansal
 142 and Tomlin, 2021]), and approximation accuracy degrades as dimension increases. ***Our sampling-
 143 based scheme addresses this bottleneck by decoupling the memory cost from dimensionality.***

144 We propose a sequential importance sampling (IS) scheme for a high-dimensional $\mathbf{x}$ and consequently
 145 v^δ can be devised for the target density $\mathbb{P}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)}$ to induce a chain-like decomposition of $v^\delta :=$
 146 $(v_1^\delta, \dots, v_d^\delta), 1 \leq d \leq T$ i.e., $v^\delta(t; \mathbf{x}) = v^\delta(d_1; \mathbf{x}_1) \prod_{d=2}^T v^\delta(d; \mathbf{x}_d) |v^\delta(d; \mathbf{x}_{[1:d-1]})$ from which
 147 we can build an importance sampling such as $q^\delta(t; \mathbf{x}) = q_1^\delta(d_1; \mathbf{x}_1) \prod_{d=2}^T q_d^\delta(d; \mathbf{x}_d) |q_d^\delta(d; \mathbf{x}_{[1:d-1]})$.
 148 Corollaries to Lemmas 2.2 and 2.3 are in the appendix. The quasi-linearization iteration scheme
 149 is stated in Algorithm Algorithm 1 for recovering the backward reachable sets/tubes based on the
 150 Gaussian expectation formulas computed in Lemmas 2.2, 2.3. Note that Algorithm 1 is essentially a
 151 Picard fixed-point iterative scheme.

Algorithm 1 Quasi-Linearization Algorithm Cole-Hopf

Fix: $\epsilon > 0$, $\mathbf{v}^{(0)}(t; \mathbf{x}) = g(\mathbf{x})$, $\mathbf{c}^{(0)} = \frac{2\mathbf{H}(t; \mathbf{x}, D\mathbf{g})}{\delta|D\mathbf{g}|^2}$.

For $k = 0, 1, 2, \dots$:

1. Freeze $\mathbf{c}^{(k)}$ at the current iterate.
 2. Solve the heat equation $\omega_t = \frac{\delta}{2}\Delta\omega$ with initial data $\omega^{(k)}(0; \mathbf{x}) = e^{-\mathbf{c}^{(k)}g(\mathbf{x})}$.
 3. Recover $\mathbf{v}^{(k+1)} = -(1/\mathbf{c}^{(k)})\log\omega^{(k+1)}$.
 4. Update $D\mathbf{v}^{(k+1)}$ and $\mathbf{c}^{(k+1)} = \frac{2\mathbf{H}(t; \mathbf{x}, D\mathbf{v}^{(k+1)})}{\delta|D\mathbf{v}^{(k+1)}|^2}$.
 5. Check convergence: $\|\mathbf{v}^{(k+1)} - \mathbf{v}^{(k)}\|/\|\mathbf{v}^{(k)}\| < \epsilon$.
-

2.5 Sampling Complexity and Convergence Guarantees

We now analyze the sample complexity of the scheme with coefficient $\mathbf{c}$ frozen.

Theorem 2.5 (Finite-sample concentration of the frozen-coefficient value estimator). *Fix phase $(t; \mathbf{x}) \in [0, T) \times \mathbb{R}^n$ and a frozen coefficient $\mathbf{c} > 0$ (satisfied when the Hamiltonian is bounded away from zero in regions of smoothness). Let $\sigma \triangleq \sqrt{\delta(T-t)}$, $\zeta \sim \mathcal{N}(\mathbf{x}, \delta(T-t)I_n)$, and assume that the terminal cost is bounded, i.e. there exist constants $g_{\min} \leq g_{\max}$ such that $g_{\min} \leq g(\zeta) \leq g_{\max}$, $\forall \zeta \in \mathbb{R}^n$.*

Let $Z \triangleq \exp(-\mathbf{c}g(\zeta))$, $\mu \triangleq \mathbb{E}[Z]$, $\mathbf{v}_c(t; \mathbf{x}) \triangleq -\frac{1}{\mathbf{c}}\log\mu$. Given i.i.d. samples $\zeta_1, \dots, \zeta_N \sim \mathcal{N}(\mathbf{x}, \delta(T-t)I_n)$, let $Z_i \triangleq \exp(-\mathbf{c}g(\zeta_i))$, $\bar{Z}_N \triangleq \frac{1}{N}\sum_{i=1}^N Z_i$, and $\hat{\mathbf{v}}_{c,N}(t; \mathbf{x}) \triangleq -\frac{1}{\mathbf{c}}\log\bar{Z}_N$. Let $a \triangleq e^{-\mathbf{c}g_{\max}}$, $b \triangleq e^{-\mathbf{c}g_{\min}}$. Then $a \leq Z_i \leq b$ almost surely, and for every $\epsilon > 0$,

$$\mathbb{P}(|\hat{\mathbf{v}}_{c,N}(t; \mathbf{x}) - \mathbf{v}_c(t; \mathbf{x})| \geq \epsilon) \leq 2 \exp\left(-\frac{2N\mu^2(1 - \exp(-\mathbf{c}\epsilon))^2}{(b-a)^2}\right). \quad (9)$$

Remark 2.6 (Tightness via Bernstein and Jensen bounds). The Hoeffding bound in (9) is asymptotically tight but ignores variance. When $\text{Var}(Z) \ll (b-a)^2/4$, a Bernstein tail bound yields strictly tighter concentration. Additionally, Corollary 2.7 below uses the loose lower bound $\mu \geq a = e^{-\mathbf{c}g_{\max}}$; when a sharper estimate (e.g., via Jensen, $\mu \geq e^{-\mathbf{c}\mathbb{E}[g(\zeta)]}$) is available, the required sample size can be substantially reduced.

Corollary 2.7 (Explicit sample size independent of μ). *In particular, since $\mu \geq a$, it is sufficient to choose*

$$N \geq \frac{(b-a)^2}{2a^2(1 - \exp(-\mathbf{c}\epsilon))^2} \log\left(\frac{2}{\alpha}\right) \quad (10)$$

to guarantee $\mathbb{P}(|\hat{\mathbf{v}}_{c,N}(t; \mathbf{x}) - \mathbf{v}_c(t; \mathbf{x})| \geq \epsilon) \leq \alpha$.

Theorem 2.8 (Contraction convergence of Algorithm Algorithm 1). *Fix evaluation states $\mathbf{x}_1, \dots, \mathbf{x}_M \in \mathbb{R}^n$ and a time $t \in [0, T)$. Let $\mathcal{G} : V \rightarrow (\mathbb{R}^n)^M$ denote a stable deterministic gradient reconstruction operator, which may arise from smoothing, interpolation, kernel regression, or variance-controlled Monte Carlo estimation used by Algorithm Algorithm 1, and write $p_m(v) \triangleq (\mathcal{G}(v))_m$. For $v \in V$, define the coefficient-update map by $(\Gamma(v))_m \triangleq 2\mathbf{H}(t; \mathbf{x}_m, p_m(v))/(\delta|p_m(v)|^2)$, $m = 1, \dots, M$, and define the frozen-coefficient heat-kernel map by*

$$(\Phi(c))_m \triangleq -\frac{1}{c_m} \log \mathbb{E}_{\zeta \sim \mathcal{N}(\mathbf{x}_m, \delta(T-t)I_n)} [\exp(-c_m g(\zeta))]. \quad (11)$$

Let $\Lambda \triangleq \Phi \circ \Gamma$ so that Algorithm Algorithm 1 is the iteration $v^{(k+1)} = \Lambda(v^{(k)})$.

Then Λ is a contraction on $\mathcal{A}$. Consequently, Λ has a unique fixed point $v^ \in \mathcal{A}$, and for every initial iterate $v^{(0)} \in \mathcal{A}$ the sequence generated by Algorithm Algorithm 1 converges linearly to v^* with*

$$\|v^{(k+1)} - v^*\|_\infty \leq q \|v^{(k)} - v^*\|_\infty \leq q^{k+1} \|v^{(0)} - v^*\|_\infty. \quad (12)$$

Moreover,

$$\begin{aligned} \|\Gamma(v^{(k)}) - \Gamma(v^*)\|_\infty &\leq \frac{2L_D}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) \|v^{(k)} - v^*\|_\infty, \\ \|v^{(k+1)} - v^{(k)}\|_\infty &\leq q^k \|v^{(1)} - v^{(0)}\|_\infty, \end{aligned} \quad (13)$$

181 and the a posteriori error estimate

$$\|v^{(k)} - v^*\|_\infty \leq \frac{q}{1-q} \|v^{(k)} - v^{(k-1)}\|_\infty \quad (14)$$

182 holds for every $k \geq 1$.

183 *Remark 2.9.* Theorem 2.8 is a convergence result for the *frozen-coefficient numerical map* imple-
 184 mented by Algorithm 1. It shows conditional linear convergence under explicit regularity,
 185 nondegeneracy, and contraction hypotheses. It does not claim unconditional global convergence for
 186 arbitrary Hamiltonians, nor does it identify the fixed point with the exact solution of the original
 187 nonlinear HJ PDE without further assumptions.

188 3 Numerical Results

189 All experiments run on a single CPU core of an Intel Core i7-14700K processor (20 physical cores,
 190 33.8 MiB L3 cache per core, base clock 5.6 GHz) with 31 GiB RAM running Ubuntu 22.04. The
 191 processor supports SIMD vectorization (AVX2, SSE4.2) which JAX and PyTorch leverage for batch
 192 operations. All reported metrics are averaged across $N_{\text{runs}} = 5$ independent trials with different
 193 random seeds; tables and figures report mean values with standard deviation error bars where
 194 applicable.

195 3.1 Territorial Defense: A Game of Two Rockets on a Plane

196 We adopt the rockets launch problem (see Fig. 7) of Dreyfus [Dreyfus, 1966] and cast it as a terminal
 197 differential game between two identical rockets, $\mathbf{P}$ and $\mathbf{E}$, on an $(x - z)$ cross-section of a Cartesian
 198 plane in [Molu, 2024a, 2025]. The game terminates when *capture* occurs, i.e., the distance $\|\mathbf{PE}\|$
 199 becomes less than a certain prespecified (scalar) quantity. The states of $\mathbf{P}$ and $\mathbf{E}$ are denoted as
 200 (x_p, x_e) respectively, driven by thrusts (u_p, u_e) in the (x, z) -plane. The motion of $\mathbf{P}$ relative to
 201 $\mathbf{E}$'s along the $(x - z)$ plane includes the relative orientation, the control input, shown in Fig. 7 as
 202 $\theta = u_p - u_e$. Full derivation of the scheme is provided in Appendix B.9. **Scope.** The proposed method
 203 is not intended to outperform structured grid solvers in low-dimensional settings ($n \leq 4$), where
 204 highly-optimized implementations exist; rather, its advantage emerges as dimensionality increases
 205 and grid storage becomes prohibitive.

206 3.2 Dimensional Scalability: A 15 Rockets System in a Pursuit-Evasion Game

207 To validate the scalability benefits of Algorithm 1, we consider a 15-rocket multi-pursuer single-
 208 evader game with state dimension $n = 45$; this dimension is intractable for grid-based solvers. State
 209 is $\mathbf{x} = (x_1, y_1, \theta_1, \dots, x_{15}, y_{15}, \theta_{15}) \in \mathbb{R}^{45}$, where each agent $i \in \{1, \dots, 15\}$ has position (x_i, y_i) ,
 210 heading θ_i , and control input $u_i \in [-1, +1]$ regulating turn rate. Dynamics are:

$$\dot{x}_i = a_i \cos(\theta_i), \quad \dot{y}_i = a_i \sin(\theta_i), \quad \dot{\theta}_i = u_i, \quad (15)$$

211 with forward speeds $a_i \in \mathbb{R}_{>0}$. The evader (agent 15) seeks escape; capture occurs when any of 14
 212 pursuers reach distance $r_{\text{capture}} \leq 1.5$ ft. Target set: $\phi(\mathbf{x}) = \min_{i \leq 14} \|\mathbf{x}_i^{\text{pos}} - \mathbf{x}_{15}^{\text{pos}}\|_2 - r_{\text{capture}}$.

213 We test three speed regimes to examine relative agent capabilities on the reachable set, (1.) evader
 214 speed $a_{\text{evader}} = 2.0$ ft/s, pursuer speeds $a_{\text{pursuers}} = 1.0$ ft/s; evader advantage; (2.) all agents equal
 215 speed $a = 1.0$ ft/s; balanced game; (3.) evader speed $a_{\text{evader}} = 1.0$ ft/s, pursuer speeds $a_{\text{pursuers}} = 2.0$
 216 ft/s; pursuer advantage.

217 This $n = 45$ problem is computationally inaccessible to grid-based reachability solvers. A 101-
 218 point grid would require $101^{45} \approx 10^{90}$ cells, exceeding all conceivable storage. Our method
 219 operates at $O(N \cdot n) \approx 7.2$ MB per iteration, where the memory cost is independent of the state
 220 dimension exponent. The quasi-linear Picard iteration (Algorithm 1) computes the relative change
 221 in value function between successive iterations; Table 1 reports the final relative residual, $\varepsilon =$
 222 $\|\mathbf{v}^{(k+1)} - \mathbf{v}^{(k)}\| / \|\mathbf{v}^{(k)}\|$ as a measure of iteration convergence.

223 Table 1 shows the results for 15-rocket pursuit-evasion across three speed regimes. Wall-clock times
 224 are reported on a 16-core Intel i7-14700K. The residual is the final relative L^2 change between Picard
 225 iterates. All three cases converge in 15 iterations with residuals below 10^{-2} , demonstrating stability

Table 1: 15-agent multi-pursuer single-evader game with three speed regimes on $n = 45$.

Case	a_{evader}	a_{pursuers}	Iterations	$\varepsilon(k)$	Wall-clock (s)
1. (Evader faster)	2.0	1.0	15	0.003121	27.6
2. (Equal speed)	1.0	1.0	15	0.002454	30.0
3. (Pursuers faster)	1.0	2.0	15	0.003462	27.6

across different game parameters. Memory consumption remains constant at 7.2 MB per iteration; wall-clock variance reflects Monte Carlo sampling fluctuations. The results demonstrate that the $O(N \cdot n)$ scaling becomes practical in 45D, where grid-based methods are intractable; the achievable fidelity (sub-1

3.3 Experimental Setup and Quantitative Results: Pursuit-Evasion Games

Viscosity parameter selection. The viscosity parameter δ controls the degree of smoothing in the Cole-Hopf transformation. By Crandall-Lions theory, the approximation error scales as $O(\sqrt{\delta})$, so smaller δ yields more accurate solutions but requires higher Monte Carlo sample counts to maintain accuracy. In our experiments, we empirically set $\delta = 0.08$ – 0.1 , which balances approximation error $O(\sqrt{\delta}) \approx 0.28$ with manageable variance in the log-sum-exp estimator. The choice $\delta \in [0.05, 0.2]$ is robust across all benchmark problems; practitioners should adjust based on the required accuracy-to-computation trade-off.

Frozen-coefficient bias. The quasi-linearization approximates $|Dv^\delta|$ by $|Dg|$ at each Picard iterate, introducing a systematic bias that is largest near the zero level set where $|Dv^\delta|$ varies most rapidly. Theorem 2.8 guarantees convergence to a fixed point of the frozen-coefficient map, not to the true viscous solution; the gap between these two objects depends on how well $|Dg|$ tracks $|Dv^\delta|$ along the optimal trajectory, a quantity we do not bound here. Empirically, this bias manifests as larger L^∞ errors near the reachability boundary where the gradient is most active; the method mitigates this through the viscosity smoothing and iterative refinement of the frozen coefficient.

Numerical realization: We discretize the spatial and temporal domains as,

$$(x, z) \in (-100, +100], \quad \theta \in (-\pi/2, \pi/2], \quad t \in (0, 1], \quad (16)$$

with 1,000 temporal grid points and spatial step $x_k - x_{k-1} = 10^{-2}$.

We set the target set as the ℓ_2 -ball $\ell(x) = \sqrt{x^2 + z^2}$ and run Algorithm 1 over the time range $(0, 1]$ with Dirichlet boundary conditions,

$$v^\delta(0; x) = g(0; x) \triangleq 0. \quad (17)$$

Spatial gradients (co-states) are computed via (6) and (7); the Hamiltonian is evaluated iteratively using $N = 14,000$ Monte Carlo samples per iteration according to Algorithm 1. The results are seen in Fig. 1.

Convergence. Algorithm 1 converges within 20 iterations across all three benchmark systems, with final relative residuals in the range 0.14–0.22. This is consistent with the linear convergence guarantee of Theorem 2.8: the contraction factor $q < 1$ under the stated regularity conditions implies that the residual sequence $\|v^{(k+1)} - v^{(k)}\|/\|v^{(k)}\|$ decays geometrically, as observed empirically.

BRT geometry. Errors concentrate near the zero level-set boundary where $|\nabla v^\delta|$ is maximal, as expected in frozen-coefficient approximations [Crandall et al., 1992].

Table 2 reports point-wise error metrics (L^∞ and L_{rel}^2) for 2D slices on the 3D Dubins and Rockets pursuit-evasion benchmarks across three representative heading angles, evaluated on a uniform 40×40 grid with reference from LevelSetPy’s 45³ solution interpolated to the same points. The L^∞ (point-wise maximum) error represents worst-case divergence, useful where outlier errors matter. The relative RMS error, L_{rel}^2 weights by solution magnitude and captures overall accuracy. Notably, the empirical L_{rel}^2 errors substantially undercut the worst-case viscosity bound $O(\sqrt{\delta}) \approx 0.283$ [Crandall et al., 1992]: at $\theta = 0$, both systems achieve $L_{\text{rel}}^2 < 0.04$, indicating that the Monte Carlo sampling error is well-controlled relative to the Cole-Hopf approximation error. In contrast, L^∞ errors

Table 2: Pursuit-evasion games: Monte Carlo vs. textttLevelSetPy error metrics on 2D (x, z) slices. Both CPU-based; MC uses $N = 14,000$ samples per quasi-linear iteration. Crandall-Lions theory gives approximation error bound $O(\sqrt{\delta})$ where $\delta = 0.08 \Rightarrow O(\sqrt{\delta}) \approx 0.283$ (viscosity-induced smoothing).

System	θ (rad)	L^∞	L_{rel}^2	MC time (s)	Iters
Rockets	$-\pi/2$	0.844	0.147	14.7	12
	0	1.044	0.134	13.9	12
	$\pi/2$	0.740	0.099	14.0	12
Dubins	$-\pi/2$	1.139	0.197	25.5	15
	0	0.821	0.032	24.9	15
	$\pi/2$	1.163	0.196	24.7	15

($\sim 0.8 - 1.2$) are larger and driven by pointwise deviations near the zero level-set boundary where $|\nabla v^\delta|$ is maximal; this is expected in frozen-coefficient approximations [Crandall et al., 1992]. The 3D isosurface computation over 15,625 sample points completes in 125 seconds (rockets) and 75.1 seconds (Dubins) on a single CPU, compared to 1.1 seconds (rockets) and 3.2 seconds (Dubins) for the dense $45^3 = 91,125$ -point grids. The zero level-set boundary is extracted via marching cubes [Lorensen and Cline, 1987]. The wall-clock overhead of the Monte Carlo method reflects the sample count ($N = 14,000$ per iteration) necessary to balance variance in the Cole-Hopf estimator against quasi-linearization residuals, a tradeoff that becomes favorable in higher dimensions where grid storage becomes prohibitive.

Scalability. Algorithm 1 requires a memory footprint per iteration of $O(N \cdot n)$ for N samples per n state dimensions. With $N = 14,000$ and $n = 3$, each iteration allocates approximately 0.6 MB for samples and intermediate values. In contrast, the dense 45^3 textttLevelSetPy grid requires approximately 1.5 MB for the value function and gradient storage. Each 2D θ -slice completes in 14 – 26 seconds on a single core, fully parallelizable across the CPU’s 20 cores. This reduces aggregate wall-clock time proportionally with available core count and makes the method well-suited to parallel architectures (GPUs, HPC clusters) where grid-based methods become memory-prohibitive in dimensions $n \geq 5$.

4 Conclusion

We have presented a quasi-linearized, frozen-coefficient sampling scheme for the viscous Hamilton-Jacobi PDE arising in safety analysis of dynamical systems. By applying a generalized Cole-Hopf-type transformation — exact when $\mathbf{H} = \frac{1}{2}|p|^2$ and iteratively approximated via Picard quasi-linearization for general Hamiltonians — the nonlinear HJ equation is reduced to a sequence of linear heat equations whose solutions are Gaussian heat-kernel expectations. The value function and its spatial gradient are recovered from these expectations via Monte Carlo sampling, yielding a storage-free, grid-free algorithm with memory cost $O(N \cdot n)$ rather than the $O(M^n)$ cost of classical grid-based solvers.

Theorem 2.5 provides a finite-sample concentration bound for the frozen-coefficient Monte Carlo estimator, establishing a standard $O(N^{-1/2})$ error rate under explicit boundedness conditions on the terminal cost. Theorem 2.8 establishes conditional linear convergence of the Picard fixed-point iteration under Lipschitz and nondegeneracy assumptions, with an explicit contraction constant and a posteriori error estimate. Neither result claims unconditional global convergence for arbitrary Hamiltonians. Numerical experiments on three benchmark reachability problems i.e. a 3D rocket pursuit-evasion game, a 3D Dubins two-car game, and a double integrator plant confirm convergence with residuals of 0.14 – 0.22.

The connection to safe reinforcement learning, policy certification, and model-based control with learned dynamics suggests several directions for future work. These include adaptive importance sampling to reduce variance in high dimensions, extension to systems with stochastic dynamics, and tighter integration with deep learning pipelines for scalable safety certification of learned policies.

References

- S. Bansal and C. J. Tomlin. Deepreach: A Deep Learning Approach to High-dimensional Reachability. In *2021 IEEE International Conference on Robotics and Automation (ICRA)*, pages 1817–1824. IEEE, 2021.
- R. Bellman. *Dynamic programming*. Princeton University Press, 1957. ISBN 0-486-42809-5.
- F. Berkenkamp, M. Turchetta, A. P. Schoellig, and A. Krause. Safe model-based reinforcement learning with stability guarantees. In *Advances in Neural Information Processing Systems 30*, pages 908–918. Curran Associates, Inc., 2017. URL <https://papers.nips.cc/paper/6692-safe-model-based-reinforcement-learning-with-stability-guarantees>.
- S. P. Bhat and D. S. Bernstein. Continuous finite-time stabilization of the translational and rotational double integrators. *IEEE Transactions on automatic control*, 43(5):678–682, 1998.
- P. Chaudhari, A. Oberman, S. Osher, S. Soatto, and G. Carlier. Deep relaxation: partial differential equations for optimizing deep neural networks. *Research in the Mathematical Sciences*, 5:1–30, 2018.
- M. G. Crandall and P.-L. Lions. Viscosity solutions of hamilton-jacobi equations. *Transactions of the American mathematical society*, 277(1):1–42, 1983.
- M. G. Crandall and P.-L. Lions. Two Approximations of Solutions of Hamilton-Jacobi Equations. *Mathematics of Computation*, 43(167):1–19, 1984.
- M. G. Crandall and A. Majda. Monotone Difference Approximations For Scalar Conservation Laws. *Mathematics of Computation*, 34(149):1–21, 1980.
- M. G. Crandall, L. C. Evans, and P. L. Lions. Some Properties of Viscosity Solutions of Hamilton-Jacobi Equations. *Transactions of the American Mathematical Society*, 282(2):487, 1984. ISSN 00029947.
- M. G. Crandall, H. Ishii, and P.-L. Lions. User’s guide to viscosity solutions of second order partial differential equations. *Bulletin of the American Mathematical Society*, 27(1):1–67, 1992.
- S. E. Dreyfus. Control Problems With Linear Dynamics, Quadratic Criterion, and Linear Terminal Constraints. Technical report, Rand Corp, Santa Monica Calif, 1966.
- L. E. Dubins. On curves of minimal length with a constraint on average curvature, and with prescribed initial and terminal positions and tangents. *American Journal of mathematics*, 79(3):497–516, 1957.
- L. Evans and P. E. Souganidis. Differential Games And Representation Formulas For Solutions Of Hamilton-Jacobi-Isaacs Equations. *Indiana Univ. Math. J*, 33(5):773–797, 1984. ISSN 0022-2518.
- L. C. Evans. *Partial Differential Equations*, volume 19. American Mathematical Society, 2022.
- E. M. Goldratt. *The Goal: Theory of Constraints*. North River Press, 1992.
- H. Heaton, S. Wu Fung, and S. Osher. Global solutions to nonconvex problems by evolution of hamilton-jacobi pdes. *Communications on Applied Mathematics and Computation*, 6(2):790–810, 2024.
- R. Isaacs. *Differential Games: A Mathematical Theory with Applications to Warfare and Pursuit, Control and Optimization*. Kreiger, Huntigton, NY, 1999.
- D. H. Jacobson and D. Q. Mayne. *Differential Dynamic Programming*. American Elsevier Publishing Company, Inc., New York, NY, 1970.
- J. F. C. Kingman. The single server queue in heavy traffic. *Mathematical Proceedings of the Cambridge Philosophical Society*, 57(4):902–904, 1961. doi: 10.1017/S0305004100036094.

347 Y. Li, N. Li, H. E. Tseng, A. Girard, D. Filev, and I. Kolmanovsky. Safe reinforcement learning using
348 robust action governor. In A. Jadbabaie, J. Lygeros, G. J. Pappas, P. A. Parrilo, B. Recht, C. J.
349 Tomlin, and M. N. Zeilinger, editors, *Proceedings of the 3rd Conference on Learning for Dynamics
350 and Control*, volume 144 of *Proceedings of Machine Learning Research*, pages 1093–1104. PMLR,
351 07–08 Jun 2021. URL <https://proceedings.mlr.press/v144/li21b.html>.

352 P.-L. Lions. *Generalized solutions of Hamilton-Jacobi equations*, volume 69. London Pitman, 1982.

353 J. D. Little. A proof for the queueing formula: $l = \lambda w$. *Operations Research*, 9(3):383–387,
354 1961.

355 W. E. Lorensen and H. E. Cline. Marching cubes: A high resolution 3d surface construction algorithm.
356 *ACM SIGGRAPH Computer Graphics*, 21(4):163–169, 1987.

357 J. Lygeros. On reachability and minimum cost optimal control. *Automatica*, 40(6):917–927, 2004.

358 A. Merz. The game of two identical cars. *Journal of Optimization Theory and Applications*, 9(5):
359 324–343, 1972.

360 I. Mitchell. A toolbox of level set methods, version 1.0. *The University of British Columbia, UBC CS
361 TR-2004-09*, pages 1–94, July 2004.

362 I. Mitchell. A Robust Controlled Backward Reach Tube with (Almost) Analytic Solution for Two
363 Dubins Cars. *EPiC Series in Computing*, 74:242–258, 2020.

364 I. M. Mitchell, A. M. Bayen, and C. J. Tomlin. A Time-Dependent Hamilton-Jacobi Formulation of
365 Reachable Sets for Continuous Dynamic Games. *IEEE Transactions on Automatic Control*, 50(7):
366 947–957, 2005. ISSN 00189286.

367 L. Molu. The Python LevelSet Toolbox (LevelSetPy). In *2024 IEEE 63rd Conference on Decision
368 and Control (CDC)*, pages 8938–8945, 2024a. doi: 10.1109/CDC56724.2024.10886640.

369 L. Molu. The python levelset toolbox (levelsetpy). In *IEEE 63rd Conference on Decision and Control
370 (CDC)*, pages 8938–8945, 2024b. doi: 10.1109/CDC56724.2024.10886640.

371 L. Molu. LevelSetPy: A GPU-Accelerated Package for Hyperbolic Hamilton-Jacobi Partial Differen-
372 tial Equations’ Solubility. *The ACM Transactions on Mathematical Software*, 2025.

373 S. Osher and R. Fedkiw. Level Set Methods and Dynamic Implicit Surfaces. *Applied Mechanics
374 Reviews*, 57(3):B15–B15, 2004. ISSN 0003-6900.

375 S. Osher and J. A. Sethian. Fronts Propagating with Curvature-Dependent Speed: Algorithms based
376 on Hamilton-Jacobi Formulations. *Journal of Computational Physics*, 79(1):12–49, 1988. ISSN
377 10902716.

378 M. L. Spearman, D. L. Woodruff, and W. J. Hopp. Conwip: a pull alternative to kanban. *The
379 International Journal of Production Research*, 28(5):879–894, 1990.

380 W. M. Wonham. Linear multivariable control: A geometric approach. *Applications of Mathematics*,
381 10, 1985.

A Background and Preliminaries (Unabridged).

This section provides a context for the contribution of this work.

B Background and Preliminaries

We first introduce the notations that are commonly used throughout the rest of this article. Reachable sets within the context of two person games [Isaacs, 1999] and their accompanying “viscous” terminal HJ PDE [Evans and Souganidis, 1984] are then introduced. This is followed by the HJ-Isaacs (HJI) PDEs commonly used to characterize reachable sets.

B.1 Notations and Terminologies

Conventions: Upper-case and lower-case bold font Roman letters are matrices and vectors, respectively; calligraphic letters are sets. Time variables e.g. t_0, t, τ, T will always be real positive numbers. The n -dimensional Euclidean space is $\mathbb{R}^n$; the set of natural numbers is $\mathbb{N}$. The natural log of T is denoted $\log T$. The state $\mathbf{x}$ belongs to the open set $\Omega \subseteq \mathbb{R}^n$. The HJ value function has a terminal value defined on the boundary of Ω , denoted $\partial\Omega$. The closure of Ω is $\bar{\Omega}$. The dot product is signified by $\langle \cdot, \cdot \rangle$. We will use D to denote the gradient of a function $\mathbf{x}$ of n variables and $\nabla \mathbf{w} = (\mathbf{w}_{x_1}, \dots, \mathbf{w}_{x_n}, \mathbf{w}_t) := (D\mathbf{w}, \mathbf{w}_t)$ as the full gradient of a function of $n+1$ variables, $\mathbf{w} = \mathbf{w}(\mathbf{x}, t)$. Where clarity is necessary, we may write $D_y \mathbf{w}$ to mean the spatial derivative of $\mathbf{w}$ with respect to the vector $\mathbf{y}$.

The unique solution $\xi \in \mathbb{R}^n$ of the dynamical system $\dot{\mathbf{x}}(\tau) = f(t; \mathbf{x}, \mathbf{u}, \mathbf{w})$, influenced by a pursuing player $\mathbf{P}$ (with control $\mathbf{w} \in \mathcal{W} \subseteq \mathbb{R}^p$) and its evading pair $\mathbf{E}$ (with control $\mathbf{u} \in \mathcal{U} \subseteq \mathbb{R}^m$), results in system trajectories $\xi_{\mathbf{x}, \tau}^{\mathbf{u}, \mathbf{w}}$ that are obtained by optimizing a scalar-valued payoff functional, $\ell(\xi_{\mathbf{x}, t}^{\mathbf{u}, \mathbf{w}}(\tau))$ i.e.

$$\mathcal{J}(\tau; \mathbf{x}, \mathbf{u}, \mathbf{w}) = \min_{\tau \in [t, T]} \ell(\xi_{\mathbf{x}, t}^{\mathbf{u}, \mathbf{w}}(\tau)) \quad (\text{Payoff function})$$

within a two-person differential game for a fixed $T \geq \tau > t$, where ℓ can be a signed distance to the (boundary of the) target set/region, for example. The controllers belong in compact sets which are measurable functions i.e. $\bar{\mathcal{U}} \equiv \mathbf{u} : [t, T] \rightarrow \mathcal{U}$, $\bar{\mathcal{W}} \equiv \mathbf{w} : [t, T] \rightarrow \mathcal{W}$.

B.2 Robustly Controlled Backward Reachable Set and Tube

Suppose that the goal of $\mathbf{P}$ is to drive the system into a user-specified target region $\mathcal{L}_0(\tau)$ within $\tau \leq T$ time steps of playing the game; while $\mathbf{E}$ simultaneously seeks to prevent this from happening. The invariant set $\mathcal{L}_0$ obtained at T , i.e.

$$\mathcal{L}_0(T) = \{\mathbf{x} \in \mathbb{R}^n \mid \mathbf{v}(0; \mathbf{x}) \leq 0\} \quad (\text{Target-Set})$$

is “robustly controlled” for the “distance-to-target-set” cost $\mathbf{g}(0; \mathbf{x})$, with constraints $|\mathbf{g}(0; \mathbf{x})| \leq k$, $|\mathbf{v}(0; \mathbf{x}) - \mathbf{v}(t; \hat{\mathbf{x}})| \leq k|\mathbf{x} - \hat{\mathbf{x}}|$ where $k > 0$ is a constant, and all $-T \leq t \leq 0$, $\{\hat{\mathbf{x}}, \mathbf{x} \in \mathbb{R}^n\}^1$. The distance to $\mathcal{L}_0(\tau)$ is typically found by optimizing $\ell(\mathbf{x}, t)$ as in (Payoff function). The HJI PDE (A.15) and Hamiltonian (A.16) for this set admit the form

$$\mathbf{v}_t(t; \mathbf{x}) + \min\{0, \mathbf{H}(t; \mathbf{x}, D\mathbf{v}(t; \mathbf{x}))\} = 0, \quad \mathbf{v}(0; \mathbf{x}) = \mathbf{g}(0; \mathbf{x}). \quad (\text{HJI-RCBRT})$$

For the safety problem setup in (A.15), the corresponding *robustly controlled backward reachable tube* (RCBRT) [Mitchell, 2020] on $(0, T]$ is the closure of the open set,

$$\mathcal{L}([-T, 0], \mathcal{L}_0) = \{\mathbf{x} \in \mathbb{R}^n \mid \exists \beta \in \mathcal{B}(t) \forall \mathbf{u} \in \mathcal{U}(t), \quad \exists \bar{\tau} \in [-T, 0], \xi(\bar{\tau}) \in \mathcal{L}_0\}, \quad (\text{Target-Tube})$$

Read: The set of states from which there exists a strategy of $\mathbf{P}$ such that, for all controls of $\mathbf{E}$, the resulting trajectory *reaches and remains in the target set* within the interval $[-T, 0]$. Following Lemma 2 of [Mitchell et al., 2005], the states in the reachable set admit the following properties w.r.t the value function $\mathbf{v}$,

$$\mathbf{x}(t) \in \mathcal{L}(\cdot) \implies \mathbf{v}(t; \mathbf{x}) \leq 0, \quad \mathbf{v}(t; \mathbf{x}) \leq 0 \implies \mathbf{x}(t) \in \mathcal{L}(\cdot). \quad (\text{Zero Levelset})$$

¹Time is reversed in BRT computational scenarios.

420 In *backward reach avoid tubes*, the agent must avoid the unsafe region at all times. We can write
 421 (HJI-RCBRT) a *robustly controlled backward reach-avoid tube* (RCBRAT) as,

$$\min\{v_t(t; \mathbf{x}) + \mathbf{H}(t; \mathbf{x}, Dv), g(t; \mathbf{x}) - \ell(t; \mathbf{x})\} \leq 0, \quad v(0; \mathbf{x}) = g(0; \mathbf{x}). \quad (\text{HJI-RCBRAT})$$

422 B.3 Dynamic Programming and Two-Person Games.

423 The formal relationships between the dynamic programming (DP) optimality condition for the *value*
 424 in differential two-person zero-sum games, and the solutions to PDEs that solve “min-max” or “max-
 425 min” type nonlinearity (the Isaacs’ equation) were presented in [Isaacs, 1999]. Essentially, Isaacs’
 426 claim was that if the *value* functions are smooth enough, then they solve certain first-order partial
 427 differential equations (PDE) problems with “max-min” or “min-max”-type nonlinearity. However,
 428 the DP value functions are seldom regular enough to admit a solution in the classical sense. “Weaker”
 429 solutions, on the other hand [Lions, 1982, Evans and Souganidis, 1984, Crandall et al., 1984, Crandall
 430 and Majda, 1980], provide generalized “viscosity” solutions to HJ PDEs under relaxed regularity
 431 conditions; these viscosity solutions are not necessarily differentiable anywhere in the state space, and
 432 the only regularity prerequisite in the definition is continuity [Crandall and Lions, 1983]. However,
 433 wherever they are differentiable, they satisfy the upper and lower values of HJ PDEs (discussed in
 434 § B.4) in a classical sense. Thus, they lend themselves well to many real-world problems existing at
 435 the interface of discrete, continuous, and hybrid systems [Lygeros, 2004, Osher and Sethian, 1988,
 436 Mitchell, 2020, Evans and Souganidis, 1984, Mitchell et al., 2005].

437 Matter-of-factly, viscosity solutions to *Cauchy-type*² HJ Equations are highly useful in backward
 438 reachability analysis [Mitchell et al., 2005]. For a state $\mathbf{x} \in \Omega$ and a fixed time t : $0 \leq t < T$, suppose
 439 that the set of all controls for players $\mathbf{P}$ and $\mathbf{E}$ are respectively

$$\bar{\mathcal{U}} \equiv \{\mathbf{u} : [t, T] \rightarrow \mathcal{U} | \mathbf{u} \text{ measurable}, \mathcal{U} \in \mathbb{R}^m\}, \quad \bar{\mathcal{W}} \equiv \{\mathbf{w} : [t, T] \rightarrow \mathcal{V} | \mathbf{w} \text{ measurable}, \mathcal{W} \subset \mathbb{R}^p\}. \quad (\text{A.2})$$

440 We are concerned with the differential equation,

$$\dot{\mathbf{x}}(\tau) = f(\tau, \mathbf{x}(\tau), \mathbf{u}(\tau), \mathbf{w}(\tau)), \quad \mathbf{x}(t) = \mathbf{x}, \quad T \leq \tau \leq t \quad (\text{A.3})$$

441 where $f(\tau, \cdot, \cdot, \cdot)$ and $\mathbf{x}(\cdot)$ are bounded and Lipschitz continuous. This bounded Lipschitz continuity
 442 property assures uniqueness of the system response $\mathbf{x}(\cdot)$ to controls $\mathbf{u}(\cdot)$ and $\mathbf{v}(\cdot)$ [Evans and
 443 Souganidis, 1984]. Associated with (A.3) is the payoff functional

$$\ell(\mathbf{u}, \mathbf{w}) = \ell(t; \mathbf{x}, \mathbf{u}(\cdot), \mathbf{w}(\cdot)) \equiv \int_t^T l(\tau, \mathbf{x}(\tau), \mathbf{u}(\tau), \mathbf{w}(\tau)) d\tau + g(\mathbf{x}(T)), \quad (\text{A.4})$$

444 where $g(\cdot) : \mathbb{R}^n \rightarrow \mathbb{R}$ satisfies

$$|g(\mathbf{x})| \leq k_1, \quad |g(\mathbf{x}) - g(\hat{\mathbf{x}})| \leq k_1 |\mathbf{x} - \hat{\mathbf{x}}| \quad (\text{A.5})$$

445 and $l : [0, T] \times \mathbb{R}^n \times \mathcal{U} \times \mathcal{W} \rightarrow \mathbb{R}$ is bounded and uniformly continuous, with

$$|l(t; \mathbf{x}, \mathbf{u}, \mathbf{w})| \leq k_2, \quad |l(t; \mathbf{x}, \mathbf{u}, \mathbf{w}) - l(t; \hat{\mathbf{x}}, \mathbf{u}, \mathbf{w})| \leq k_2 |\mathbf{x} - \hat{\mathbf{x}}| \quad (\text{A.6})$$

446 for constants k_1, k_2 and all $0 \leq t \leq T$, $\hat{\mathbf{x}}, \mathbf{x} \in \mathbb{R}^n$, $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$. We call T the *terminal time*
 447 (it may be infinity!) and the integral, when it does not depend on the control laws, is the *performance*
 448 *index*. The evader’s goal is to maximize the payoff (A.4) and pursuer’s goal is to minimize it.

449 B.4 Upper and Lower Values of the Differential Game.

450 Suppose that the pursuer’s mapping strategy (starting at t) is $\beta : \bar{\mathcal{U}}(t) \rightarrow \bar{\mathcal{W}}(t)$ provided for
 451 each $t \leq \tau \leq T$ and $\mathbf{u}, \hat{\mathbf{u}} \in \bar{\mathcal{U}}(t)$; then $\mathbf{u}(\bar{t}) = \hat{\mathbf{u}}(\bar{t})$ a.e. on $t \leq \bar{t} \leq \tau$ implies $\beta[\mathbf{u}](\bar{t}) =$
 452 $\beta[\hat{\mathbf{u}}](\bar{t})$ a.e. on $t \leq \bar{t} \leq \tau$. The differential game’s lower value for a solution $\mathbf{x}(t)$ that solves (A.3)
 453 for $\mathbf{u}(t)$ and $\mathbf{v}(t) = \beta[\mathbf{u}](\cdot)$ is

$$v^-(t; \mathbf{x}) = \inf_{\beta \in \mathcal{B}(t)} \sup_{\mathbf{u} \in \bar{\mathcal{U}}(t)} \ell(\mathbf{u}, \beta[\mathbf{u}]) \triangleq \inf_{\beta \in \mathcal{B}(t)} \sup_{\mathbf{u} \in \bar{\mathcal{U}}(t)} \int_t^T l(\tau, \mathbf{x}(\tau), \mathbf{u}(\tau), \beta[\mathbf{u}](\tau)) d\tau + g(\mathbf{x}(T)). \quad (\text{A.7})$$

²Cauchy-type HJ equations are time-dependent versions of the HJ PDE.

Similarly, suppose that the evader's mapping strategy (starting at t) is $\alpha : \bar{\mathcal{W}}(t) \rightarrow \bar{\mathcal{U}}(t)$ provided for each $t \leq \tau \leq T$ and $\mathbf{w}, \hat{\mathbf{w}} \in \bar{\mathcal{W}}(t)$; then $\mathbf{w}(\bar{t}) = \hat{\mathbf{w}}(\bar{t})$ a.e. on $t \leq \bar{t} \leq \tau$ implies $\alpha[\mathbf{w}](\bar{t}) = \alpha[\hat{\mathbf{w}}](\bar{t})$ a.e. on $t \leq \bar{t} \leq \tau$. The differential game's upper value for a solution $\mathbf{x}(t)$ that solves (A.3) for $\mathbf{u}(t) = \alpha[\mathbf{w}](\cdot)$ and $\mathbf{w}(t)$ is

$$\mathbf{v}^+(t; \mathbf{x}) = \sup_{\alpha \in \mathcal{A}(t)} \inf_{\mathbf{w} \in \bar{\mathcal{W}}(t)} \ell(\alpha[\mathbf{w}], \mathbf{w}) \triangleq \sup_{\alpha \in \mathcal{A}(t)} \inf_{\mathbf{w} \in \bar{\mathcal{W}}(t)} \int_t^T l(\tau, \mathbf{x}(\tau), \alpha[\mathbf{w}](\tau), \mathbf{w}(\tau)) d\tau + \mathbf{g}(\mathbf{x}(T)). \quad (\text{A.8})$$

These non-local PDEs ((A.7) and (A.8)) are hardly smooth throughout the state space so that they lack classical solutions even for smooth Hamiltonian and boundary conditions. However, these two values are "viscosity" (generalized) solutions [Lions, 1982, Crandall and Lions, 1983] of the associated HJ-Isaacs (HJI) PDE, i.e. solutions which are *locally Lipschitz* in $\Omega \times [0, T]$, and with at most first-order partial derivatives in the Hamiltonian.

B.5 Viscosity Solution of HJ-Isaac's Equations.

For any optimal control problem a value function is constructed based on the optimal cost (or payoff) of any input phase $(\mathbf{x}, T)$. In reachability analysis, typically this is defined using a terminal cost function $g(\cdot) : \mathbb{R}^n \rightarrow \mathbb{R}$ that satisfies

$$|g(\mathbf{x})| \leq k, \quad |g(\mathbf{x}) - g(\hat{\mathbf{x}})| \leq k|\mathbf{x} - \hat{\mathbf{x}}| \quad (\text{A.9})$$

for constant k and all $T \leq t \leq 0$, $\hat{\mathbf{x}}, \mathbf{x} \in \mathbb{R}^n$, $\mathbf{u} \in \mathcal{U}$ and $\mathbf{w} \in \mathcal{W}$. The zero sublevel set of $g(\mathbf{x})$ i.e.

$$\mathcal{L}_0 = \{\mathbf{x} \in \bar{\Omega} \mid g(\mathbf{x}) \leq 0\}, \quad (\text{A.10})$$

Lemma B.1. *The lower value $\mathbf{v}^-$ in (A.15) is the viscosity solution to the lower Isaac's equation*

$$\mathbf{v}_t^- + \mathbf{H}^-(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, D\mathbf{v}^-) = 0, \quad t \in [0, T], \quad \mathbf{x} \in \mathbb{R}^n, \quad \mathbf{v}^-(T; \mathbf{x}) = g(\mathbf{x}(T)), \quad \mathbf{x} \in \mathbb{R}^n \quad (\text{A.11})$$

with lower Hamiltonian,

$$\mathbf{H}^-(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, p) = \max_{\mathbf{u} \in \mathcal{U}} \min_{\mathbf{w} \in \mathcal{W}} \langle \mathbf{f}(t; \mathbf{x}, \mathbf{u}, \mathbf{w}), p \rangle. \quad (\text{A.12})$$

where p , the co-state, is the spatial derivative of $\mathbf{v}^-$ w.r.t $\mathbf{x}$.

Lemma B.2. *The upper value $\mathbf{v}^+$ in (A.8) is the viscosity solution of the upper Isaac's equation*

$$\mathbf{v}_t^+ + \mathbf{H}^+(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, D\mathbf{v}^+) = 0, \quad t \in [0, T], \quad \mathbf{x} \in \mathbb{R}^n, \quad \mathbf{v}^+(T; \mathbf{x}) = g(\mathbf{x}(T)), \quad \mathbf{x} \in \mathbb{R}^n \quad (\text{A.13a})$$

with upper Hamiltonian,

$$\mathbf{H}^+(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, p) = \min_{\mathbf{w} \in \mathcal{W}} \max_{\mathbf{u} \in \mathcal{U}} \langle \mathbf{f}(t; \mathbf{x}, \mathbf{u}, \mathbf{w}), p \rangle, \quad (\text{A.14})$$

with p being appropriately defined.

Corollary B.3. (i) $\mathbf{v}^- \leq \mathbf{v}^+$ over $(t \in [0, T], \mathbf{x} \in \mathbb{R}^n)$ (ii) if for all $t \in [0, T], (\mathbf{x}, p) \in \mathbb{R}^n$, the minimax condition is satisfied i.e. $\mathbf{H}^+(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, p) = \mathbf{H}^-(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, p)$, then $\mathbf{v}^- \equiv \mathbf{v}^+$.

B.6 Reachability for Systems Verification.

Reachability analysis is one of many verification methods that allows us to reason about (control-affine) dynamical systems. The verification problem may consist in finding a *set of reachable states* that lie along the trajectory of the solution to a first order nonlinear partial differential equation that originates from some initial state $\mathbf{x}_0 = \mathbf{x}(0)$ up to a specified time bound, $t = t_f$. From a set of initial and unsafe state sets, the time-bounded safety verification problem is to determine if there is an initial state and a particular time within the bound that the solution to HJI PDE enters the unsafe set.

Reachability could be analyzed in a (i) *forward* sense, whereupon system trajectories are examined to determine if they enter certain states from an *initial set*; (ii) *backward* sense, whereupon system trajectories are examined to determine if they enter certain *target sets*; (iii) *reach set* sense, in which

they are examined to see if states reach a set at a *particular time*; or (iv) *reach tube* sense, in which they are evaluated that they reach a set at a point *during a time interval*.

Backward reachability consists in avoiding an unsafe set of states under the worst-possible disturbance at all times; relying on nonanticipative control strategies. Backward reachable sets (BRS) and backward reachable tubes (BRTs) are popularly analyzed in a game of two vehicles with non-stochastic dynamics [Merz, 1972]. Such BRTs possess discontinuity at cross-over points (which exist at edges) on the surface of the tube, and may be non-convex. Therefore, treating the end-point constraints under these discontinuity characterizations need careful consideration and analysis when switching control laws if the underlying PDE does not have continuous partial derivatives (we discuss this further in § 2).

B.7 A Differential Game with Two Inputs

This section provides a cursory background for a minimal problem construction. Readers interested in a more detailed background can consult the thorough treatment in appendix A.

For a (possibly infinite) *terminal time* T , a constant k , controls $\mathbf{u} \in \bar{\mathcal{U}}$ (of $\mathbf{E}$), $\mathbf{w} \in \bar{\mathcal{W}}$ (of $\mathbf{P}$), and all $\hat{\mathbf{x}}, \mathbf{x} \in \mathbb{R}^n$, let the system possess a bounded and uniformly continuous (BUC) payoff functional $\mathbf{g}(\mathbf{x}(T))$ that satisfies $\mathbf{g} : \mathbb{R}^n \rightarrow \mathbb{R}$. The differential game's *lower value* [Evans and Souganidis, 1984],

$$v(t; \mathbf{x}) = \inf_{\beta \in \mathcal{B}(t)} \sup_{\mathbf{u} \in \mathcal{U}(t)} \mathcal{J}(t; \mathbf{x}, \mathbf{u}, \mathbf{w}) \triangleq \inf_{\beta \in \mathcal{B}(t)} \sup_{\mathbf{u} \in \mathcal{U}(t)} \min_{\tau \in [t, T]} \ell(\xi_{\mathbf{x}, t}^{\mathbf{u}, \mathbf{w}}(\tau)) \quad (\text{A.15})$$

is the viscosity solution to (HJ-IVP) [Evans and Souganidis, 1984, Lemma 2.1] with Hamiltonian,

$$\mathbf{H}(t; \mathbf{x}, \mathbf{u}, \mathbf{w}, p) = \max_{\mathbf{u} \in \mathcal{U}} \min_{\mathbf{w} \in \mathcal{W}} \langle f(t; \mathbf{x}, \mathbf{u}, \mathbf{w}), p \rangle. \quad (\text{A.16})$$

Here, p is the co-state and $\beta \in \mathcal{B}(t)$ is a *non-anticipative strategy* of the pursuer $\mathbf{P}$, i.e. a map $\beta : \bar{\mathcal{U}}(t) \rightarrow \bar{\mathcal{W}}(t)$ such that for every $t \leq \tau \leq T$ and every $\mathbf{u}, \hat{\mathbf{u}} \in \bar{\mathcal{U}}(t)$, the condition $\mathbf{u}(\bar{t}) = \hat{\mathbf{u}}(\bar{t})$ a.e. on $[t, \tau]$ implies $\beta[\mathbf{u}](\bar{t}) = \beta[\hat{\mathbf{u}}](\bar{t})$ a.e. on $[t, \tau]$ [Evans and Souganidis, 1984] (see § B.4 for a detailed treatment). *Note that the $O(N \cdot n)$ scaling applies to the evaluation of this max min Hamiltonian, which is often the most expensive part of grid-based solvers.*

B.8 Reachability from Differential Games Optimal Control

For any admissible control-disturbance pair $(\mathbf{u}(\cdot), \mathbf{w}(\cdot))$ and initial phase $(\mathbf{x}_0, t_0)$, Crandall [Crandall and Lions, 1983] and Evan's [Evans and Souganidis, 1984] claim is that there exists a unique function

$$\xi(t) = \xi(t; t_0, \mathbf{x}_0, \mathbf{u}(\cdot), \mathbf{w}(\cdot)) \quad (\text{A.17})$$

that satisfies (A.3) a.e. with the property that

$$\xi(t_0) = \xi(t_0; t_0, \mathbf{x}_0, \mathbf{u}(\cdot), \mathbf{w}(\cdot)) = \mathbf{x}_0. \quad (\text{A.18})$$

Read (A.17): the motion of (A.3) passing through phase $(\mathbf{x}_0, t_0)$ under the action of control $\mathbf{u}$, and disturbance $\mathbf{w}$, and observed at a time t afterwards. One way to design a system verification problem is compute the reachable set of states that lie along the trajectory (A.17) such that we evade the unsafe sets up to a time e.g. t_f within a given time bound e.g. $[t_0, t_f]$. In this regard, we discard the *cost-to-go*, $l(t; \mathbf{x}(\tau), \mathbf{u}(\tau), \mathbf{w}(\tau))$ in (A.4), (A.15), or (A.8) and certify safety as resolving the terminal value, $g(\mathbf{x}(T))$.

In backward reachability analysis, the lower value of the differential game is used in constructing an analysis of the backward reachable set (or tube). Therefore, we can cast a target set as the time-resolved terminal value $v^-(\mathbf{x}, T) = \mathbf{g}(\mathbf{x}(T))$ so that given a time bound, and an unsafe set of states, the time-bounded safety verification problem consists in certifying that there is no phase within the target set (Target-Set) such that the solution to (A.3) enters the unsafe set. Following the backward reachability formulation of [Mitchell et al., 2005], we say the zero sublevel set of $g(\cdot)$ in (A.11) i.e.

$$\mathcal{L}_0 = \{\mathbf{x} \in \bar{\Omega} \mid g(\mathbf{x}) \leq 0\}, \quad (\text{A.19})$$

is the *target set* in the phase space $\Omega \times \mathbb{R}$ for a backward reachability problem (proof in [Mitchell et al., 2005]). This target set can represent the failure set, regions of danger, or obstacles to be avoided

etc in the state space. Note that the target set, $\mathcal{L}_0$, is a closed subset of $\mathbb{R}^n$ and is in the closure of Ω . And the *robustly controlled backward reachable tube* for $\tau \in [-T, 0]^3$ is the closure of the open set

$$\mathcal{L}([\tau, 0], \mathcal{L}_0) = \{x \in \Omega \mid \exists \beta \in \bar{\mathcal{W}}(t) \forall u \in \mathcal{U}(t), \exists \bar{t} \in [-T, 0], \xi(\bar{t}) \in \mathcal{L}_0\}, \bar{t} \in [-T, 0]. \quad (\text{A.20})$$

Read: the set of states from which the strategies of P , and for all controls of E imply that we reach the target set within the interval $[T, 0]$. More specifically, following Lemma 2 of [Mitchell et al., 2005], the states in the reachable set admit the following properties w.r.t the value function v

$$x \in \mathcal{L}_0 \implies v^-(x, t) \leq 0, \quad v^-(x, t) \leq 0 \implies x \in \mathcal{L}_0. \quad (\text{A.21a})$$

Observe: (a) The goal of the pursuer, or P , is to drive the system's trajectories into the unsafe set i.e., P has w at will and aims to minimize the termination time of the game (c.f. (Target-Set)); (b) the evader, or E , seeks to avoid the unsafe set i.e., E has controls u at will and seeks to maximize the termination time of the game (c.f. (Target-Set)); (c) E has regular controls, u , drawn from a Lebesgue measurable set, $\mathcal{U}$ (c.f. (A.15)). (d) P possesses *nonanticipative strategies* (c.f. (A.15)) i.e. $\beta[u](\cdot)$ such that for any of the ordinary controls, $u(\cdot) \in \mathcal{U}$ of E , P knows how to optimally respond to E 's inputs. This is a classic reachability problem on the resolution of the infimum-supremum over the *strategies* of P and *controls* of E with the time of capture resolved as an extremum of a cost functional) over a time interval.

The terminal value $v^-(t; x)$, that characterizes the target set $\mathcal{L}_0$ is the viscosity solution to the HJI PDE

$$v_t^-(t; x) + \min\{0, H^-(t; x, u, w, Dv^-)\} = 0, \quad v^-(0; x) = g(x), \quad (\text{A.22})$$

where the vector field v_x^- is known in terms of the game's terminal conditions so that the overall game is akin to a two-point boundary-value problem. For ease of readability, we will remove the negative superscript on the lower value and Hamiltonian (A.16).

B.9 Detailed Hamiltonian Derivation for Two-Rockets Game

For the two-rockets pursuit-evasion game in relative coordinates, we derive the simplified Hamiltonian from first principles. The general Hamiltonian for a two-player zero-sum game is

$$H(t; x, p) = - \max_{u_e \in [\underline{u}_e, \bar{u}_e]} \min_{u_p \in [\underline{u}_p, \bar{u}_p]} [p_1 \quad p_2 \quad p_3] \begin{bmatrix} a_p \cos \theta + u_e x \\ a_p \sin \theta + a_e + u_p x - g \\ u_p - u_e \end{bmatrix}, \quad (\text{A.23})$$

where p_1, p_2, p_3 are the co-states (spatial derivatives of the value function), and $[\underline{u}_e, \bar{u}_e]$ and $[\underline{u}_p, \bar{u}_p]$ are the evader's and pursuer's control bounds, respectively. The game structure reflects the evader's ability to choose controls first (outer max over u_e) followed by the pursuer's response (inner min over u_p).

Setting symmetric control bounds $\underline{u}_e = \underline{u}_p = -1$ and $\bar{u}_e = \bar{u}_p = +1$, and using the symmetry assumption $a_e = a_p = a$, we optimize over the control variables:

$$H = - \max_{u_e \in [-1, 1]} \min_{u_p \in [-1, 1]} [p_1(a \cos \theta + u_e x) + p_2(a \sin \theta + a + u_p x - g) + p_3(u_p - u_e)]. \quad (\text{A.24})$$

Rearranging by control terms:

$$H = - \max_{u_e \in [-1, 1]} \min_{u_p \in [-1, 1]} [p_1 a \cos \theta + p_2 a (1 + \sin \theta) - p_2 g + (p_1 x + p_3) u_p + (p_2 x - p_3) u_e]. \quad (\text{A.25})$$

Since the pursuer minimizes over u_p , it chooses $u_p = -\text{sign}(p_1 x + p_3)$ to minimize the linear term. Similarly, the evader maximizes over u_e , choosing $u_e = \text{sign}(p_2 x - p_3)$. For bang-bang optimal controls, these become:

$$\min_{u_p} (p_1 x + p_3) u_p = -|p_1 x + p_3|, \quad (\text{A.26})$$

$$\max_{u_e} (p_2 x - p_3) u_e = |p_2 x - p_3|. \quad (\text{A.27})$$

Thus the simplified Hamiltonian becomes:

$$H(t; x, p) = - [ap_1 \cos \theta + p_2(a + a \sin \theta - g) - |p_1 x + p_3| - |p_2 x - p_3|], \quad (\text{A.28})$$

which, when substituted into the HJ PDE yields the simplified rockets HJI equation used in the numerical experiments (cf. Section 3.3).

³The (backward) horizon, $-T$ is negative for $T > 0$.

C HJ PDE Linearization

In this section, we construct the Cole-Hopf-like linearization of the viscous Hamilton-Jacobi (HJ) PDE and propose a sampling machinery we use in building our results. We show that the transformation is *exact* only when the Hamiltonian is quadratic in the co-state i.e. $\mathbf{H} = \frac{1}{2}\langle \mathbf{p}^\top, \mathbf{p} \rangle$, and we derive the precise residual that for general Hamiltonians. We then formulate the quasi-linearization iterative algorithm for computing (HJI-RCBRT-Visc) and (HJI-RCBRAT-Visc). We finish this appendix with the correct Gaussian expectation formulas for the value function and its spatial gradient.

C.1 The viscous HJ Equation's solution

In these proofs that follow, we will express the solutions to the viscous HJ equation as the logarithm of the expectation of Gaussian kernels that parameterize the state space. Let us construct the spatial and time derivatives to the solution to the viscous initial-value-HJ problem (HJ-Visc).

Proof of Proposition 2.1. Recall from (HJ-Visc) that

$$\mathbf{v}_t^\delta + \mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta) = \frac{\delta}{2} \Delta \mathbf{v}^\delta \quad \text{in } \mathbb{R}^n \times (0, T], \quad \mathbf{v}^\delta(0; \mathbf{x}) = \mathbf{g}(\mathbf{x}) \text{ on } \mathbb{R}^n \times \{t = 0\}, \quad (\text{B.1})$$

where $\delta > 0$ is the viscosity parameter, $\mathbf{H} : \mathbb{R} \times \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ is the Hamiltonian, and $\mathbf{g} : \mathbb{R}^n \rightarrow \mathbb{R}$ is the terminal/initial cost. Let $\omega^\delta = \varphi(\mathbf{v}^\delta(t; \mathbf{x}))$, where $\varphi : \mathbb{R} \rightarrow \mathbb{R}$ is a smooth function. Our goal is to choose φ so that ω^δ solves a linear partial differential equation. Differentiating, we find that

$$\omega_t^\delta = \varphi'(\mathbf{v}^\delta) \mathbf{v}_t^\delta \text{ and } \Delta \omega^\delta = \varphi''(\mathbf{v}^\delta) |D\mathbf{v}^\delta|^2 + \varphi'(\mathbf{v}^\delta) \Delta \mathbf{v}^\delta. \quad (\text{B.2})$$

Thus,

$$\omega_t^\delta = \varphi' \mathbf{v}_t^\delta \triangleq \varphi'(\mathbf{v}^\delta) \left(\frac{\delta}{2} \Delta \mathbf{v}^\delta - \mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta) \right), \quad (\text{B.3a})$$

$$= \frac{\delta}{2} \Delta \omega^\delta - \left(\frac{\delta}{2} \varphi''(\mathbf{v}^\delta) |D\mathbf{v}^\delta|^2 + \varphi'(\mathbf{v}^\delta) \mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta) \right). \quad (\text{B.3b})$$

Suppose that we choose,

$$\varphi^\delta = \exp \left(\frac{-2}{\delta} \cdot \frac{\mathbf{H}^\delta}{|D\mathbf{z}|^2} \mathbf{z} \right), \quad (\text{B.4})$$

as the solution to the residual $\mathbf{R}(t; \mathbf{x}) := \frac{\delta}{2} \varphi''(\mathbf{v}^\delta) |D\mathbf{v}^\delta|^2 + \varphi'(\mathbf{v}^\delta) \mathbf{H}^\delta = 0$, it follows that

$$\omega_t^\delta \triangleq \frac{\delta}{2} \Delta \omega^\delta. \quad (\text{B.5})$$

Set

$$\mathbf{c}(t; \mathbf{x}) = \frac{2}{\delta} \cdot \mathbf{H}^\delta / |D\mathbf{v}^\delta|^2 \quad (\text{B.6})$$

and let $\mathbf{c}(t; \mathbf{x})$ be locally frozen as a spatially-varying coefficient rather than a functional of $\mathbf{v}^\delta$.

Corollary C.1. *Note that $\mathbf{R}(t; \mathbf{x})$ becomes zero when φ^δ is a constant. This holds when $\mathbf{H}^\delta / |D\mathbf{v}^\delta|^2$ is independent of the phase $(\mathbf{x}, t)$. In these situations, $\mathbf{H}^\delta = \frac{1}{2}\langle \mathbf{p}^\top, \mathbf{p} \rangle$, $\mathbf{c} = 1/\delta = \text{const.}$ thus recovering Proposition 2.1.*

Then, dropping the templated arguments for ease of readability, if $\mathbf{v}^\delta$ solves (HJ-Visc), then we have the following approximate local transformation

$$\omega^\delta = \exp(-\mathbf{c} \mathbf{v}^\delta) \quad (\text{B.7})$$

solves the initial value problem for the differential equation

$$\omega_t^\delta - \frac{\delta}{2} \Delta \omega^\delta = \mathbf{R}(t; \mathbf{x}) \equiv 0 \quad \text{in } \Omega \times (0, T], \text{ and } \omega^\delta(0; \mathbf{x}) = \exp(-\mathbf{c}(0; \mathbf{x}) \mathbf{g}(\mathbf{x})) \quad \text{on } \Omega \times \{t = 0\}, \quad (\text{B.8})$$

589 where the **residual** is

$$\mathbf{R}(t; \mathbf{x}) = \left(\frac{\delta}{2} \varphi''(\mathbf{v}^\delta) |D\mathbf{v}^\delta|^2 + \varphi'(\mathbf{v}^\delta) \mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta) \right) \equiv 0. \quad (\text{B.9})$$

590 Equation (B.7) is a heat-like equation, similar to the Cole-Hopf transformation [Evans, 2022] for the
591 Eikonal version of (HJ-Visc)⁴.

592 It follows that the unique bounded fundamental solution of (B.8) is given by the Green's function
593 convolution,

$$\omega^\delta(t; \mathbf{x}) = \frac{1}{(\sqrt{2\pi\delta t})^n} \int_{\Omega} \exp\left(-\frac{1}{2\delta t} |\mathbf{x} - \mathbf{y}|^2\right) \exp(-c\mathbf{g}(\mathbf{y})) d\mathbf{y}, \quad (\mathbf{x} \in \Omega, t > 0) \quad (\text{B.10a})$$

$$\triangleq \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-c\mathbf{g}(\mathbf{y}))] := \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))] \quad (\text{B.10b})$$

594 where (B.10a) is the standard heat kernel convolution and the kernel has been rewritten as a Gaussian
595 density with mean $\mathbf{x}$ and variance δt . Note that $\omega^\delta(t; \mathbf{x}) = 0$ elsewhere. $\square$

596 *Proof of Lemma 2.2.* Going by equation (B.7), we can write

$$\mathbf{v}^\delta = -(\delta/2) \frac{|D\mathbf{v}^\delta|^2}{\mathbf{H}(t; \mathbf{x}, D\mathbf{v}^\delta)} \log \omega^\delta := -\frac{1}{c} \log \omega^\delta, \quad (\text{B.11})$$

597 so that the unique bounded solution to the initial-value (HJ-Visc) i.e. (5) becomes

$$\mathbf{v}^\delta(t; \mathbf{x}) = -\frac{1}{c} \cdot \log \left\{ \frac{1}{(\sqrt{2\pi\delta t})^n} \int_{\Omega} \exp\left(-\frac{1}{2\delta t} \cdot |\mathbf{x} - \mathbf{y}|^2\right) \exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y})) d\mathbf{y} \right\}, \quad (\text{B.12a})$$

$$\triangleq -\frac{1}{c} \cdot \log \left\{ \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))] \right\}. \quad (\text{B.12b})$$

598 **Corollary C.2.** For the (HJI-RCBRT-Visc) and (HJI-RCBRAT-Visc), $\mathbf{g}(\mathbf{y})$ the above equation
599 transforms to $\mathbf{g}(\mathbf{x} + \sqrt{\delta(T-t)}\mathbf{y})$. For numerical stability, we use the log-sum-exp identity

$$\mathbf{v}^\delta(t; \mathbf{x}) = \frac{-1}{c} \log \frac{1}{N} \sum_{i=1}^N \exp(-c\mathbf{g}(\mathbf{x} + \sqrt{\delta(T-t)}\mathbf{y}_i)) \quad \text{with } \mathbf{y}_i \sim \mathcal{N}(\mathbf{x}, \delta t).$$

600 *A fortiori*, we have the solution to the viscous HJ equation as the log of the expectation of a Gaussian
601 density with mean $\mathbf{x}$ and variance δt . $\square$

602 C.2 Spatial Gradient of the HJ Payoff

603 *Proof of Lemma 2.3.* From (B.10a) observe,

$$D\omega^\delta(t; \mathbf{x}) = -\frac{1}{(\sqrt{2\pi\delta t})^n} \int_{\Omega} \frac{(\mathbf{x} - \mathbf{y})}{\delta \cdot t} \exp\left(-\frac{|\mathbf{x} - \mathbf{y}|^2}{2\delta t}\right) \exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y})) d\mathbf{y}, \quad (\text{B.13a})$$

$$= -\frac{1}{\delta t} \mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [(\mathbf{x} - \mathbf{y}) \exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))]. \quad (\text{B.13b})$$

604 Inspecting (B.11), we may write

$$D\mathbf{v}^\delta(t; \mathbf{x}) = -\frac{1}{c} D[\log \omega^\delta(t; \mathbf{x})] = -\frac{1}{c} \frac{D\omega^\delta(t; \mathbf{x})}{\omega^\delta(t; \mathbf{x})} = \frac{1}{\delta \cdot t \cdot c} \cdot \frac{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [(\mathbf{x} - \mathbf{y}) \exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))]}{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))]}, \quad (\text{B.14})$$

$$\triangleq \frac{1}{t \cdot \delta \cdot c} \cdot \left(\mathbf{x} - \frac{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\mathbf{y} \exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))]}{\mathbb{E}_{\mathbf{y} \sim \mathcal{N}(\mathbf{x}, \delta t)} [\exp(-\mathbf{c} \cdot \mathbf{g}(\mathbf{y}))]} \right). \quad (\text{B.15})$$

605 $\square$

⁴For the Eikonal version of (HJ-Visc), we set $\mathbf{H}(D\mathbf{v}(\mathbf{x})) = 0$ in Ω , and $\mathbf{v}(0; \mathbf{x}) = 0$ on $\partial\Omega$.

606 **Corollary C.3** (Log-sum-exp estimator for the value function). *For (HJI-RCBRT-Visc) and*
 607 *(HJI-RCBRAT-Visc), the terminal cost argument transforms to $\mathbf{g}(\mathbf{x} + \sqrt{\delta(T-t)}\mathbf{y}_i)$. For numerical*
 608 *stability the log-sum-exp identity gives,*

$$\mathbf{v}^\delta(t; \mathbf{x}) = \frac{-1}{\mathbf{c}^{(k)}} \log \frac{1}{N} \sum_{i=1}^N \exp\left(-\mathbf{c}^{(k)} \mathbf{g}\left(\mathbf{x} + \sqrt{\delta(T-t)}\mathbf{y}_i\right)\right), \quad (\text{B.16})$$

609 with $\mathbf{y}_i \stackrel{iid}{\sim} \mathcal{N}(\mathbf{0}, I_n)$.

610 **Corollary C.4** (Monte Carlo gradient estimator). *With samples $\mathbf{s}_i = \mathbf{x} + \sqrt{\delta(T-t)}\mathbf{y}_i$, $\mathbf{y}_i \stackrel{iid}{\sim}$*
 611 *$\mathcal{N}(\mathbf{0}, I_n)$, the importance-weighted estimator for (7) is,*

$$D\mathbf{v}^\delta(t; \mathbf{x}) = \frac{1}{t \cdot \delta \cdot \mathbf{c}^{(k)}} \left(\mathbf{x} - \frac{\frac{1}{N} \sum_{i=1}^N \mathbf{s}_i \cdot \exp(-\mathbf{c}^{(k)} \mathbf{g}(\mathbf{s}_i))}{\frac{1}{N} \sum_{i=1}^N \exp(-\mathbf{c}^{(k)} \mathbf{g}(\mathbf{s}_i))} \right). \quad (\text{B.17})$$

612 C.3 Sampling Complexity

613 **Assumption C.5.** Let $V = \mathbb{R}^M$ with the sup norm $\|v\|_\infty = \max_{1 \leq m \leq M} |v_m|$, where $v_m \approx$
 614 $\mathbf{v}(t; \mathbf{x}_m)$. Assume the following hold on a closed admissible set $\mathcal{A} \subset V$.

- 615 1. There exists $G > 0$ such that $|\mathbf{g}(\mathbf{x})| \leq G$ for all $\mathbf{x} \in \mathbb{R}^n$.
- 616 2. There exist constants $0 < m_0 \leq P_*$ and $0 < c_{\min} \leq c_{\max} < \infty$ such that for every $v \in \mathcal{A}$
 617 and each m ,

$$m_0 \leq |p_m(v)| \leq P_*, \quad c_{\min} \leq (\Gamma(v))_m \leq c_{\max}. \quad (\text{B.18})$$

- 618 3. The Hamiltonian is Lipschitz in the co-state on the ball $|p| \leq P_*$: there exist constants
 619 $H_*, L_H > 0$ such that

$$|\mathbf{H}(t; \mathbf{x}_m, p)| \leq H_*, \quad |\mathbf{H}(t; \mathbf{x}_m, p) - \mathbf{H}(t; \mathbf{x}_m, q)| \leq L_H |p - q| \quad (\text{B.19})$$

620 for all $|p|, |q| \leq P_*$ and all m .

- 621 4. The reconstruction map is Lipschitz on $\mathcal{A}$: there exists $L_D > 0$ such that

$$\|\mathcal{G}(v) - \mathcal{G}(w)\|_\infty \leq L_D \|v - w\|_\infty, \quad \forall v, w \in \mathcal{A}. \quad (\text{B.20})$$

- 622 5. The admissible set is invariant under Λ , and the quantity

$$q \triangleq \frac{2G}{c_{\min}} \cdot \frac{2L_D}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_* P_*}{m_0^4} \right) \quad (\text{B.21})$$

623 satisfies $0 \leq q < 1$.

624 **Remark C.6.** Note that item 2 of Assumption C.5 excludes regions where the value-gradient degenerates,
 625 such as flat reachable interiors or singular shocks. Furthermore, under item 5, the contraction
 626 constant q may become unbounded in the inviscid limit, i.e. as $\delta \rightarrow 0$ and/or $m_0 \rightarrow 0$. In practice,
 627 numerical stabilization mechanisms such as viscosity regularization, coefficient clipping, or solu-
 628 tion ‘‘smearing’’ [Osher and Fedkiw, 2004] may be introduced to maintain stability of the discrete
 629 integration scheme.

630 *Proof of Theorem 2.5.* Recall that,

$$Z \triangleq \exp(-\mathbf{c} \mathbf{g}(\zeta)), \quad \mu \triangleq \mathbb{E}[Z], \quad \mathbf{v}_c(t; \mathbf{x}) \triangleq -\frac{1}{\mathbf{c}} \log \mu. \quad (\text{B.22})$$

631 For i.i.d. samples $\zeta_1, \dots, \zeta_N \sim \mathcal{N}(\mathbf{x}, \delta(T-t)I_n)$, $Z_i \triangleq \exp(-\mathbf{c} \mathbf{g}(\zeta_i))$, so that $\bar{Z}_N \triangleq \frac{1}{N} \sum_{i=1}^N Z_i$,
 632 and $\hat{\mathbf{v}}_{c,N}(t; \mathbf{x}) \triangleq -\frac{1}{\mathbf{c}} \log \bar{Z}_N$.

633 Let $a \triangleq e^{-\mathbf{c} g_{\max}}$, $b \triangleq e^{-\mathbf{c} g_{\min}}$. Since $g_{\min} \leq \mathbf{g}(\zeta) \leq g_{\max}$ and $\mathbf{c} > 0$, monotonicity of the
 634 exponential implies

$$e^{-\mathbf{c} g_{\max}} \leq e^{-\mathbf{c} \mathbf{g}(\zeta)} \leq e^{-\mathbf{c} g_{\min}} \quad (\text{B.23})$$

635 almost surely. By definition of a and b , this is exactly $a \leq Z \leq b$ almost surely. Hence each Z_i is
 636 bounded in $[a, b]$.

637 **Value residuals, $\bar{Z}_N$.** Since $v_c(t; \mathbf{x}) = -\frac{1}{c} \log \mu$ and $\hat{v}_{c,N}(t; \mathbf{x}) = -\frac{1}{c} \log \bar{Z}_N$, we find that

$$\hat{v}_{c,N}(t; \mathbf{x}) - v_c(t; \mathbf{x}) = -\frac{1}{c} \log \left(\frac{\bar{Z}_N}{\mu} \right). \quad (\text{B.24})$$

638 Hence, $|\hat{v}_{c,N}(t; \mathbf{x}) - v_c(t; \mathbf{x})| \geq \varepsilon$ implies that

$$\left| \log \left(\frac{\bar{Z}_N}{\mu} \right) \right| \geq c\varepsilon, \quad (\text{B.25})$$

639 or that $\bar{Z}_N \geq \mu \exp(c\varepsilon)$ or $\bar{Z}_N \leq \mu \exp(-c\varepsilon)$. Hence

$$\mathbb{P}(|\hat{v}_{c,N}(t; \mathbf{x}) - v_c(t; \mathbf{x})| \geq \varepsilon) \leq \mathbb{P}(\bar{Z}_N - \mu \geq \mu(\exp(c\varepsilon) - 1)) + \mathbb{P}(\bar{Z}_N - \mu \leq -\mu(1 - \exp(-c\varepsilon))). \quad (\text{B.26})$$

640 **Hoeffding bounds for tails.** Because $Z_1, \dots, Z_N$ are i.i.d. and each lies in $[a, b]$, Hoeffding's
641 inequality implies that for every $s > 0$,

$$\mathbb{P}(\bar{Z}_N - \mu \geq s) \leq \exp\left(-\frac{2Ns^2}{(b-a)^2}\right), \quad \mathbb{P}(\bar{Z}_N - \mu \leq -s) \leq \exp\left(-\frac{2Ns^2}{(b-a)^2}\right). \quad (\text{B.27})$$

642 Apply the first inequality with $s_+ \triangleq \mu(\exp(c\varepsilon) - 1)$ and the second with $s_- \triangleq \mu(1 - \exp(-c\varepsilon))$.
643 We obtain

$$\mathbb{P}(\bar{Z}_N - \mu \geq \mu(\exp(c\varepsilon) - 1)) \leq \exp\left(-\frac{2N\mu^2(\exp(c\varepsilon) - 1)^2}{(b-a)^2}\right), \quad (\text{B.28})$$

$$\mathbb{P}(\bar{Z}_N - \mu \leq -\mu(1 - \exp(-c\varepsilon))) \leq \exp\left(-\frac{2N\mu^2(1 - \exp(-c\varepsilon))^2}{(b-a)^2}\right). \quad (\text{B.29})$$

644 **Union bound and simplification.** Combining the previous bounds gives,

$$\mathbb{P}(|\hat{v}_{c,N}(t; \mathbf{x}) - v_c(t; \mathbf{x})| \geq \varepsilon) \leq \exp\left(-\frac{2N\mu^2(\exp(c\varepsilon) - 1)^2}{(b-a)^2}\right) + \exp\left(-\frac{2N\mu^2(1 - \exp(-c\varepsilon))^2}{(b-a)^2}\right). \quad (\text{B.30})$$

645 Since $\exp(c\varepsilon) - 1 \geq 1 - \exp(-c\varepsilon)$ for every $\varepsilon > 0$, the first exponential is no larger than the second.
646 Therefore

$$\mathbb{P}(|\hat{v}_{c,N}(t; \mathbf{x}) - v_c(t; \mathbf{x})| \geq \varepsilon) \leq 2 \exp\left(-\frac{2N\mu^2(1 - e^{-c\varepsilon})^2}{(b-a)^2}\right). \quad (\text{B.31})$$

647 This proves the claim of Theorem 2.5. □

648 *Proof of Corollary 2.7.* Since $a \leq Z$ almost surely, taking expectations yields $\mu = \mathbb{E}[Z] \geq a$. Hence
649 it is enough to require

$$2 \exp\left(-\frac{2Na^2(1 - e^{-c\varepsilon})^2}{(b-a)^2}\right) \leq \alpha. \quad (\text{B.32})$$

650 Taking logarithms and solving for N gives

$$N \geq \frac{(b-a)^2}{2a^2(1 - e^{-c\varepsilon})^2} \log \frac{2}{\alpha}, \quad (\text{B.33})$$

651 which proves the second claim. □

652 *Proof of Theorem 2.8.* We proceed in several steps.

653 **Step 1: Lipschitz continuity of the frozen-coefficient solve map Φ .** Fix an index m and define

$$\psi_m(s) \triangleq -\frac{1}{s} \log M_m(s), \quad M_m(s) \triangleq \mathbb{E}_{\zeta \sim \mathcal{N}(\mathbf{x}_m, \delta(T-t)I_n)} [\exp(-s\mathbf{g}(\zeta))]. \quad (\text{B.34})$$

654 Then $(\Phi(c))_m = \psi_m(c_m)$. Since $M_m(s) > 0$,

$$\psi'_m(s) = \frac{1}{s^2} \log M_m(s) - \frac{1}{s} \frac{M'_m(s)}{M_m(s)}. \quad (\text{B.35})$$

655 Differentiating under the expectation gives

$$M'_m(s) = \mathbb{E}_{\zeta \sim \mathcal{N}(\mathbf{x}_m, \delta(T-t)I_n)} [-\mathbf{g}(\zeta) \exp(-s\mathbf{g}(\zeta))]. \quad (\text{B.36})$$

656 Hence

$$-\frac{M'_m(s)}{M_m(s)} = \frac{\mathbb{E}[\mathbf{g}(\cdot)e^{-s\mathbf{g}(\cdot)}]}{\mathbb{E}[e^{-s\mathbf{g}(\cdot)}]}, \quad (\text{B.37})$$

657 which is a weighted expectation of $\mathbf{g}$ which has absolute value of at most G by assumption (i). In
658 addition, $|\mathbf{g}| \leq G$ implies

$$e^{-sG} \leq M_m(s) \leq e^{sG}. \quad (\text{B.38})$$

659 We can therefore write $|\log M_m(s)| \leq sG$. Substituting these two bounds into the derivative formula
660 yields

$$|\psi'_m(s)| \leq \frac{1}{s^2}(sG) + \frac{1}{s}G = \frac{2G}{s} \leq \frac{2G}{c_{\min}} \quad (\text{B.39})$$

661 for every $s \in [c_{\min}, c_{\max}]$. By the mean-value theorem,

$$|\psi_m(c_m) - \psi_m(\tilde{c}_m)| \leq \frac{2G}{c_{\min}} |c_m - \tilde{c}_m|. \quad (\text{B.40})$$

662 Taking the maximum over m , we must have

$$\|\Phi(c) - \Phi(\tilde{c})\|_{\infty} \leq \frac{2G}{c_{\max}} \|c - \tilde{c}\|_{\infty}. \quad (\text{B.41})$$

663 **Step 2: Lipschitz continuity of the coefficient-update map Γ .** Fix m and write $p \triangleq p_m(v)$,

664 $q \triangleq p_m(w)$. Then

$$|(\Gamma(v))_m - (\Gamma(w))_m| = \frac{2}{\delta} \left| \frac{\mathbf{H}(t; \mathbf{x}_m, p)}{|p|^2} - \frac{\mathbf{H}(t; \mathbf{x}_m, q)}{|q|^2} \right|. \quad (\text{B.42})$$

665 Add and subtract $\mathbf{H}(t; \mathbf{x}_m, q)/|p|^2$ to obtain

$$\left| \frac{\mathbf{H}(t; \mathbf{x}_m, p)}{|p|^2} - \frac{\mathbf{H}(t; \mathbf{x}_m, q)}{|q|^2} \right| \leq \frac{|\mathbf{H}(t; \mathbf{x}_m, p) - \mathbf{H}(t; \mathbf{x}_m, q)|}{|p|^2} + |\mathbf{H}(t; \mathbf{x}_m, q)| \left| \frac{1}{|p|^2} - \frac{1}{|q|^2} \right|. \quad (\text{B.43})$$

666 By assumption (ii), $|p|, |q| \geq m_0$, so that $|p|^2, |q|^2 \geq m_0^2$. Using assumption (iii), the first term is
667 bounded by

$$\frac{|\mathbf{H}(t; \mathbf{x}_m, p) - \mathbf{H}(t; \mathbf{x}_m, q)|}{|p|^2} \leq \frac{L_H}{m_0^2} |p - q|. \quad (\text{B.44})$$

668 For the second term, note that

$$\left| \frac{1}{|p|^2} - \frac{1}{|q|^2} \right| = \frac{||q|^2 - |p|^2|}{|p|^2 |q|^2} \leq \frac{||q| - |p||(|q| + |p|)}{m_0^4} \leq \frac{2P_*}{m_0^4} |p - q|, \quad (\text{B.45})$$

669 where the final inequality uses $|p|, |q| \leq P_*$ from assumption (ii). Using $|\mathbf{H}(t; \mathbf{x}_m, q)| \leq H_*$ from
670 assumption (iii), we conclude that

$$\left| \frac{\mathbf{H}(t; \mathbf{x}_m, p)}{|p|^2} - \frac{\mathbf{H}(t; \mathbf{x}_m, q)}{|q|^2} \right| \leq \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) |p - q|. \quad (\text{B.46})$$

671 Therefore

$$|(\Gamma(v))_m - (\Gamma(w))_m| \leq \frac{2}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) |p - q|. \quad (\text{B.47})$$

672 Finally, by assumption (iv),

$$|p - q| \leq \|\mathcal{G}(v) - \mathcal{G}(w)\|_\infty \leq L_D \|v - w\|_\infty. \quad (\text{B.48})$$

673 Taking the maximum over m yields

$$\|\Gamma(v) - \Gamma(w)\|_\infty \leq \frac{2L_D}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) \|v - w\|_\infty. \quad (\text{B.49})$$

674 **Step 3: $\Lambda = \Phi \circ \Gamma$ is a contraction.** Combining (B.41) and (B.49), we obtain

$$\|\Lambda(v) - \Lambda(w)\|_\infty = \|\Phi(\Gamma(v)) - \Phi(\Gamma(w))\|_\infty \leq \frac{2G}{c_{\min}} \|\Gamma(v) - \Gamma(w)\|_\infty, \quad (\text{B.50})$$

$$\leq \frac{2G}{c_{\min}} \cdot \frac{2L_D}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) \|v - w\|_\infty = q \|v - w\|_\infty. \quad (\text{B.51})$$

675 Assumption (v) states that $q < 1$, hence Λ is a contraction on $\mathcal{A}$.

676 **Step 4: Existence, uniqueness, and linear convergence.** Since $\mathcal{A}$ is closed in the Banach space
 677 $(V, \|\cdot\|_\infty)$ and invariant under Λ , Banach's fixed-point theorem applies. Therefore there exists a
 678 unique $v^* \in \mathcal{A}$ such that $\Lambda(v^*) = v^*$. Moreover, for every $v^{(0)} \in \mathcal{A}$, the sequence generated by
 679 $v^{(k+1)} = \Lambda(v^{(k)})$ converges to v^* and satisfies

$$\|v^{(k+1)} - v^*\|_\infty \leq q \|v^{(k)} - v^*\|_\infty \leq q^{k+1} \|v^{(0)} - v^*\|_\infty. \quad (\text{B.52})$$

680 This proves the first convergence statement.

681 **Step 5: Convergence of the coefficients.** Applying (B.49) with $w = v^*$ gives

$$\|\Gamma(v^{(k)}) - \Gamma(v^*)\|_\infty \leq \frac{2L_D}{\delta} \left(\frac{L_H}{m_0^2} + \frac{2H_*P_*}{m_0^4} \right) \|v^{(k)} - v^*\|_\infty. \quad (\text{B.53})$$

682 **Step 6: Residual decay.** Since $v^{(k+1)} = \Lambda(v^{(k)})$ and Λ is a contraction,

$$\|v^{(k+1)} - v^{(k)}\|_\infty = \|\Lambda(v^{(k)}) - \Lambda(v^{(k-1)})\|_\infty \leq q \|v^{(k)} - v^{(k-1)}\|_\infty. \quad (\text{B.54})$$

683 Applying this recursively yields

$$\|v^{(k+1)} - v^{(k)}\|_\infty \leq q^k \|v^{(1)} - v^{(0)}\|_\infty. \quad (\text{B.55})$$

684 **Step 7: A posteriori error estimate.** Because $v^{(k)} \rightarrow v^*$, we may write

$$v^* - v^{(k)} = \sum_{j=k}^{\infty} (v^{(j+1)} - v^{(j)}). \quad (\text{B.56})$$

685 Taking sup norms and using the triangle inequality,

$$\|v^* - v^{(k)}\|_\infty \leq \sum_{j=k}^{\infty} \|v^{(j+1)} - v^{(j)}\|_\infty \leq \sum_{j=k}^{\infty} q^{j-k} \|v^{(k+1)} - v^{(k)}\|_\infty \quad (\text{B.57})$$

$$= \frac{1}{1-q} \|v^{(k+1)} - v^{(k)}\|_\infty. \quad (\text{B.58})$$

686 Using once more that $\|v^{(k+1)} - v^{(k)}\|_\infty \leq q \|v^{(k)} - v^{(k-1)}\|_\infty$, we conclude that

$$\|v^{(k)} - v^*\|_\infty \leq \frac{q}{1-q} \|v^{(k)} - v^{(k-1)}\|_\infty. \quad (\text{B.59})$$

687 *A fortiori, this completes the proof.* □




Figure 2: Two Dubins' vehicles in relative Cartesian coordinates. Reprinted from Mitchell [2020].

D Further Numerical Results

This example was originally proposed by Merz [1972] as an iteration upon Isaacs [1999]'s homicidal chauffeur game, whereupon a pursuit-evasion game between two players with similar speeds and minimum turn radii, is thoroughly analyzed. In Mitchell [2020], this problem was established as a benchmark for testing the solubility of capturable set of states (the backward reachable tube) in Merz's classical pursuit-evasion game. In this example, we solve the problem with our LevelSetPy toolbox and establish that the approximated barrier surface to the two-player game conforms with standard results.

The game is that of two cars sharing similar Dubins dynamics Dubins [1957]: P and E both have a positive minimum turn radii, w , and constant speeds v – with motion restricted to a plane as we have for the rocket launch differential game above. In relative coordinates, the diagrammatic structure of the motion is as depicted in Fig. 2. Choosing the Cartesian coordinate for motion representation, the state vector of the game with E at the origin can be characterized by its x_1, x_2 position relative to P and the angle θ between the two vehicles. Capture occurs when the distance $\|PE\|_2$ between the pursuer and the evader becomes less than a specified radius.

The relative equations of motion, going by Fig. 2, is

$$\begin{pmatrix} \dot{x}_1 \\ \dot{x}_2 \\ \dot{x}_3 \end{pmatrix} = \begin{pmatrix} -v_e + v_p \cos x_3 + w_e x_2 \\ v_p \sin x_3 - w_e x_1 \\ w_p - w_e \end{pmatrix}. \quad (\text{C.1})$$

We adopt specialization to a case where the two vehicles only possess a unit velocity and unit maximum turn rates. Here, as Merz notes, if the initial velocities are parallel such as $x_3 = 0$, then the equations of relative motion imply that E can be separated from P forever by the initial radial separation if it replicates P 's strategy. Whence, the barrier surface is closed and we are presented with Isaacs [1999]'s game of kind where we must determine the nature of the surface. This terminal surface possesses a closed-form solution and we refer readers to the treatment by Merz [1972]. In this example, our chief concern is to judge the efficacy of our toolbox with respect to the analytical solution of the barrier surface.

The backward reachable tube that consists of the paths taken by the trajectories of either player is defined as in the rockets pursuit-evasion game so that we have

$$\ell(0, x) = \{x \in \mathcal{X} | x_1^2 + x_2^2 \leq r^2\}, \quad (\text{C.2})$$

where again r is the capture radius. The target set is a cylinder as ℓ above excludes the heading, x_3 . It is represented as shown in Fig. 3.

For a detailed treatment of the barrier surface, we refer readers to a proper analysis as elucidated in Mitchell [2020]. Here, we focus on the construction of the BUP. The set of states that constitute the useable part and its boundary are respectively a function of the implicit surface function representation




Figure 3: The target set (left) and the boundary of the useable part of the state space after the differential game between P and E .




Figure 4: Dubins Vehicle BRT Slices

719 $\ell : [-T, 0] \times \mathcal{X} \rightarrow \mathbb{R}$ so that for a $t \in [0, T]$, where $T > 0$ is

$$\mathcal{T} = \{x \in \mathcal{X} | \ell(0, x) \leq 0\} \quad (\text{C.3})$$

$$R([-t, 0], \mathcal{T}) = \{x \in \mathcal{X} | \ell(t, x) \leq 0\}, \quad (\text{C.4})$$

720 When $t > 0$, the implicit surface representation is the following HJI PDE

$$\frac{\partial}{\partial t} \ell(t, x) + \min(0, \mathbf{H}(x, \nabla_x \ell(t, x))) = 0. \quad (\text{C.5})$$

721 It is easy to verify that the Hamiltonian is

$$H(x, p) = p_1(v_p \cos x_3 - v_e) - p_2(v_p \sin x_3) + w|p_1 x_2 - p_2 x_1 - p_3| - w|p_3|. \quad (\text{C.6})$$

722 Since we are concerned with the special case that the linear and angular speeds are equal, we set

723 $v_e = v_p = w \triangleq +1$ in the foregoing so that the Hamiltonian, in the final analysis is

$$H(x, p) = p_1(\cos x_3 - 1) - p_2(\sin x_3) + |p_1 x_2 - p_2 x_1 - p_3| - |p_3|. \quad (\text{C.7})$$

724 The results and comparisons are provided in Fig. 4 and Fig. 10.

725 D.1 The Double Integral Plant

726 Here, we analyze a time-optimal control problem to determine what admissible control⁵ can “trans-
727 port” the system under consideration to a desired “origin” in the shortest possible time. We consider
728 the double integral plant Bhat and Bernstein [1998], Wonham [1985] as an illustrative example of our

⁵A control law is admissible when its range belongs in the admissible input set where it is bounded.




Figure 5: State trajectories of the double integral plant. The solid curves with upward-pointing arrows are trajectories generated under $u = +1$ (switching curve γ_+); dashed curves with downward-pointing arrows are trajectories under $u = -1$ (switching curve γ_-).

729 objective, which is to compute the points in the state space that can reach the origin in *finite-time*
 730 under the influence of a time-optimal controller.

731 We shall leverage standard necessary conditions from the principle of optimality Bellman [1957]
 732 to obtain a time-optimal feedback control design; introduce the notion of isochrones and switching
 733 surfaces; and discuss the analytic and approximate solutions (with our library) to the time-optimal
 734 control problem for a double integrator. We shall conclude the section by comparing the analytic
 735 and the overapproximated numerical solution (using the LevelSetPy toolbox) to the time to reach the
 736 origin problem.

737 D.1.1 Dynamics and Problem Setup

738 The double integrator is controllable, so that open-loop strategies may be employed in driving specific
 739 states to the origin in finite time Wonham [1985]. The plant has the following second-order dynamics

$$\ddot{x}(t) = u(t) \quad (C.8)$$

740 and admits bounded control signals $|u(t)| \leq 1$ for all time t . After a change of variables, we have
 741 the following system of first-order differential equations

$$\begin{aligned} \dot{x}_1(t) &= x_2(t), \\ \dot{x}_2(t) &= u(t), \quad |u(t)| \leq 1. \end{aligned} \quad (C.9)$$

742 The *reachability problem* that we consider is to address the question of what states can reach a
 743 certain point (here, the origin) in a transient manner. That is, we would like to find point sets on the
 744 state space, at a particular time step, such that we can bring the system to the equilibrium, $(0, 0)$.

745 D.1.2 Time-optimal control scheme

746 This is an H -minimal control problem whereupon we must find the control law that minimizes the
 747 Hamiltonian

$$H(x, p) = p_1 \dot{x}_1 + p_2 \dot{x}_2. \quad (C.10)$$

748 The necessary optimality condition stipulates that the minimizing control law be

$$u(t) = -\text{sign}(p_2(t)) \triangleq \pm 1. \quad (C.11)$$




Figure 6: The time to reach the origin as 2D slice comparison for the double integrator plant

Table 3: Double integrator BRS; Algorithm 1 convergence at backward times $t \in \{0.8, 0.5, 0.2, 0.0\}$ with $N = 8,000$ and $\delta = 0.08$.

t	Iterations	Final Residual
0.80	15	0.0143
0.50	15	0.0334
0.20	15	0.0596
0.00	15	0.0828

For the co-states in question, suppose that their initial values (for constants k_1 and k_2) are $p_1(t_0) = k_1$ and $p_2(t_0) = k_2$, only four candidates can serve as time-optimal control sequences i.e. $\{[+1], [-1], [+1, -1], [-1, +1]\}$. On a finite time interval, $t \in [t_0, t_f]$, the time-optimal $\mathbf{u}(t)$ is a constant $k \equiv \pm 1$ so that for initial conditions $\mathbf{x}_1(t_0) = \boldsymbol{\xi}_1$ and $\mathbf{x}_2(t_0) = \boldsymbol{\xi}_2$, it can be verified that the state trajectories obey the relation

$$\mathbf{x}_1(t) = \boldsymbol{\xi}_1 + \frac{1}{2}k(\mathbf{x}_2^2 - \boldsymbol{\xi}_2^2), \text{ for } t = k(\mathbf{x}_2(t) - \boldsymbol{\xi}_2). \quad (\text{C.12})$$

The trajectories of (C.12) traced out over a finite time horizon $t = [-1, 1]$ with *piecewise constant control laws*, $u = \pm 1$ on a state space and under the control laws $\mathbf{u}(t) = \pm 1$ is depicted in Fig. 5. Curves with arrows that point upwards denote trajectories under the control law $\mathbf{u} = +1$; call these trajectories γ_+ ; while the trajectories marked by dashed arrows pointing downward on the curves were executed under $\mathbf{u} = -1$; call these trajectories γ_- .

Table 3 reports convergence at four backward times; all cases achieve 15 iterations with residuals well below the $O(\sqrt{\delta}) \approx 0.28$ bound.

D.2 The Game of Two rockets on a Plane

We adopt the rocket launch problem of Dreyfus Dreyfus [1966] which is to launch a rocket in fixed time to a desired altitude, given a final vertical velocity component and a maximum final horizontal component as constraints. The rocket's motion is dictated by the following differential equations (under Dreyfus' assumptions)

$$\dot{x}_1 = x_3; \quad x_1(t_0) = 0; \quad (\text{C.13a})$$

$$\dot{x}_2 = x_4, \quad x_2(t_0) = 0; \quad (\text{C.13b})$$

$$\dot{x}_3 = a \cos u, \quad x_3(t_0) = 0; \quad (\text{C.13c})$$

$$\dot{x}_4 = a \sin u - g, \quad x_4(t_0) = 0; \quad (\text{C.13d})$$

where, (x_1, x_2) are respectively the horizontal and vertical range of the rockets (in feet), (x_3, x_4) are respectively the horizontal and vertical velocities of the rockets (in feet per second), while a and g are respectively the acceleration and gravitational accelerations (in feet per square second). Being a free endpoint problem, we transform it into a game between two players (C.13) without the terminal time constraints as defined in Jacobson and Mayne [1970]. The states of $\mathbf{P}$ and $\mathbf{E}$ are now denoted as (x_p, x_e) respectively which are driven by their thrusts (u_p, u_e) respectively in the xz -plane (see




Figure 7: *Left*: Motion of two rockets on a Cartesian xz -plane with a thrust inclination in relative coordinates given by $\theta := u_p - u_e$. *Right*: Representation of the initial values as a heat map for the two rockets described in (16).

Figure 7). The relevant kinematic equations are (C.13b) and (C.13d).

$$\dot{x}_{2e} = x_{4e}; \quad \dot{x}_{2p} = x_{4p}, \quad (\text{C.14a})$$

$$\dot{x}_{4e} = a \sin u_e - g, \quad \dot{x}_{4p} = a \sin u_p - g \quad (\text{C.14b})$$

where a and g are respectively the acceleration and gravitational accelerations (in feet per square second) i.e. $a = 64 \text{ ft/sec}^2$ and $g = 32 \text{ ft/sec}^2$. We reformulate the problem as a two-player differential game where the optimization operations remain in the interior, avoiding discontinuous switches in control. This formulation implicitly bounds each rocket's speed through the gravitational dynamics; the natural speed limit is $a \sin u/g$, which corresponds to each rocket's asymptotic speed when launched vertically.

Therefore, we rewrite (C.13) with P 's motion relative to E 's along the (x, z) plane so that the relative orientation as shown in Fig. 7 is $\theta = u_p - u_e$. The coordinates of P are freely chosen; however, the coordinates of E are chosen a distance r away from (x, z) so that the EP vector's inclination measured counterclockwise from the x -axis is θ . Following the conventions in Fig. 7, the game's relative equations of motion in reduced space is $\mathcal{X} = (x, z, \theta)$ where $\theta \in [-\frac{\pi}{2}, \frac{\pi}{2})$ and $(x, z) \in \mathbb{R}^2$ are

$$\dot{x} = a_p \cos \theta + u_e x, \quad (\text{C.15a})$$

$$\dot{z} = a_p \sin \theta + a_e + u_e x - g, \quad (\text{C.15b})$$

$$\dot{\theta} = u_p - u_e. \quad (\text{C.15c})$$

The payoff, Φ , is the distance of P from E when capture occurs denoted as $\|PE\|_2$. Capture occurs when $\|PE\|_2 \leq r$ for a pre-specified capture radius, $r > 0$. In (??), we say P controls u_p and is minimizing Φ , and E controls u_e and is maximizing P . The boundary of the *usable part* of the origin-centered circle of radius r^6 is $\|PE\|_2$ so that

$$r^2 = x^2 + z^2, \quad (\text{C.16a})$$

and all capture points are specified by

$$\dot{r}(x, t) + \min \left[0, \mathbf{H}(\mathbf{x}, \frac{\partial r(x, t)}{\partial \mathbf{x}}) \right] \leq 0, \quad (\text{C.17})$$

with the corresponding Hamiltonian

$$\mathbf{H}(\mathbf{x}, p) = - \max_{u_e \in \mathcal{U}_e} \min_{u_p \in \mathcal{U}_p} \begin{bmatrix} p_1 \\ p_2 \\ p_3 \end{bmatrix}^T \begin{bmatrix} a_p \cos \theta + u_e x \\ a_p \sin \theta + a_e + u_p x - g \\ u_p - u_e \end{bmatrix}. \quad (\text{C.18})$$

⁶We set $r = 1.5 \text{ ft}$ in our evaluations.




Figure 8: Initial and final backward reachable tubes for the rocket system (cf. Fig. 7) computed using the method outlined in Evans and Souganidis [1984], Osher and Sethian [1988], Mitchell [2004]. We set $a_e = a_p = 64 \text{ ft/sec}^2$ and $g = 32 \text{ ft/sec}^2$ as in Dreyfus' original example. We compute the reachable set by optimizing for the paths of slowest-quickest descent in equation (??).

Suppose that the maximizing u_e is $\bar{u}_e$ and the minimizing u_p is $\bar{u}_p$. We have at the point of slowest-quickest descent on the capture surface, that

$$\bar{u}_e = p_1 x - p_3, \quad (\text{C.19a})$$

$$\bar{u}_p = p_3 - p_2 x. \quad (\text{C.19b})$$

We set the linear velocities and accelerations equal to one another i.e. $u_e = u_p$ and $a_e = a_p$. Thus, the Hamiltonian takes the form

$$\begin{aligned} H(x, p) = & -\cos(u)|ap_1| + \cos(u)|ap_1| - \sin(u)|ap_2| - \\ & \sin(u)|ap_2| + u|p_3| - u|p_3|. \end{aligned} \quad (\text{C.20})$$

Using a distributed version of the levelset toolbox Molu [2024a], the backward reachable tube of the game is depicted in Fig. 8. A game between the two players was run over 11 global optimization time steps. The initial value function (left inset of Fig. 8) is represented as a dynamic implicit surface over all point sets in the state space with a signed distance function. We made the third coordinate axis of the state space (here the common heading of the two rockets) to align with the third cylinder axis. The final BRT at the end of the optimization run is shown in the right inset of Fig. 8.

Two-Rockets 3D Backward Reachable Tube (zero level set)
levelsetpy: 45³ grid | MC: 25³ grid, 14000 samples
levelsetpy (inviscid) MC Cole-Hopf ($\delta=0.08$)




Figure 9: Initial and final backward reachable tubes for the rocket system (cf. Fig. 7) computed using the method outlined in Evans and Souganidis [1984], Osher and Sethian [1988], Mitchell [2004]. We set $a_e = a_p = 64 \text{ ft/sec}^2$ and $g = 32 \text{ ft/sec}^2$ as in Dreyfus’ original example. We compute the reachable set by optimizing for the paths of slowest-quickest descent in equation (??).

E Limitations

Discrete bang-bang control: a cautionary case study in manufacturing dispatch. During development, we attempted to extend the method to manufacturing flow control; a 2D problem arising from a six-station serial production line where the bottleneck (constraint) station S_2 has much lower throughput than downstream station S_3 . The core question is: when should the factory dispatch (release) parts to each station to avoid deadlock (queue overflow at S_2) and starvation (queue depletion at S_3)? This requires solving a 2D backward reachability problem in state space (q_c, q_d) , where q_c is the queue length at the constraint station and q_d is the queue length at the downstream station. Figure 11 illustrates the six-station structure; the bottleneck S_2 (shaded red) creates the coupled queue dynamics that our method must solve.

The discrete control structure emerges naturally; at each station, dispatch is binary; enable processing ($u = 1$) or hold ($u = 0$). With two stations, this yields four bang-bang control combinations $(u_c, u_d) \in \{0, 1\}^2$. In principle, our algorithm should handle this; we simply enumerate the four controls and compute the Hamiltonian as $H(x, p) = \min_{(u_c, u_d)} [p_1(\lambda - \mu_c u_c) + p_2(\mu_c u_c - \mu_d u_d)]$.

However, the bang-bang structure creates a fundamental incompatibility with the frozen-coefficient scheme. The difficulty is a sharp phenomenon in the Hamiltonian landscape; as one moves across the boundary separating two control regimes (e.g., crossing the surface where u_c switches from 0 to 1), the optimal control law exhibits *shocks*; discontinuous jumps in the control value and its co-state sensitivities. The four bang-bang combinations partition the (q_c, q_d) space into four linear regions, separated by switching surfaces where the optimal control is discontinuous. Near these boundaries, the Hamiltonian and its gradient $|Dv^\delta|$ exhibit sharp concentration and oscillation.

Our frozen-coefficient scheme, which fixes the spatial gradient $|Dg|$ over an entire Picard iteration, is fundamentally ill-suited to track this piecewise-linear switching structure. In continuous-control problems (e.g., thrust magnitude in pursuit-evasion), the Hamiltonian landscape is smooth and the frozen coefficient remains a reasonable local approximation over 15 Picard steps. But in the manufacturing dispatch case, the fixed gradient becomes increasingly stale as the iterate approaches a switching boundary; the quasi-linear residuals plateau and then diverge. In practice, we observed NaN blow-up in the log-sum-exp evaluator by Picard iteration 7–10, and when we forced continuation, the value-function estimates disagreed 79

Root-cause analysis reveals that Theorems 2.5 and 2.8 both presuppose regularity conditions; Lipschitz continuity of the Hamiltonian, bounded second derivatives, and nondegeneracy of the control-to-dynamics mapping; that bang-bang controls violate precisely at switching surfaces. The theoretical




Figure 10: Dubins Vehicles' BRT Slices: Levelsetpy vs Monte-Carlo (Ours). The bottom row shows a heat map of the errors at various relative vehicle orientations (the control input) between the Monte Carlo sampling scheme and the LevelSetPy scheme.

scaffolding does not apply. We therefore recommend restricting the method to continuous control systems, or discrete problems with *smooth* bang-bang envelopes (e.g., sigmoid approximations to step functions that avoid true discontinuities). Discrete combinatorial control with genuine switching remains an open challenge for sampling-based HJ methods and may require fundamentally different machinery; hybrid dynamical systems theory, piecewise-smooth analysis, or augmented Lagrangian formulations that explicitly track switching surfaces.

Broader scope limitations. Beyond the bang-bang case, several other fundamental limitations should be acknowledged:

(i) *Sampling overhead in moderate dimensions:* Although the $O(N \cdot n)$ scaling becomes favorable as n increases, the absolute cost of $N = 14,000$ samples per iteration can be higher than structured grid methods for $n \leq 4$; practitioners should use grid-based solvers for low-dimensional problems where memory is not the bottleneck.

(ii) *Viscosity parameter sensitivity:* The choice of δ controls the smoothing level and approximation error $O(\sqrt{\delta})$; suboptimal δ selection can degrade accuracy or require excessive samples. Grid-search or Bayesian optimization is recommended during tuning.




Figure 11: Six-station manufacturing line with bottleneck at S_2 . Queue state (q_c, q_d) represents lengths at constraint and downstream stations; binary controls $u_c, u_d \in \{0, 1\}$ enable/disable dispatch.

- 849 (iii) *Convergence dependent on regularity*: The Picard iteration convergence guarantee (Theorem 2.8)
850 requires Lipschitz Hamiltonians and bounded second derivatives; nonconvex or highly nonlinear
851 Hamiltonians may converge slowly or not at all.
- 852 (iv) *Gradient estimation noise*: The spatial gradient estimate (7) relies on Monte Carlo variance in
853 the log-sum-exp; in high-viscosity regimes or when samples become concentrated in narrow regions,
854 gradient estimation can become unstable.
- 855 (v) *Limited applicability to learned approximators*: The theoretical guarantees (Theorems 2.5 and 2.8)
856 assume exact Hamiltonian evaluation. Extension to neural-network-approximated Hamiltonians (e.g.,
857 in learned dynamics models) introduces additional approximation error not quantified by our theory.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: [Yes] The main claims of the paper as mentioned in the abstract and introduction concern a grid-free and storage-free computation of reachable sets via Monte Carlo sampling of Gaussian expectations. Theoretical insights that justify this numerical contribution are introduced with rigorous proofs and there are several baseline experiments that signal the efficacy of our introduced $O(N \cdot n)$ rather than $O(M^n)$ memory requirements for existing solvers, where M , N , and n respectively denote number of discretization points for a grid, Monte Carlo samples, and state space dimension.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: [Yes] Yes, the limitations section in the paper explains the efficacy of the algorithm only in quadratic cases. For frozen coefficients Picard-iteration, there is a tradeoff in the resolution of the reachable set owing to the Monte Carlo rollouts employed. In the future, a better approximator that makes the MC sampling tight with respect to the levelset benchmark is worth an exploration.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.

- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: [Yes] All assumptions are stated explicitly within each theorem and lemma environment. Full proofs appear in Appendix A (background) and Appendix B (main proofs), including proofs of Lemmas 1–2, Corollaries 1–2, and Theorems 1–2. Proof sketches are provided inline where appropriate.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: [Yes] The algorithm that specified the results reproduction is delineated in the paper. The authors are willing to release the codebase for the experiments once a decision is made regarding our paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.

- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No] The code is not yet publicly hosted, but will be released upon acceptance. To preserve anonymity at submission time, an anonymized ZIP archive of the implementation is included in the supplementary materials; reviewers may use it to reproduce all reported experiments.

Justification: [No] Anonymized supplementary code is provided as a ZIP file accompanying this submission.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: [Yes] Experimental hyperparameters (number of Monte Carlo samples N , viscosity δ , time horizon T , grid resolution for baselines) are specified in the results section and captions. All experiments were run on a standard workstation with an Intel Core i9 CPU and 64 GB RAM; no GPU was required. Software dependencies (NumPy, SciPy, levelsetpy) are cited in the paper and listed in the supplementary code archive.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.

- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: [Yes] Quantitative comparisons report L^2 and Hausdorff errors against the levelsetpy ground truth. Where Monte Carlo estimates are shown, the stochastic variability is controlled by the sample-complexity bound of Theorem 1, which provides explicit (ε, δ) -guarantees on estimation error as a function of N . No separate confidence-interval bars are plotted because each experiment uses a fixed, deterministic evaluation grid; the theoretical bound serves as the statistical certificate.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.).
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: [Yes] All experiments were executed on a single CPU workstation (Intel Core i9, 64 GB RAM). Individual runs range from under one second (low-dimensional cases, $n \leq 3$) to a few minutes (6-dimensional Dubins car / pursuit-evasion). No GPU or cluster was used. The total compute for all reported experiments is under two CPU-hours, in contrast to the hours-to-days required by grid-based solvers at equivalent resolution. These figures are summarized in the results section.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: [Yes] This work is purely foundational, developing mathematical and algorithmic tools for safety-critical reachability analysis. It involves no human subjects, no personal data, no deception, and no dual-use risks. The authors have reviewed the NeurIPS Code of Ethics and confirm full compliance.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: [Yes] Positive impacts: scalable backward reachability analysis directly enables formal safety verification for autonomous vehicles, robotic systems, and aircraft, applications where current grid-based solvers are computationally infeasible. Negative impacts: the method could in principle be applied to model adversarial pursuit strategies; however, the same analysis applies to existing reachability solvers. No foreseeable misuse risk is specific to this approach. A brief broader-impacts discussion is included in the paper's conclusion.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

1123 Justification: [N/A]

1124 Guidelines:

- 1125 • The answer [N/A] means that the paper poses no such risks.
- 1126 • Released models that have a high risk for misuse or dual-use should be released with
- 1127 necessary safeguards to allow for controlled use of the model, for example by requiring
- 1128 that users adhere to usage guidelines or restrictions to access the model or implementing
- 1129 safety filters.
- 1130 • Datasets that have been scraped from the Internet could pose safety risks. The authors
- 1131 should describe how they avoided releasing unsafe images.
- 1132 • We recognize that providing effective safeguards is challenging, and many papers do
- 1133 not require this, but we encourage authors to take this into account and make a best
- 1134 faith effort.

1135 **12. Licenses for existing assets**

1136 Question: Are the creators or original owners of assets (e.g., code, data, models), used in

1137 the paper, properly credited and are the license and terms of use explicitly mentioned and

1138 properly respected?

1139 Answer: [Yes]

1140 Justification: [Yes] All third-party software used for experiments (NumPy, SciPy,

1141 levelsetpy) is properly cited in the references with version numbers where applicable.

1142 NumPy and SciPy are BSD-licensed; levelsetpy is cited per its original publication. No

1143 proprietary datasets were used.

1144 Guidelines:

- 1145 • The answer [N/A] means that the paper does not use existing assets.
- 1146 • The authors should cite the original paper that produced the code package or dataset.
- 1147 • The authors should state which version of the asset is used and, if possible, include a
- 1148 URL.
- 1149 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1150 • For scraped data from a particular source (e.g., website), the copyright and terms of
- 1151 service of that source should be provided.
- 1152 • If assets are released, the license, copyright information, and terms of use in the
- 1153 package should be provided. For popular datasets, paperswithcode.com/datasets
- 1154 has curated licenses for some datasets. Their licensing guide can help determine the
- 1155 license of a dataset.
- 1156 • For existing datasets that are re-packaged, both the original license and the license of
- 1157 the derived asset (if it has changed) should be provided.
- 1158 • If this information is not available online, the authors are encouraged to reach out to
- 1159 the asset's creators.

1160 **13. New assets**

1161 Question: Are new assets introduced in the paper well documented and is the documentation

1162 provided alongside the assets?

1163 Answer: [N/A]

1164 Justification: [N/A]

1165 Guidelines:

- 1166 • The answer [N/A] means that the paper does not release new assets.
- 1167 • Researchers should communicate the details of the dataset/code/model as part of their
- 1168 submissions via structured templates. This includes details about training, license,
- 1169 limitations, etc.
- 1170 • The paper should discuss whether and how consent was obtained from people whose
- 1171 asset is used.
- 1172 • At submission time, remember to anonymize your assets (if applicable). You can either
- 1173 create an anonymized URL or include an anonymized zip file.

1174 **14. Crowdsourcing and research with human subjects**

1175 Question: For crowdsourcing experiments and research with human subjects, does the paper
1176 include the full text of instructions given to participants and screenshots, if applicable, as
1177 well as details about compensation (if any)?

1178 Answer: [N/A]

1179 Justification: [N/A]

1180 Guidelines:

- 1181 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1182 with human subjects.
- 1183 • Including this information in the supplemental material is fine, but if the main contribu-
1184 tion of the paper involves human subjects, then as much detail as possible should be
1185 included in the main paper.
- 1186 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,
1187 or other labor should be paid at least the minimum wage in the country of the data
1188 collector.

1189 **15. Institutional review board (IRB) approvals or equivalent for research with human**
1190 **subjects**

1191 Question: Does the paper describe potential risks incurred by study participants, whether
1192 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
1193 approvals (or an equivalent approval/review based on the requirements of your country or
1194 institution) were obtained?

1195 Answer: [N/A]

1196 Justification: [N/A]

1197 Guidelines:

- 1198 • The answer [N/A] means that the paper does not involve crowdsourcing nor research
1199 with human subjects.
- 1200 • Depending on the country in which research is conducted, IRB approval (or equivalent)
1201 may be required for any human subjects research. If you obtained IRB approval, you
1202 should clearly state this in the paper.
- 1203 • We recognize that the procedures for this may vary significantly between institutions
1204 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
1205 guidelines for their institution.
- 1206 • For initial submissions, do not include any information that would break anonymity (if
1207 applicable), such as the institution conducting the review.

1208 **16. Declaration of LLM usage**

1209 Question: Does the paper describe the usage of LLMs if it is an important, original, or
1210 non-standard component of the core methods in this research? Note that if the LLM is used
1211 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
1212 scientific rigor, or originality of the research, declaration is not required.

1213 Answer: [N/A]

1214 Justification: [N/A]

1215 Guidelines:

- 1216 • The answer [N/A] means that the core method development in this research does not
1217 involve LLMs as any important, original, or non-standard components.
- 1218 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
1219 be described.