PROCEEDINGS A

royalsocietypublishing.org/journal/rspa

HJ-Gauss

Article submitted to journal

Subject Areas:
dynamical systems, game theory, control systems, runtime safety assurance

Keywords:
reachability analysis, differential games, dynamical systems, verification, rigorous approximation

Correspondence:
lekanmolu@scriptedonachip.com

Sampling-based HJ Reachability

Lekan Molu

Seattle, WA 98109

Viscous Hamilton–Jacobi (HJ) equations provide regularized value fields for backward reachability, but tensor-grid solvers have exponential storage in the state dimension. We apply a variable-coefficient Cole–Hopf transformation that maps the viscous HJ equation to a semilinear heat equation with coupled Gaussian initial-condition and nonlinear source integrals. A depth-first Monte Carlo Picard algorithm evaluates the complete transformed value-and-gradient update at individual states. On an analytic advection benchmark, increasing state dimension from 8 to 2048 at fixed sampling and Picard depth changes median point-query runtime from 0.0436 to 0.0760 seconds, while the counted live coordinate payload is linear in dimension. Increasing samples per update from 4 to 48 raises runtime from 0.0052 to 0.9355 seconds. The measurements therefore support grid-free point-query storage rather than an $O(Nn)$ end-to-end runtime claim. Separate deterministic transformed-PDE computations provide same-viscosity comparisons for rockets, Dubins vehicles, a double integrator, and a starling-inspired relative-state model. The $100{,}000$-state experiment queries one precomputed reduced field; it is not a coupled population HJ solve. The associated zero-sublevel geometries and topology counts are numerical observations rather than continuum safety certificates.

Animations

Rockets comparison | Dubins comparison | Double integrator | Starling-inspired geometry | Monte Carlo scaling | Code

© The Authors. Published by the Royal Society under the terms of the Creative Commons Attribution License http://creativecommons.org/licenses/by/4.0/, which permits unrestricted use, provided the original author and source are credited.

THE ROYAL SOCIETY PUBLISHING

1. Introduction

Modern complex systems — with their diverse interconnection of actuators, sensors, control systems, and the software modules that coordinate their various processes and events — require software verification tools capable of geometrically reasoning about their future behaviors. Given the continuous and/or hybrid systems characterization of their dynamics, simulation and testing, verification, observation, and further testing loop cycles are the standard norm for detecting unforeseen behavior ahead of deployment. Owing to their scale, however, analyzing system behavior based on closed-form analytic mathematical tools is often infeasible. This has stimulated efforts that complement traditional mathematical analysis tools from dynamical systems with emerging techniques in reachability [1], viability [2], or invariance [19] theories. These new tools allow exploring all possible scenarios where faults may arise for safety assurance; that is, ascertaining the freedom of a system from harm by generating evidence that a system, or any its components satisfy all specified requirements including functional and allocated baselines (i.e., *verification*) [12]. Examples of complex systems include smart manufacturing, ground [27] and air traffic management systems [39], flight control systems [31], or distributed systems.

Differential optimal control and game theories are useful formal verification tools for certifying the runtime safety assurance [19] of multiple agents interacting over a shared space in a *geometric* sense [21]. Both rely on a resolution of the Hamilton-Jacobi-Bellman (HJB) and/or the Hamilton-Jacobi-Isaacs (HJI) equation [11]. Being partial differential equations (PDEs), general high-dimensional HJB/HJI equations rarely admit closed-form solutions. Thus, approximation techniques have become the default mechanism for formal systems verification in these settings. Reachability theory [32, 39] is one such approximation technique. It entails computing the essentially non-oscillating (ENO) [38] Lax-Friedrichs [10] approximations to the HJ PDE via levelset methods [33, 34, 36]. Propagating the fronts of a conservation law on a mesh, the inviscid Euler equations' solution reduces to the derivative of the HJ problem's solution [14]; such that integrating this solution with total variation diminishing Runge-Kutta (TVD-RK) [37] produces generalized *viscosity solutions* to the original HJ problem.

HJ equation systems are employed within reachability analysis to evaluate a dynamical system's behavioral (or trajectory) evolution throughout a state space — for example, given its initial point sets (states). The problem is typically setup so that the initial set must satisfy *input or state constraints* in a *least restrictive sense* and under a *worst-possible disturbance* [28]. Reachability is thus a useful geometric reasoning tool for determining the states that can be reached from an initial condition in a finite number of steps [32, 34] . *Reachability* can be analyzed in a (i) *forward* sense, where system trajectories are examined to determine if they enter certain states from an *initial set*; (ii) *backward* sense, where system trajectories are examined to determine if they enter certain *target sets*; (iii) *reach set* sense, in which they are examined to see if states reach a *set* at a *particular time*; or (iv) *reach tube* sense, in which they are examined to determine that they reach a set *during a time interval*.

This article studies the Cauchy equations associated with terminal reachability and backward reachable tubes. A terminal reachable set concerns target entry at a fixed horizon; a tube permits entry during an interval. Reach–avoid problems additionally constrain the trajectory before target entry and require an obstacle or stopping formulation. Classical tensor-grid discretizations require storage that grows exponentially with the number of state coordinates at fixed resolution [4]; avoiding that storage does not alone establish numerical accuracy or favorable complexity.

Consider the evolution-form of the Cauchy-type HJ PDE

$$\boldsymbol{v}_t + \boldsymbol{H}(t; \boldsymbol{x}, D\boldsymbol{v}) = 0 \text{ in } \Omega \times (0, T], \quad \boldsymbol{v}(0; \boldsymbol{x}) = \boldsymbol{g}(0; \boldsymbol{x}) \text{ on } \Omega \times \{t = 0\} \tag{HJ-IVP}$$

where $\boldsymbol{x} \in \Omega \subseteq \mathbb{R}^n$, the Hamiltonian $\boldsymbol{H}$ is continuous, and $D\boldsymbol{v}$ is the spatial gradient. The initial condition $\boldsymbol{g}$ is bounded and uniformly continuous. Crossing characteristics can destroy classical differentiability even when the Hamiltonian and initial condition are smooth [8]. Viscosity solutions provide generalized solutions. We study a fixed, positive-viscosity regularization before considering its approximation error.

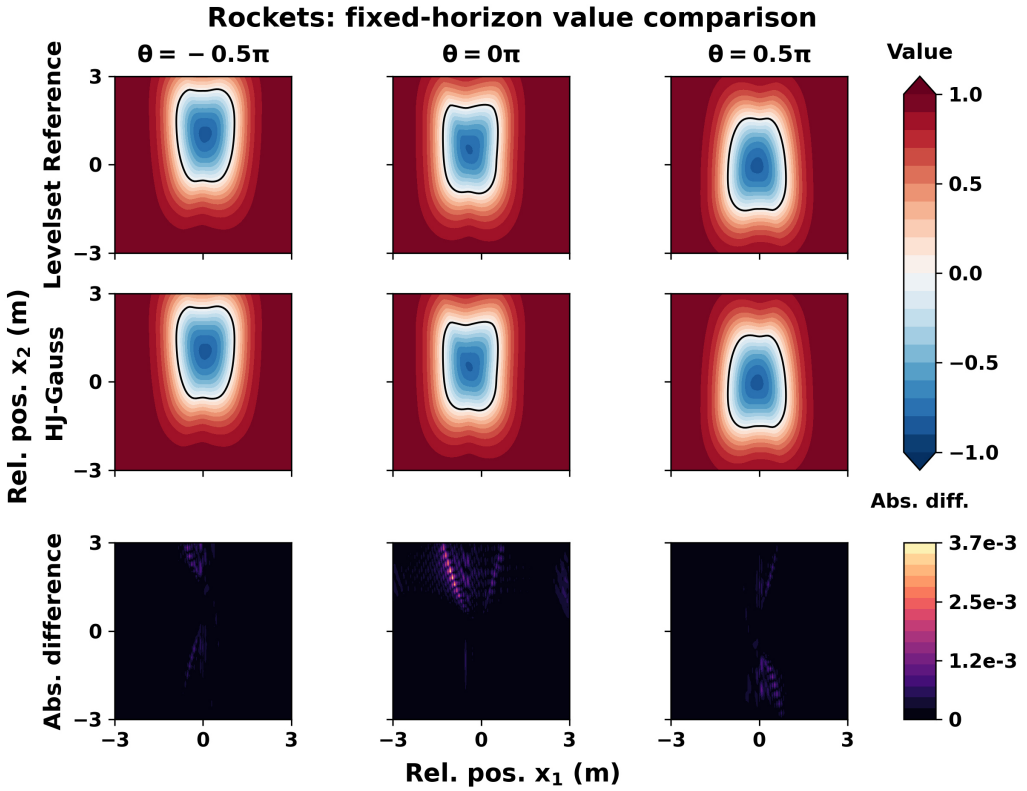


Figure 1: Normalized rockets benchmark at elapsed horizon $T = 0.5$, with $\delta = 0.08$ with fixed-horizon values. Rows show the `Levelset` Reference [33, 34], `HJ-Gauss`, and their absolute difference, respectively. Columns show the slices of the BRTs at $\theta = -\pi/2$, 0, and $\pi/2$. Both solvers use the same initial condition and viscosity parameter. Black contours mark the zero level set.

In this paper, we introduce `HJ-Gauss`, a generalized Cole–Hopf transformation of the viscous form of (HJ-IVP) followed by a quasilinearized fixed-point iteration. The transformation yields a semilinear heat equation under positive logarithmic inversion while retaining the Hamiltonian and value derivatives. We evaluate its coupled value-and-gradient mild updates with a depth-first nonlinear Monte Carlo Picard algorithm. Our construction uses the semilinear parabolic framework of [18, 20]; it does not imply a new dimension-independent complexity result. In essence, our results concern (i) the transformation identity and its differentiated mild representation; (ii) a finite-time source iteration that contracts on a positive invariant set under explicit value-and-gradient Lipschitz assumptions; (iii) a complete nested Monte Carlo realization of finite Picard levels; and (iv) measured point-query storage and runtime scaling with explicit claim limits.

The rest of this paper is structured as follows: §2 discusses background material and presents our theoretical machinery, §3 provides numerical validation on moderate to high-dimensional problems, and we conclude the paper in §4. Comprehensive proofs to our results are presented in Appendix A, convergence, residual certification and robustness analyses are given in Appendix B.

2. Background and Quasilinearization Framework

Backward reachability associates a differential game with a Hamilton–Jacobi–Isaacs equation [32]. Let $\mathcal{U} \subset \mathbb{R}^m$ and $\mathcal{W} \subset \mathbb{R}^p$ be compact control sets. At physical time $s \in [0, T]$, let $\overline{\mathcal{U}}(s)$ and $\overline{\mathcal{W}}(s)$ contain the measurable control signals on $[s, T]$. Let $\mathcal{B}(s)$ contain the nonanticipative maps

$\beta : \overline{\mathcal{U}}(s) \to \overline{\mathcal{W}}(s)$. The trajectory ξ starts in state $\boldsymbol{x}$ and satisfies

$$\dot{\xi}(r) = f(r, \xi(r), \boldsymbol{u}(r), \beta[\boldsymbol{u}](r)), \qquad \xi(s) = \boldsymbol{x}, \quad s \le r \le T, \tag{2.1}$$

where f is bounded, continuous and uniformly Lipschitz in $\boldsymbol{x}$. For bounded uniformly continuous running cost l and terminal cost $\boldsymbol{g}$, the lower value of the differential game is

$$\boldsymbol{v}_{\text{phys}}(s, \boldsymbol{x}) = \inf_{\beta \in \mathcal{B}(s)} \sup_{\boldsymbol{u} \in \overline{\mathcal{U}}(s)} \underbrace{\left\{ \int_s^T l(r, \xi(r), \boldsymbol{u}(r), \beta[\boldsymbol{u}](r))\, dr + \boldsymbol{g}(\xi(T)) \right\}}_{\Phi(s,\boldsymbol{x}) := \text{Payoff}}. \tag{2.2}$$

The target is represented by the zero sublevel set of $\boldsymbol{g}$. We assign $\boldsymbol{u}$ to the evader, which maximizes the payoff to avoid this set, and $\boldsymbol{w}$ to the pursuer, which minimizes the payoff to force target entry. The Hamiltonian is

$$\boldsymbol{H}_{\text{phys}}(s, \boldsymbol{x}, \boldsymbol{p}) = \max_{\boldsymbol{u} \in \mathcal{U}} \min_{\boldsymbol{w} \in \mathcal{W}} \{\boldsymbol{p} \cdot f(s, \boldsymbol{x}, \boldsymbol{u}, \boldsymbol{w}) + l(s, \boldsymbol{x}, \boldsymbol{u}, \boldsymbol{w})\}. \tag{2.3}$$

For pure target-reach problems, the running cost is set to zero and the target is $\boldsymbol{g}$. The elapsed time $t = T - s$ converts the terminal condition into an initial condition so that if we write $\boldsymbol{v}(t, \boldsymbol{x}) = \boldsymbol{v}_{\text{phys}}(T - t, \boldsymbol{x})$ and define $\boldsymbol{H}(t, \boldsymbol{x}, \boldsymbol{p}) = -\boldsymbol{H}_{\text{phys}}(T - t, \boldsymbol{x}, \boldsymbol{p})$, the elapsed equation is $\boldsymbol{v}_t + \boldsymbol{H}(t, \boldsymbol{x}, D\boldsymbol{v}) = 0$. Fixing $\delta > 0$ gives the viscous parabolic regularization

$$\boldsymbol{v}_t^\delta + \boldsymbol{H}(t, \boldsymbol{x}, D\boldsymbol{v}^\delta) = \frac{\delta}{2} \Delta \boldsymbol{v}^\delta, \qquad \boldsymbol{v}^\delta(0, \boldsymbol{x}) = \boldsymbol{g}(\boldsymbol{x}). \tag{2.4}$$

Equation (2.4) supplies approximation results under the Hamiltonian and initial condition [8, 9]. The transformation below concerns classical solutions of this regularized equation.

(a) HJ Linearization

Write $\boldsymbol{v} = \boldsymbol{v}^\delta$ within this subsection and choose a prescribed coefficient $\boldsymbol{c} \in C_b^{1,2}([0, T] \times \mathbb{R}^n)$ with $0 < \boldsymbol{c}_- \le \boldsymbol{c} \le \boldsymbol{c}_+$. Define the transformed value and its initial condition by

$$\boldsymbol{\omega} = e^{-\boldsymbol{c}\boldsymbol{v}}, \qquad \boldsymbol{\omega}_0(\boldsymbol{x}) = e^{-\boldsymbol{c}(0,\boldsymbol{x})\boldsymbol{g}(\boldsymbol{x})}. \tag{2.5}$$

For a scalar $z > 0$ and vector $q \in \mathbb{R}^n$, introduce

$$V(t, \boldsymbol{x}, z) = -\frac{\log z}{\boldsymbol{c}(t, \boldsymbol{x})}, \qquad P(t, \boldsymbol{x}, z, q) = -\frac{q}{\boldsymbol{c}z} + \frac{\log z}{\boldsymbol{c}^2} D\boldsymbol{c}, \tag{2.6}$$

$$\boldsymbol{R}_{\text{alg}}[\boldsymbol{c}, V, P] = \boldsymbol{c} \left(\boldsymbol{H}(t, \boldsymbol{x}, P) - \frac{\delta \boldsymbol{c}}{2} |P|^2 \right), \tag{2.7}$$

$$\boldsymbol{R}_{\text{der}}[\boldsymbol{c}, V, P] = -\boldsymbol{c}_t V + \frac{\delta}{2} (\Delta \boldsymbol{c}) V - \frac{\delta}{2} |D\boldsymbol{c}|^2 V^2 + \delta (1 - \boldsymbol{c}V) D\boldsymbol{c} \cdot P, \tag{2.8}$$

$$F(t, \boldsymbol{x}, z, q) = z \{\boldsymbol{R}_{\text{alg}}[\boldsymbol{c}, V, P] + \boldsymbol{R}_{\text{der}}[\boldsymbol{c}, V, P]\}. \tag{2.9}$$

Then, the transformed equation is

$$\boldsymbol{\omega}_t - \frac{\delta}{2} \Delta \boldsymbol{\omega} = F(t, \boldsymbol{x}, \boldsymbol{\omega}, D\boldsymbol{\omega}), \qquad \boldsymbol{\omega}(0) = \boldsymbol{\omega}_0. \tag{2.10}$$

For positive solutions, logarithmic inversion recovers the original viscous equation — reducing to a heat equation when the sum in (2.9) vanishes along the solution. Let S_t denote the heat semigroup with covariance $\delta t I_n$. The transformed equation admits the *mild representation*

$$\boldsymbol{\omega}(t) = S_t \boldsymbol{\omega}_0 + \int_0^t S_{t-s} F(s, \cdot, \boldsymbol{\omega}(s), D\boldsymbol{\omega}(s))\, ds, \tag{2.11}$$

which couples values and gradients. Each mild update substitutes the preceding iterate into both source arguments.

Algorithm 1 Depth-first nonlinear Monte Carlo Picard point evaluation

1: Choose a query $(t, \boldsymbol{x})$, level $K \geq 0$, viscosity $\delta > 0$, positive coefficient $\boldsymbol{c}$, even sample count N, and random seed.
2: Estimate $(S_t, DS_t)\boldsymbol{\omega}_0$ using $N/2$ antithetic Gaussian pairs.
3: **if** $K = 0$ **then**
4: Return the heat value and score gradient.
5: **end if**
6: Draw $N/2$ independent $U \sim \mathcal{U}(0,1)$ and $Z \sim \mathcal{N}(0, I_n)$; set $S = t(1-U^2)$.
7: **for** each antithetic displacement $\pm\sqrt{\delta(t-S)}Z$ **do**
8: Recursively evaluate (z_{K-1}, Dz_{K-1}) at both displaced source states.
9: Evaluate the complete source $F(S, \cdot, z_{K-1}, Dz_{K-1})$ from (2.9).
10: **end for**
11: Add the importance-weighted Duhamel value and score-gradient estimates to the heat pair.
12: Reject a nonfinite or nonpositive transformed estimate; otherwise recover $(\boldsymbol{v}_K, D\boldsymbol{v}_K)$ using (2.13).

For $f_k(s, \boldsymbol{x}) = F(s, \boldsymbol{x}, z_k, Dz_k)$, the value and gradient update as an ordered pair:

$$(z_{k+1}, Dz_{k+1})(t) = (S_t, DS_t)\boldsymbol{\omega}_0 + \int_0^t (S_{t-s}, DS_{t-s}) f_k(s)\, ds. \tag{2.12}$$

$$\boldsymbol{v}_{k+1}(t, \boldsymbol{x}) = -\frac{\log z_{k+1}(t, \boldsymbol{x})}{\boldsymbol{c}(t, \boldsymbol{x})}, \qquad D\boldsymbol{v}_{k+1} = -\frac{Dz_{k+1}}{\boldsymbol{c} z_{k+1}} - \frac{\boldsymbol{v}_{k+1}}{\boldsymbol{c}} D\boldsymbol{c}. \tag{2.13}$$

Appendix A derives the Gaussian forms of the source integral and coefficients. Algorithm 1 implements the coupled value-and-gradient mild update at one requested state. It evaluates the nonlinear source at recursively sampled intermediate states rather than assuming those values are already known.

(b) Relations to Backward Reachable Sets/Tubes

It is often desirable to consider conditions under which system trajectories may enter a *target set* within a time interval $(0, T]$. This could be desirable in goal-regions of the phase space (safe sets) or undesirable configurations (unsafe sets). Suppose that the goal of $\boldsymbol{P}$ is to drive the $\boldsymbol{E}$ into a target region $\mathcal{L}_0(\tau)$ within $\tau \leq T$ time steps of game play. The target set and the corresponding *robustly controlled backward reachable tube* for $\tau \in [-T, 0]$ in the closure of the open set [31] are

$$\begin{aligned} \mathcal{L}_0(T) &= \{\boldsymbol{x} \in \mathbb{R}^n \,|\, \boldsymbol{v}(0; \boldsymbol{x}) \leq 0\}, \\ \mathcal{L}([\tau, 0], \mathcal{L}_0) &= \{\boldsymbol{x} \in \Omega \,|\, \exists\, \beta \in \bar{\mathcal{V}}(t)\, \forall\, \boldsymbol{u} \in \mathcal{U}(t), \exists\, \tau \in [-T, 0], \boldsymbol{\xi}(\bar{t}) \in \mathcal{L}_0\}. \end{aligned} \tag{2.14a}$$

Read $\mathcal{L}([\tau, 0], \mathcal{L}_0)$: The set of states from which the strategies of $\boldsymbol{P}$ and for all controls of $\boldsymbol{E}$ imply that we *reach and remain in the target set* in the interval $[-T, 0]$. Following Lemma 2 of [32], the states in the reachable set admit the following properties w.r.t the value function $\boldsymbol{v}$:

$$\boldsymbol{x}(t) \in \mathcal{L}(\cdot) \implies \boldsymbol{v}(\boldsymbol{x}, t) \leq 0, \boldsymbol{v}(\boldsymbol{x}, t) \leq 0 \implies \boldsymbol{x}(t) \in \mathcal{L}(\cdot). \tag{2.15}$$

Player $\boldsymbol{P}$ is minimizing (the game's termination time c.f. $\boldsymbol{\ell}(0, \boldsymbol{x})$), seeking to drive system trajectories into the unsafe set; and $\boldsymbol{E}$ is maximizing (the game's termination time) i.e., is seeking to avoid the unsafe set-[1]. A differential game's value function admits a "viscosity" (generalized) solution [8, 26] for the associated HJ-Isaacs (HJI) PDE which are *locally Lipschitz* in $\Omega \times [-T, 0]$, and with at most first-order partial derivatives in the Hamiltonian. Equation (2.2) admits the viscosity

[1]For the goal-satisfaction (or *liveness*) problem setups, the strategies are reversed and the backward reachable tube are the states from which the evader $\boldsymbol{E}$ can successfully reach the target set despite worst-case efforts of the pursuer $\boldsymbol{P}$.

solution

$$\frac{\partial \boldsymbol{v}}{\partial t}(\boldsymbol{x},t)+\min\{0,\boldsymbol{H}(t;\boldsymbol{x},p)\}=0,\;\; V(\boldsymbol{x},0)=g(\boldsymbol{x}) \tag{2.16}$$

where the vector field $p\,(\equiv \boldsymbol{v}_{\boldsymbol{x}})$ is known in terms of the game's terminal conditions so that the overall game is akin to a two-point boundary-value problem; and the Hamiltonian $\boldsymbol{H}(t;\boldsymbol{x},\boldsymbol{u},\boldsymbol{v},\boldsymbol{v}_{\boldsymbol{x}})$ is defined as

$$\boldsymbol{H}(t;\boldsymbol{x},\boldsymbol{u},\boldsymbol{v},\boldsymbol{v}_{\boldsymbol{x}})=\max_{u\in\mathcal{U}}\min_{v\in\mathcal{V}}\langle f(t;\boldsymbol{x},\boldsymbol{u},\boldsymbol{v}),\boldsymbol{v}_{\boldsymbol{x}}\rangle. \tag{2.17}$$

For BRTs, the agent must avoid the unsafe region at all times. Thus, the *robustly controlled backward reach-avoid tube* (RCBRAT) i.e.,

$$\min\{\boldsymbol{v}_t(t;\boldsymbol{x})+\boldsymbol{H}\,(t;\boldsymbol{x},D\boldsymbol{v})\,,\boldsymbol{g}(t;\boldsymbol{x})-\boldsymbol{\ell}(t;\boldsymbol{x})\}\leq 0,\quad \boldsymbol{v}(0;\boldsymbol{x})=\boldsymbol{g}(0;\boldsymbol{x}), \tag{HJI-RCBRAT}$$

for a signed-distance (from target) cost $\ell(t;x)$. Equation (2.2) only allows the game to determine that a trajectory belongs in the target set at exactly time zero. That is, the evader can chase a trajectory into the target set and escape it before reaching the final time, 0. The min operation in (2.16) lets the pursuer "freeze" trajectories that may be under the evader's control when the evader tries to evolve such trajectories outside of the target set. For more details on the construction of this P.D.E., see [32]. Introducing the viscosity parameter, $\delta>0$, the viscous variational HJ PDE is

$$\partial_t \boldsymbol{v}^\delta+\min\{0,\boldsymbol{H}(\boldsymbol{x},D\boldsymbol{v}^\delta)\}=\frac{\delta}{2}\Delta\boldsymbol{v}^\delta,\qquad \boldsymbol{v}^\delta(0)=\boldsymbol{g}. \tag{2.18}$$

Considering (2.18), we have the diffusion coefficient as $\delta/2$, so that $\boldsymbol{c}$ satisfies the bounds in §2(a). Suppose we write $\boldsymbol{v}^\delta:=\boldsymbol{v}$ for ease of readability. We now derive the Gaussian representation of the BRT/BRAT.

Lemma 2.1 (Transformation Identity)**.** *Suppose that $\boldsymbol{v}(t;\boldsymbol{x}),\boldsymbol{c}(t;\boldsymbol{x})\in C^{1,2}$ and set $\boldsymbol{\omega}=e^{-\boldsymbol{c}\boldsymbol{v}}$. Define the residual $r_v=\boldsymbol{v}_t+\boldsymbol{H}(t,\boldsymbol{x},D\boldsymbol{v})-\delta\Delta\boldsymbol{v}/2$. Then*

$$\boldsymbol{\omega}_t-\frac{\delta}{2}\Delta\boldsymbol{\omega}=\boldsymbol{\omega}\{\boldsymbol{R}_{\rm alg}[\boldsymbol{c},\boldsymbol{v},D\boldsymbol{v}]+\boldsymbol{R}_{\rm der}[\boldsymbol{c},\boldsymbol{v},D\boldsymbol{v}]-\boldsymbol{c}r_v\}. \tag{2.19}$$

Proof. Set $\phi=\boldsymbol{c}\boldsymbol{v}$. Differentiation gives $\boldsymbol{\omega}_t=-\boldsymbol{\omega}\phi_t$, $D\boldsymbol{\omega}=-\boldsymbol{\omega}D\phi$, and $\Delta\boldsymbol{\omega}=\boldsymbol{\omega}(|D\phi|^2-\Delta\phi)$. Expand $D\phi=\boldsymbol{c}D\boldsymbol{v}+\boldsymbol{v}D\boldsymbol{c}$ and $\Delta\phi=\boldsymbol{c}\Delta\boldsymbol{v}+2D\boldsymbol{c}\cdot D\boldsymbol{v}+\boldsymbol{v}\Delta\boldsymbol{c}$. Substitution cancels the two terms containing $\boldsymbol{c}\Delta\boldsymbol{v}$. The remaining terms give (2.19). Since $\boldsymbol{c}\boldsymbol{\omega}>0$, the original residual vanishes if and only if the transformed residual vanishes. □

Corollary 2.1 (Homogeneous Heat Flow Reduction)**.** *Considering (2.19), homogeneous heat flow holds if and only if $\boldsymbol{R}_{\rm alg}+\boldsymbol{R}_{\rm der}=0$. A sufficient condition is a positive constant $\boldsymbol{c}_0>0$ and $\boldsymbol{H}(t,\boldsymbol{x},D\boldsymbol{v}):=\delta\boldsymbol{c}_0|D\boldsymbol{v}|^2/2$. The Hamiltonian $\boldsymbol{H}(p)=|p|^2/2$ satisfies this condition whenever $\boldsymbol{c}_0=1/\delta$.*

Proof. Apply Lemma 2.1 with $r_v=0$. A constant $\boldsymbol{c}$ eliminates the derivative residual, and the quadratic relation eliminates the algebraic residual. □

Corollary 2.2 (Source representation and value recovery)**.** *Suppose $\boldsymbol{\omega}_0$ is bounded and uniformly continuous, and the source $f(s,\boldsymbol{x})=F(s,\boldsymbol{x},\boldsymbol{\omega}(s,\boldsymbol{x}),D\boldsymbol{\omega}(s,\boldsymbol{x}))$ is bounded and continuous. Every bounded classical solution satisfies*

$$\boldsymbol{\omega}(t,\boldsymbol{x})=\mathbb{E}[\boldsymbol{\omega}_0(\boldsymbol{x}+\sqrt{\delta t}\,Z)]+\int_0^t\mathbb{E}[f(s,\boldsymbol{x}+\sqrt{\delta(t-s)}\,Z)]\,ds, \tag{2.20a}$$

$$\boldsymbol{v}(t,\boldsymbol{x})=-\frac{1}{\boldsymbol{c}(t,\boldsymbol{x})}\log\boldsymbol{\omega}(t,\boldsymbol{x}),\qquad Z\sim\mathcal{N}(0,I_n). \tag{2.20b}$$

Proof. Variation of constants for the heat equation gives the mild representation. The whole-space heat kernel has unit mass and equals the indicated Gaussian density. Invert the transformation at the evaluation point. □

The source contains the unknown solution. If $|\boldsymbol{v}| \le G$, then $e^{-\boldsymbol{c}_+ G} \le \boldsymbol{\omega} \le e^{\boldsymbol{c}_+ G}$.

Proposition 2.1 (Gradient for `HJ-Gauss`)**.** *Under the preceding bounded-source hypotheses, for $t > 0$ the gradient satisfies*

$$D\boldsymbol{\omega}(t,\boldsymbol{x}) = \frac{1}{\sqrt{\delta t}}\mathbb{E}[Z\boldsymbol{\omega}_0(\boldsymbol{x}+\sqrt{\delta t}\,Z)] + \int_0^t \frac{\mathbb{E}[Zf(s,\boldsymbol{x}+\sqrt{\delta(t-s)}\,Z)]}{\sqrt{\delta(t-s)}}\,ds, \tag{2.21a}$$

$$D\boldsymbol{v}(t,\boldsymbol{x}) = -\frac{D\boldsymbol{\omega}(t,\boldsymbol{x})}{\boldsymbol{c}(t,\boldsymbol{x})\boldsymbol{\omega}(t,\boldsymbol{x})} - \frac{\boldsymbol{v}(t,\boldsymbol{x})}{\boldsymbol{c}(t,\boldsymbol{x})}D\boldsymbol{c}(t,\boldsymbol{x}). \tag{2.21b}$$

Proof. Differentiate the Gaussian kernel with respect to its center. Its derivative has integrable L^1 norm bounded by $\sqrt{n/(\delta(t-s))}$. This bound justifies differentiation of the source integral. The product rule applied to $\boldsymbol{v} = -\boldsymbol{c}^{-1}\log\boldsymbol{\omega}$ gives the second formula. □

(c) Weighted Importance Sampling of Reachable Sets

One can apply importance sampling to transform the Gaussian representation while preserving the integral through a density ratio. For a displacement $\vartheta \in \mathbb{R}^n$, variance $\sigma^2 > 0$, and sample $\boldsymbol{y} \sim \mathcal{N}(\boldsymbol{x}+\vartheta, \sigma^2 I_n)$, the "importance weight" is

$$w_\vartheta(\boldsymbol{y}) = \exp\left(\frac{|\vartheta|^2 - 2\vartheta\cdot(\boldsymbol{y}-\boldsymbol{x})}{2\sigma^2}\right). \tag{2.22}$$

Let $\sigma^2 = \delta t$ be the initial-condition and $\sigma^2 = \delta(t-s)$ represent the variance. Suppose that the evaluation count M denotes queried states and sample count N denotes independent random draws at each query and update. Conditional on the preceding iterate, suppose $a_0 \le \boldsymbol{\omega}_0 \le b_0$ and $|f_k| \le B_F$. Uniform source-time sampling gives the unbiased transformed-value sample

$$A = \boldsymbol{\omega}_0(\boldsymbol{x}+\sqrt{\delta t}Z) + tf_k(tU, \boldsymbol{x}+\sqrt{\delta t(1-U)}Z), \qquad U\sim\mathcal{U}(0,1), \quad Z\sim\mathcal{N}(0,I_n), \tag{2.23}$$

with independent U and Z. Each sample lies in $[a_0 - tB_F, b_0 + tB_F]$, of width $D = b_0 - a_0 + 2tB_F$.

Theorem 2.1 (Finite-query concentration for a fixed source)**.** *Fix positive integers M, K, N, representing the number of queries, the maximum number of updates, and the samples in each batch, and let $\eta \in (0,1)$. Assume the preceding source bounds hold conditionally at every update with common sample-width bound D_*. If each batch contains N conditionally independent samples with the stated common conditional law, then all transformed-value sampling errors are at most*

$$h_N = D_*\sqrt{\frac{\log(2MK/\eta)}{2N}} \tag{2.24}$$

with probability at least $1-\eta$.

Proof in Appendix B.

Corollary 2.3 (Sample count and logarithmic recovery)**.** *For transformed-value tolerance $\varepsilon_\omega > 0$, it suffices to choose*

$$N \ge \frac{D_*^2}{2\varepsilon_\omega^2}\log\frac{2MK}{\eta}. \tag{2.25}$$

If the target mild update is at least $a > 0$ and $h_N < a$, then its sampled logarithmic inverse differs by at most $h_N/[\boldsymbol{c}_-(a-h_N)]$ on the same event.

The proof is given in Appendix B.

Theorem 2.2 (Conditional finite-time convergence of the `HJ-Gauss` source map)**.** *Let* $X_\tau = C([0,\tau];\mathrm{BUC}^1)$ *with the sum of the uniform value and gradient norms. Suppose a nonempty closed set* $\mathcal{A} \subset X_\tau$ *is invariant under the mild map, its functions satisfy* $a \le z \le b$ *and* $|Dz| \le Q$ *with* $a > 0$*, and the bounded source is uniformly Lipschitz in* (z, Dz) *with constant* L*. If*

$$\theta = L\left(\tau + 2\sqrt{n\tau/\delta}\right) < 1, \tag{2.26}$$

then the `HJ-Gauss` *source map has a unique fixed point in* $\mathcal{A}$ *and its iterates starting in* $\mathcal{A}$ *converge geometrically in* X_τ*. If that fixed point is classical, its inverse solves the viscous Hamilton–Jacobi equation.*

The proof is given with Theorem B.1 in Appendix B, together with an invariant-ball criterion and a posteriori stopping bound.

(d) Notes on Monte Carlo and Viscosity Approximation Errors

For a bounded smooth approximation U with bounded residual

$$r_U = U_t + \boldsymbol{H}(t, \boldsymbol{x}, DU) - \frac{\delta}{2}\Delta U, \tag{2.27}$$

comparison gives a value-error bound by the initial-condition error plus $\int_0^t \|r_U(s)\|_\infty ds$. If an independent vanishing-viscosity theorem also gives a bound $C_{\mathrm{vis}}\sqrt{\delta}$, the combined error bound is

$$E(t) = \|U(0) - \boldsymbol{g}\|_\infty + \int_0^t \|r_U(s)\|_\infty ds + C_{\mathrm{vis}}\sqrt{\delta}. \tag{2.28}$$

The transformed residual satisfies $d = -\boldsymbol{c}\exp(-\boldsymbol{c}U)r_U$ (see Appendix B).

Corollary 2.4 (Conditional sign certification)**.** *Suppose* $|U - \boldsymbol{v}^0| \le E$ *on a specified set, with* $\boldsymbol{v}^0 \le 0$ *denoting the target-reaching class. Then* $U > E$ *implies* $\boldsymbol{v}^0 > 0$*, and* $U < -E$ *implies* $\boldsymbol{v}^0 < 0$*. The remaining states form the abstention band* $|U| \le E$*.*

The proof is given in Appendix B.

3. Sampling Scaling and Numerical Validation

The experiments separate three questions: point-query scaling of Algorithm 1, discretization accuracy of deterministic transformed-PDE solvers, and repeated evaluation of a precomputed relative-state field. Deterministic comparisons share the viscosity, Hamiltonian, bounded initial condition, and horizon, separating their numerical discrepancy from viscosity bias. The model descriptions connect the game experiments to their dynamical and geometric interpretation. All recomputed fields use the elapsed-time convention of Section 2.

(a) Nonlinear Monte Carlo Picard scaling

The sampling benchmark uses the bounded initial value $g(\boldsymbol{x}) = 0.1n^{-1}\sum_{j=1}^n \cos x_j$, viscosity $\delta = 0.4$, and linear-advection Hamiltonian $\boldsymbol{H}(p) = a \cdot p$. The velocity components range from $0.05/\sqrt{n}$ to $0.15/\sqrt{n}$. The viscous HJ solution is available analytically, while the Cole–Hopf source evaluated by Algorithm 1 remains nonlinear. We query the zero state at horizon $t = 0.15$. Each reported runtime and memory value is the median of 3 fresh-process runs.

At Picard level 2 and 12 samples per update, increasing dimension from 8 to 2048 changes median runtime from 0.0436 to 0.0760 seconds. The fitted log–log runtime slope is 0.085, and

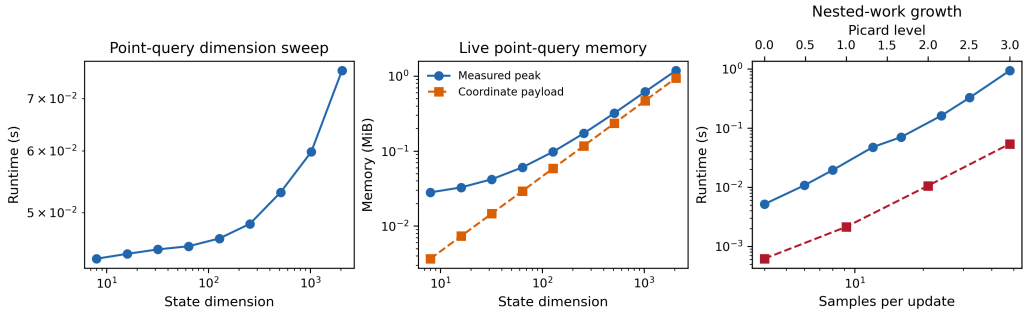


Figure 2: Scaling of one nonlinear Monte Carlo Picard point query. Left: runtime at fixed sample count and Picard level. Center: Python-traced peak memory and the counted depth-linear coordinate payload. Right: runtime growth with samples per update and Picard level. Markers are medians of three fresh-process runs; connecting lines guide the eye.

the maximum absolute value error in this sweep is 5.53×10^{-4}. The measured Python-traced-memory slope is 0.694; fixed interpreter costs dominate the smaller cases, while the directly counted coordinate payload is linear in dimension.

The sample and depth sweeps expose the nested work. At dimension eight and Picard level two, increasing samples per update from 4 to 48 increases median runtime from 0.0052 to 0.9355 seconds, with fitted log–log slope 2.05. With six samples, increasing the Picard level from 0 to 3 increases runtime from 0.00062 to 0.05399 seconds. Figure 2 therefore supports grid-free point-query storage for fixed Picard depth. It does not support an $O(Nn)$ end-to-end runtime claim, favorable scaling in Picard depth, or continuum field reconstruction.

(b) Starlings' Murmuration as a Scalability Example

This subsection studies starling-inspired numerical geometry and repeated queries of a reduced value field. The sampling scalability of Algorithm 1 is measured separately in Section 3(a). Through empirical [3, 6, 7, 17, 41] and theoretical findings [23], evidence abounds that natural species exhibiting collective behavior have their convergence and group cohesion characterized by simple topological interaction rules for formation motions such as collision avoidance, anisotropic preservation of density and structure. European Starlings for example exhibit flock splitting, vacuole nucleation, cordon, and flash expansion isotropically [15] through local anisotropic synchrony, which taken together among possibly hundreds of thousands of local interactions, keep these animals whirling, swooping, and flying in isotropic formations [3]. Individual agents aggregate into substructures within the overall system where motion is synergized via simple local topological interactions that preserve stable global heading and cohesion [23].

When an individual within a subgroup of Starlings murmuration senses danger (e.g. an attack from a Peregrine Falcon), immediate *nearest neighbors* change course. This is propagated across the entire group within the fraction of a second [3], resulting in elegant formations. This motivates our large-scale reachable safety framework. Overall motion is characterized by dynamically evolving interfaces [3, 7] that separate the respective multiple regions [24, 25, 43].

Figure 4 shows the value functions of three murmuration events. Cordon formation, flash expansion, and vacuole nucleation each produce a distinct zero-sublevel geometry, and `HJ-Gauss` reproduces the levelset reference [34] in each case. The absolute difference is largest near the zero contour, which is where the sign of the value matters for the safe set.

(i) Notations and abstractions

Let $\Omega \subseteq \mathbb{R}^n$ denote the state domain and $v : [0, T] \times \Omega \to \mathbb{R}$ a scalar value field. A subgroup occupies a region $\mathcal{S}_i \subseteq \Omega$, and the collection occupies $\mathcal{C} = \bigcup_i \mathcal{S}_i$. We assume regions are separable

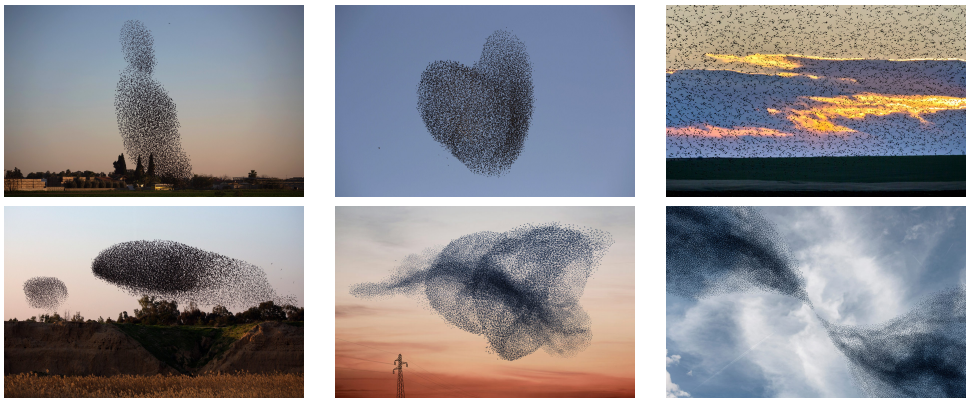


Figure 3: Photographs illustrate collective behaviors of European starlings (*Sturnus vulgaris*). Reading across each row. (i) A flock of starlings rises into the air, in a dense structure (Reuters/Amir Cohen). (ii) Starlings migrating over an Israeli village (AP Photo/Oded Balilty). (iii) Starlings feeding on laid seeds on the ground in Romania. (iv) Two flocks of migrating starlings (Menahem Kahana/AFP/Getty Images). (v) A concentric conical formation of starlings (Courtesy of The Gathering Site.). (vi) Splitting and joining of a flock of starlings.

and disjoint so that $\mathcal{S}_i \cap \mathcal{S}_j = \emptyset$ for $i \neq j$. An individual bird has position and heading state $\boldsymbol{x}^{(i)} = (x^{(i)}, y^{(i)}, z^{(i)}, \theta^{(i)})$. with neighbors $\mathcal{N}_i(t)$ at time t — the set of all birds that lie within a predefined radius, r_i. A *flock*, F, of birds labeled $\{1, 2, \cdots, n_a\}$ within a murmuration is a collection of birds $(\mathcal{X}, T)$ such that all birds within the flock interact with their nearest neighbors in a topological sense. Every bird within a flock has similar dynamics to that of its neighbor(s) and they travel at the same linear speed, v; the angular headings, w, however, may be different. Each bird's continuous-time dynamics, $\dot{\boldsymbol{x}}^{(i)}(t)$, evolves with constant forward speed $v > 0$, vertical control $u_z^{(i)}$, and steering rate $w^{(i)}$ as

$$\dot{\boldsymbol{x}}^{(i)} = \left[v\cos\theta^{(i)}, v\sin\theta^{(i)}, u_z^{(i)}, \overline{w}^{(i)}\right]^\top, \qquad \overline{w}^{(i)} = \frac{w^{(i)} + \sum_{j\in\mathcal{N}_i(t)} w^{(j)}}{1 + |\mathcal{N}_i(t)|}. \tag{3.1}$$

(ii) Murmuration as evolving level sets

Each agent within a flock interacts with a fixed number of neighbors, n_c, within a fixed topological range, r_c — reinforcing *group cohesion* [3]. We are interested in *robust group cohesion*; therefore we let a pursuer, $\boldsymbol{P}$, with a worst-possible disturbance attack the flock, and flocks constitute an evading player, $\boldsymbol{E}$ — so that each bird in a flock is always in relative coordinates to $\boldsymbol{P}^j$. The zero sublevel set of $\boldsymbol{v}_j$ represents a murmuration region, and its zero level set denotes its interface with other regions. In this specialized case, the $\boldsymbol{E}$ and $\boldsymbol{P}$'s speeds and maximum turn radii are equal: if both players start the game with the same initial velocity and orientation, the relative equations of motion show that $\boldsymbol{E}$ can mimic $\boldsymbol{P}$'s strategy by forever keeping the starting radial separation. As such, the *barrier* is closed and *the central theme in this game of kind is to determine the surface* [29]. Owing to the high-dimensionality of the state space, we cannot resolve this barrier analytically, hence we resort to our numerical approximation method.

Suppose that a murmuration's global heading is predetermined and that each bird i_j within each flock, F_j, $(j = \{1, \cdots, n_f\})$ has a constant linear velocity, v_j^i. Level-set methods accommodate connectivity changes without explicitly reconnecting tracked surfaces. Let $\cup$ describe the union operator — the aggregation symbol of individual bird payoffs in a flock — which may be implicitly constructed as the isocontour of some signed distance function. We can thus aggregate the payoff

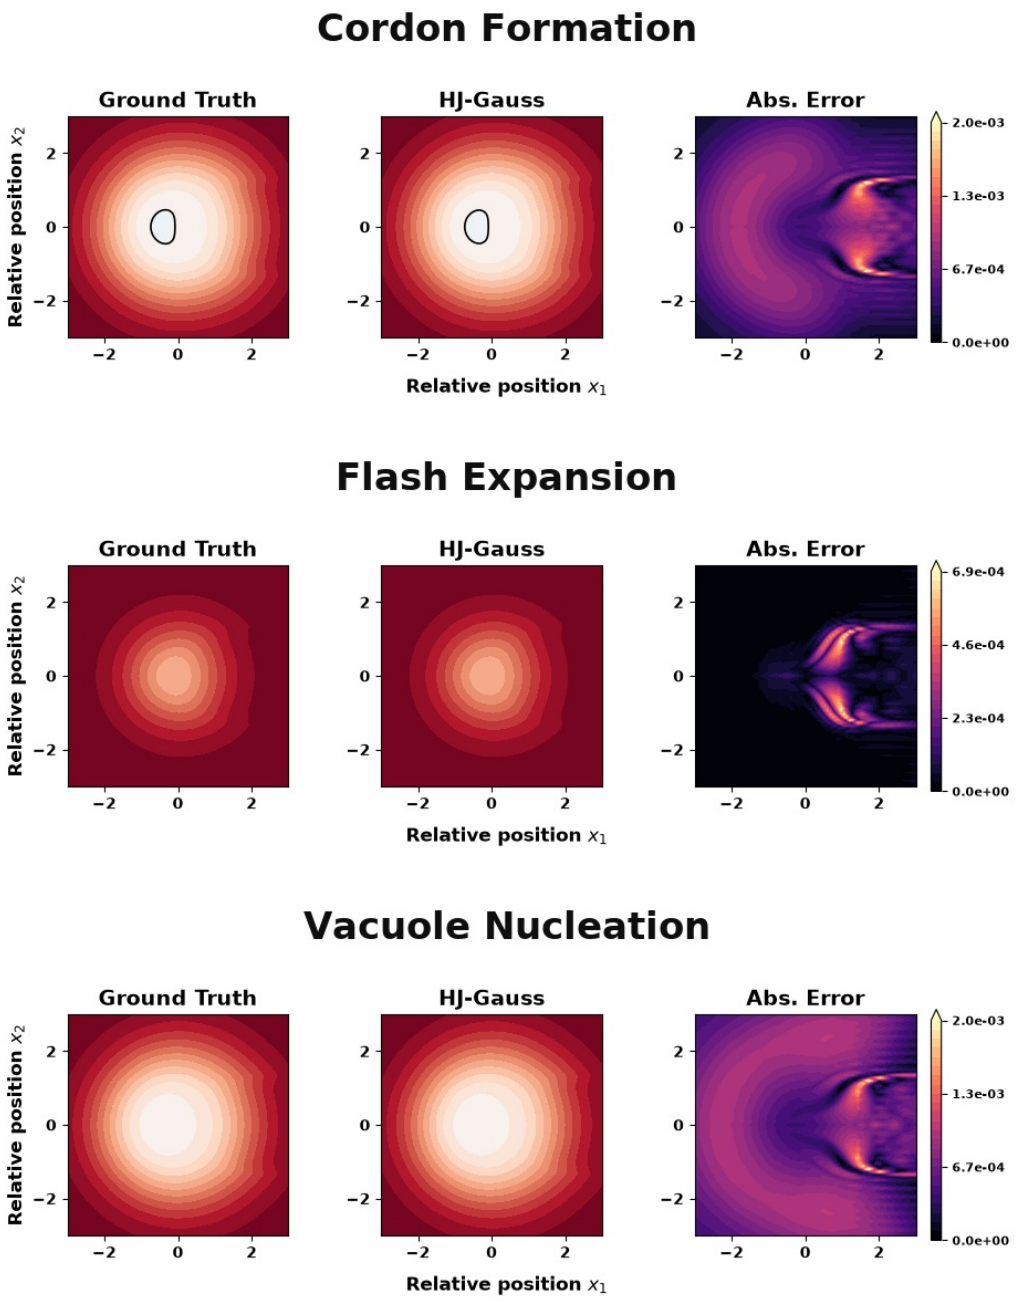


Figure 4: Value functions for three murmuration events, shown at the final horizon: cordon formation, flash expansion, and vacuole nucleation. Each panel shows the `Levelset` Reference (left), the `HJ-Gauss` solution (center), and their absolute difference (right).

of birds 1 and 2, for example, as

$$\cup\{\boldsymbol{v}_1(\boldsymbol{x},t),\boldsymbol{v}_2(\boldsymbol{x},t)\} \equiv \boldsymbol{v}_1(\boldsymbol{x},t) \cup \boldsymbol{v}_2(\boldsymbol{x},t) = \min(\boldsymbol{v}_1(\boldsymbol{x},t),\boldsymbol{v}_2(\boldsymbol{x},t)). \tag{3.2}$$

The terminal motion of a flock coincide with the extremal of the payoff functional, given as

$$\boldsymbol{v}(\boldsymbol{x},t) = \inf_{\beta^{(1)}\in\mathcal{B}^{(1)}} \sup_{\boldsymbol{u^{(1)}}\in\mathcal{U}^{(1)}} g^{(1)}(\boldsymbol{x}(T)) \cup \cdots \inf_{\beta^{(n_f)}\in\mathcal{B}^{(n_f)}} \sup_{\boldsymbol{u}^{(n_f)}\in\mathcal{U}^{(n_f)}} .g^{(n_f)}(\boldsymbol{x}(T)) \tag{3.3}$$

where n_f is the total number of distinct flocks in a murmuration. The resolution of this equation admits a viscosity solution to the following variational terminal HJI PDE [32]

$$\begin{aligned} &\cup_{j=1}^{n_f}\left[\cup_{i=1}^{n_a}\left(\frac{\partial \boldsymbol{v}_i}{\partial t}(\boldsymbol{x},t) + \min\left[0, \boldsymbol{H}^{(i)}(\boldsymbol{x}^{(i)}, \boldsymbol{v}_x(\boldsymbol{x},t))\right]\right)\right] \\ &\qquad = 0. \end{aligned} \tag{3.4}$$

with Hamiltonian,

$$\boldsymbol{H}^{(i)}(t;\boldsymbol{x}^{(i)},\boldsymbol{u}^{(i)},\boldsymbol{v}^{(i)},p^{(i)}) = \max_{u^{(i)}\in\mathcal{U}^{(i)}} \min_{v^{(i)}\in\mathcal{V}^{(i)}} \langle f^{(i)}(t;\boldsymbol{x},\boldsymbol{u}^{(i)},\boldsymbol{v}^{(i)}), p^{(i)}\rangle. \tag{3.5}$$

For a flock, F_j, the Hamiltonian is the total energy given by a summation of the exerted energy by each agent i so that we can verify a realization of the *main equation* (3.5) as

$$\boldsymbol{H}(\boldsymbol{x},p) = \max_{w_e^{(k)_j}\in[\underline{w}_e^j,\bar{w}_e^j]} \min_{w_p^{(k)_j}\in[\underline{w}_p^j,\bar{w}_p^j]} \cup_{j=1}^{n_f}\left[H_a^{(k)_j}(\boldsymbol{x},p) \cup \left(\cup_{i=1}^{n_a-1} H_f^{(i)_j}(\boldsymbol{x},p)\right)\right] \tag{3.6}$$

$$\begin{aligned} &= \cup_{j=1}^{n_f}\left(\cup_{i=1}^{n_a-1}\left[p_1^{(i)_j} v^{(i)_j}\cos\boldsymbol{x}_4 + p_2^{(i)_j} v^{(i)_j}\sin\boldsymbol{x}_4 + p_3^{(i)_j} u_z^{(i)_j} + p_4^{(i)_j}\langle w_e^{(i)_j}\rangle_r\right]\right. \\ &\quad \cup\left[p_1^{(k)_j}\left(v^{(k)_j} - v^{(k)_j}\cos\boldsymbol{x}_4^{(k)_j}\right) - p_2^{(k)_j} v^{(k)_j}\sin\boldsymbol{x}_4^{(k)_j} - \underline{w}_p^j|p_4^{(k)_j}| \right. \\ &\quad \left.\left. + \bar{w}_e^j\left|p_2^{(k)_j}\boldsymbol{x}_1^{(k)_j} - p_1^{(k)_j}\boldsymbol{x}_2^{(k)_j} + p_3^{(k)_j}\right|\right]\right). \end{aligned} \tag{3.7}$$

where $\boldsymbol{H}_a^{(k)_j}(\boldsymbol{x},p)$ is the Hamiltonian of the individual under attack by a pursuing agent, $\boldsymbol{P}$, and $H_f^{(i)_j}(\boldsymbol{x},p)$ are the respective Hamiltonians of the free agents, $i=\{1,\cdots,n_f\}$, within an evading flock in a murmuration, and not under the direct influence of capture or attack by $\boldsymbol{P}$; we denote by $w_e^{(i)_j}$ the heading of an evader i within a flock j and $w_p^{(j)}$ the heading of a pursuer aimed at flock j; $\underline{w}_e^{(k)_j}$ is the orientation that corresponds to the orientation of the agent with minimum turn radius among all the neighbors of agent k, inclusive of agent k at time t; similarly, $\bar{w}_e^{(k)_j}$ is the maximum orientation among all of the orientation of agent k's neighbors.

Equation (3.7) can be verified by writing the Hamiltonian of the free agents in absolute coordinates and those of the agent under attack in relative coordinates with respect to the pursuer. Fixing the pursuer's minimum and maximum turn and and climb rates yields the equation. Note that when the linear speeds of the evading agents and pursuer are equal i.e., $v_e^{(i)_j}(t) = v_p(t) = +1m/s$, a murmuration's Hamiltonian reduces to

$$\begin{aligned} \boldsymbol{H}(\boldsymbol{x},p) = \cup_{j=1}^{n_f}\left(\cup_{i=1}^{n_a-1}\left[p_1^{(i)_j}\cos\boldsymbol{x}_4 + p_2^{(i)_j}\sin\boldsymbol{x}_4 + p_3^{(i)_j}u_{z_e}^{(i)_j} + p_4^{(i)_j}\langle w_e^{(i)_j}\rangle_r\right]\cup_{k_j}\right. \\ \left[p_1^{(k)_j}\left(1-\cos\boldsymbol{x}_4^{(k)_j}\right) - p_2^{(k)_j}\sin\boldsymbol{x}_4^{(k)_j} - \underline{w}_p^j|p_4^{(k)_j}| + \bar{w}_e^j\left|p_2^{(k)_j}\boldsymbol{x}_1^{(k)_j} - p_1^{(k)_j}\boldsymbol{x}_2^{(k)_j}\right.\right. \\ \left.\left.\left. + p_4^{(k)_j}\right| - \gamma|p_3^{(k)_j}| - \gamma|p_3^{(k)_j}|\right]\right). \end{aligned} \tag{3.8}$$

We now describe four characterizing features of a murmuration which we use to validate the scalability of Algorithm 1. Owing to paper length reasons, all proofs in the events described below are in the online version of our paper [35, §D].

Vacuole Nucleation Topology: When an attacking predator penetrates the interior of a flock's BRT, a topological hole or a *vacuole* [40] nucleates in the zero-sublevel set $\mathcal{L}(t) \triangleq \{x : v(t,x) \leq 0\}$. Define this *Euler characteristic* [16] event as $\chi(t) \triangleq V - E + F \triangleq \beta_0(t) - \beta_1(t)$ with vertex, edge, and face counts V, E, F of the discretization of $\mathcal{L}(t)$ on an $m \times m$ grid; here β_0 is the number of connected components and β_1 is the first Betti number (number of independent loops/holes).

Theorem 3.1 (Vacuole Nucleation Realization on the BRT)**.** *Suppose that the evader agent's Hamiltonian* $\boldsymbol{H}_f(\boldsymbol{x}, p)$ *satisfies* $\boldsymbol{H}_f(\boldsymbol{x}^*, p^*) = 0$ *for some interior state* $\boldsymbol{x}^* \in \Omega$ *with costate* $p^* = D\boldsymbol{v}(t^*, \boldsymbol{x}^*)$*. Then at time* t^* *when the zero-crossing occurs, if the predator penetrates the interior of the flock's reachable set, the topology of* $\mathcal{L}(t^\star)$ *undergoes a transition characterized by,* $\chi(t^{*+}) = \chi(t^{*-}) - 1$*, where* χ *is the Euler characteristic. This jump indicates the nucleation of a vacuole (topological hole) in the flock's safe set, marking the transition from simply connected (*$\beta_1 = 0$*) to multiply connected (*$\beta_1 \geq 1$*) topology.*

Proof. By Morse theory for sublevel sets of smooth functions, the topology of $\Omega^-(t)$ changes only at critical values of $\boldsymbol{v}(\cdot, t)$. Specifically, when the parameter t crosses a value t^* at which a saddle point of $\boldsymbol{v}(\cdot, t^*)$ lies on the boundary of $\Omega^-(t^*)$, the sublevel set attaches a 1-handle, which increases the first Betti number as $\beta_1(t^{*+}) = \beta_1(t^{*-}) + 1$. Hence,

$$\chi(t^{*+}) = \beta_0(t^{*+}) - \beta_1(t^{*+}) = \beta_0(t^{*-}) - (\beta_1(t^{*-}) + 1) \triangleq \chi(t^{*-}) - 1. \tag{3.9}$$

□

Flock Splitting and the BRT: When a predator applies sufficient control to reduce the evader's optimal escape velocity, the BRT splits into disjoint components.

Proposition 3.1 (Flock Splitting)**.** *Let* $\Omega^-(t)$ *be the zero-sublevel set of the BRT at time t. The reachable set undergoes an* $n_{\text{comp}}(t) > 1$ *splitting into connected components if and only if the evading bird's Hamiltonian satisfies the control-constrained minimax condition,*

$$H_{\text{att}}(\boldsymbol{x}^*, p^*) = \min_{u_p \in \mathcal{U}_p} \max_{u_e \in \mathcal{U}_e} \left[p_1(u_p - v_e) + p_2 u_e + p_3(u_p^z - u_e^z) + p_4(\omega_p - \omega_e) \right] < -\delta_c, \tag{3.10}$$

where $\delta_c > 0$ *is a separation margin,* $\mathcal{U}_p, \mathcal{U}_e$ *are admissible control sets, and* (u_p, ω_p)*,* (u_e, ω_e) *are the pursuer and evader acceleration and angular velocity controls respectively.*

Proof. We reduce the split to the passage of a saddle through the zero level, and then identify the saddle's rate with the minimax in (3.10).

Let $\boldsymbol{x}^\star$ be the lowest saddle separating two candidate basins. The components separate exactly when the saddle value is positive,

$$n_{\text{comp}}(t) > 1 \iff v(t, \boldsymbol{x}^\star) > 0. \tag{3.11}$$

Along the ridge the spatial gradient vanishes, so the HJI equation gives the rate at which the saddle value rises,

$$\frac{d}{dt} v(t, \boldsymbol{x}^\star) = -\min\big(0, H_{\text{att}}(\boldsymbol{x}^\star, p^\star)\big). \tag{3.12}$$

The saddle lifts above zero within one comparison step exactly when this rate exceeds the margin, that is, when $-H_{\text{att}}(\boldsymbol{x}^\star, p^\star) > \delta_c$. Combining this with (3.11) and (3.12) gives the equivalence. The minimax in (3.10) is attained because the integrand is affine in each control, so the extremizers are bang–bang. □

Remark 3.1 (Temporal Splitting Pattern)**.** The split occurs at a critical time t_{split} determined by when the predator's control law first permits the pursuer to violate the co-state constraint in the evader's optimal strategy. The sequence of splitting times is monotone increasing: $t_{\text{split}}^{(1)} < t_{\text{split}}^{(2)} < \cdots$, with the m-th bifurcation occurring when isolation of the m-th agent becomes achievable under the game-theoretic equilibrium. The flock's resilience to splitting depends on the coupling strength in the risk-reduction term (the p_4 coefficient in the Hamiltonian), which represents the predation risk shared across connected components.

Cordon geometry and the BRT: A flock executing a defensive *cordon formation* creates an annular perimeter around threatened agents, preventing isolated predator incursions into the interior. This configuration is geometrically characterized by a multiply connected reachable set with one independent loop ($\beta_1 = 1$).

Proposition 3.2 (Annular BRT and Barrier Certificate)**.** *Let $\mathcal{C}(r_{\text{in}}, r_{\text{out}}) \triangleq \{x \in \Omega : r_{\text{in}} \le \|x_{1:2}\| \le r_{\text{out}}\}$ denote the annular barrier region ("cordon") with inner radius r_{in} and outer radius r_{out}. If the BRT restricted to $\mathcal{C}$ satisfies,*

$$v(t,x) > 0 \quad \text{for all } x \in \{y \in \mathcal{C} : \|y_{1:2}\| < r_{\text{in}}\}, \tag{3.13}$$

then the inner region is a barrier-certificate-protected domain*, and any trajectory of the attacked agent passing through $\mathcal{C}(r_{\text{in}}, r_{\text{out}})$ remains captured (i.e., reaches the terminal set) within finite time. The topology of $\Omega^-(t) \cap \mathcal{C}$ is annular, with Betti number $\beta_1 = 1$ and Euler characteristic,*

$$\chi\big(\Omega^-(t) \cap \mathcal{C}\big) = 0. \tag{3.14}$$

Proof. We prove the barrier property and the topology separately.

Barrier. Along an optimal motion the value is nonincreasing, so v cannot increase along the trajectory. By (3.13), the inner region has $v > 0$, so a trajectory inside the core stays outside the capture tube and cannot reach the terminal set. Conversely, a trajectory that meets the annulus with $v \le 0$ stays in the sublevel set, which excludes the core. It is therefore driven to the terminal set in finite time.

Topology. Under (3.13), the sublevel set $\Omega^-(t) \cap \mathcal{C}$ excludes the open core and is bounded outside by r_{out}. Because $g(\boldsymbol{x}) = \|\boldsymbol{x}_{1:2}\| - r$ is radially monotone, this set deformation-retracts onto the closed annulus, which is homotopy equivalent to a circle. Hence $\beta_0 = 1$ and $\beta_1 = 1$, and

$$\chi\big(\Omega^-(t) \cap \mathcal{C}\big) = \beta_0 - \beta_1 = 0. \tag{3.15}$$

□

Remark 3.2 (Barrier Certificate Interpretation)**.** The annular topology (Eq. (3.14)) physically represents a mutual-defense configuration where the safe set (valued region, $\boldsymbol{v} \le 0$) wraps around the predator or protected interior, creating a topological obstruction to escape. From the reachability perspective, any agent trapped inside the cordon cannot escape; from outside, no agent can cross the barrier without incurring catastrophic cost. This geometric property is independent of the detailed dynamics and depends only on the level-set structure of v.

The numerical validation of the cordon configuration is visualised through the reachability level sets of our simulations, which show the barrier topology explicitly.

Flash expansion and the BRT: The *flash expansion* is characterized by a rapid, isotropic outward motion of the flock boundary in response to a predator approach. This corresponds to a linear growth of the zero-level-set radius.

Lemma 3.1 (Bounded Flash Expansion Rate)**.** *Under the 4D aerial dynamics with evader linear speed v_e, the maximum radial extent of the BRT zero-level-set in the horizontal plane grows at most linearly,*

$$\max_i \|\boldsymbol{x}_i(t+1)_{1:2}\| - \max_i \|\boldsymbol{x}_i(t)_{1:2}\| \le v_e \, \Delta t, \tag{3.16}$$

per backward time step of size Δt. The altitude component $\boldsymbol{x}_3$ does not contribute to the radial growth rate because the capture set is a cylinder — the terminal cost excludes the heading and altitude coordinates from the capture metric.

(iii) Phase Transition Markers in Numerical Simulations

We detect phase transitions computationally by monitoring the topological invariants $\chi(t)$, $\beta_1(t)$, and $n_{\text{comp}}(t)$ of the zero-sublevel set across backward time steps. The detection algorithm is: (i) At each step t, compute the 2D slice $v(t, x_1, x_2; \bar{x}_3, 0)$ on a grid of resolution $m \times m$. (ii) Label connected components of $\{v \le 0\}$. (iii) Estimate β_1 by counting encircled complementary components. (iv) Compute $\chi(t) = n_{\text{comp}}(t) - \beta_1(t)$. The threshold-crossing rules that trigger an

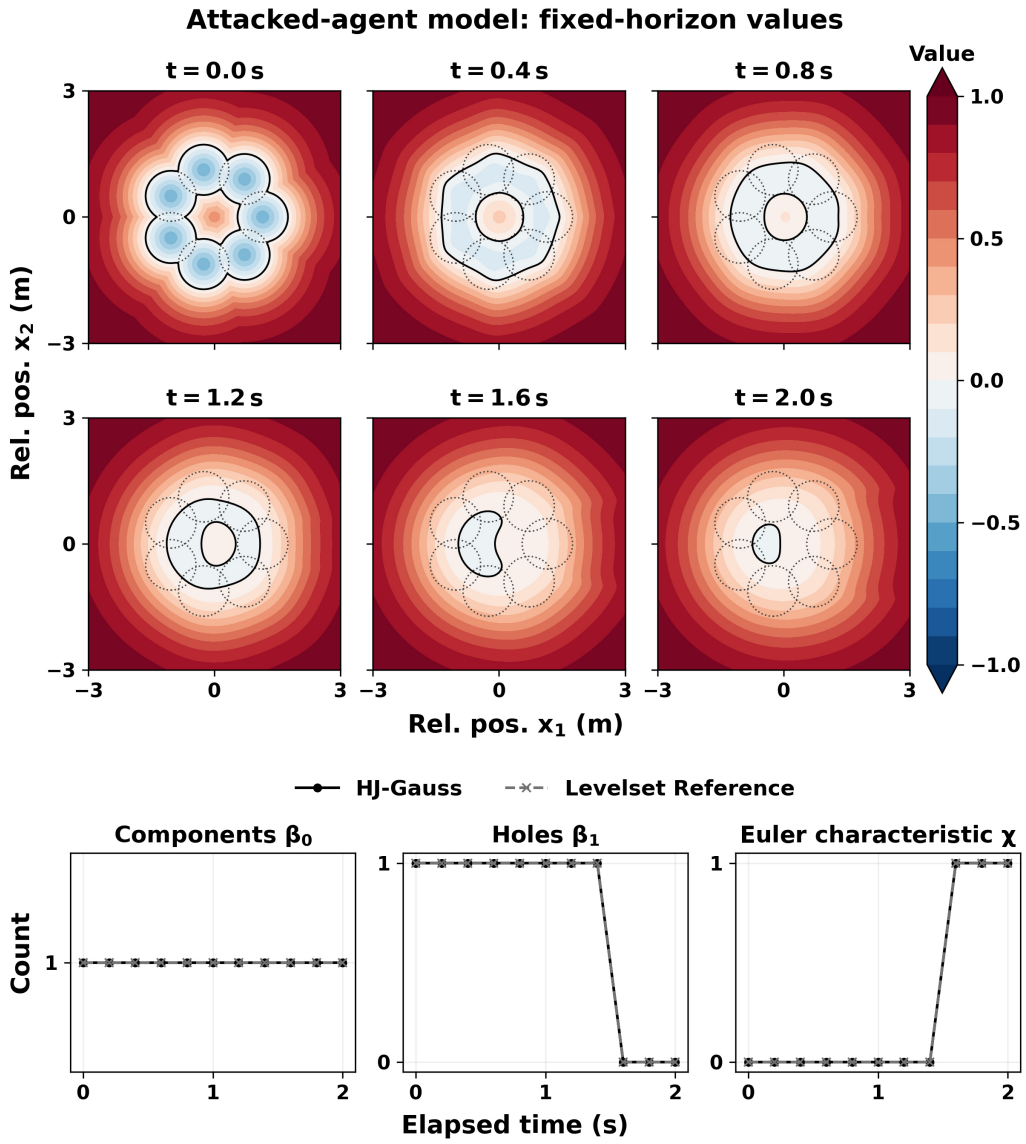


Figure 5: Numerical geometry of the flock value approximation. The upper panels show values at six elapsed horizons in seconds, with black zero contours and dotted initial target disks. The lower panel displays the component count, hole count, and Euler characteristic used to describe the slice topology.

vacuole nucleation, cordon formation, and flock fragmentation events are

$$\Delta\chi_t \triangleq \chi_{t+1} - \chi_t < -0.5, \beta_1(t) \geq 1 \text{ and } \beta_1(t-1) = 0, n_{\text{comp}}(t) > n_{\text{comp}}(t-1), \tag{3.17}$$

respectively. The phase-transition classification partitions the state space into five canonical murmuration phases: (i) cohesion, (ii) evasion, (iii) cordon, (iv) expansion, and (v) fragmentation. Each detected event is logged with its time index, event type, and χ, β_1, n_{comp} values, so that the topology metrics are traced across the full backward horizon.

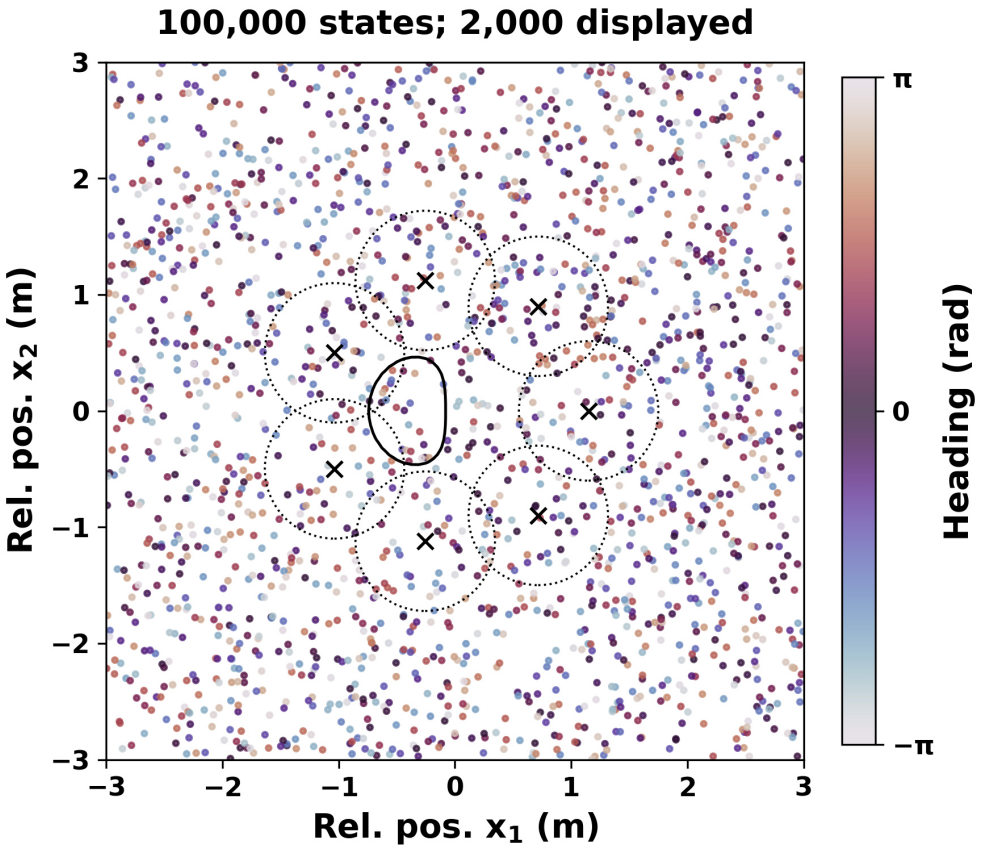


Figure 6: Planar projection of 2,000 displayed bird-model states, colored by heading. Crosses and dotted circles mark target centers and initial disks; positions are in meters.

(iv) Numerical geometry

Figure 5 compares the zero sublevel geometry with the levelset reference [34]. On a 193×193 planar grid at $\theta = 0$, both solvers give the same component, hole, and Euler counts. The displayed set has one connected component throughout. Figure 6 displays a reproducible subset of 2,000 states from the state population. Color encodes each state's heading. The overlaid zero contour is the $\theta = 0$ slice at $t = 2$; each population value instead uses that state's own heading.

(c) The Double Integrator Time-optimal Control Problem

Here, we analyze a time-optimal control problem on the double integral plant [5, 42] to determine what admissible control can "transport" the system under consideration to a desired "origin" in the shortest possible time. The goal is to compute points in the state space that can reach the origin in *finite-time* under the influence of a time-optimal controller. We formulate its value function as the *time-to-reach-the-origin* minimization problem. With second-order dynamics $\ddot{\boldsymbol{x}} = \boldsymbol{u}$, change of variables yields

$$\dot{\boldsymbol{x}}_1(t) = \boldsymbol{x}_2(t), \qquad \dot{\boldsymbol{x}}_2(t) = \boldsymbol{u}(t), \qquad |\boldsymbol{u}(t)| \leq 1. \tag{3.18}$$

with Hamiltonian $\boldsymbol{H}_{\text{phys}}(\boldsymbol{x}, p, \boldsymbol{u}) = p_1 \boldsymbol{x}_2 + p_2 \boldsymbol{u}$. The *reachability problem addresses states that can reach the origin in a transient manner*, such that we can bring the system to the equilibrium state, $(0, 0)$. This is an $\boldsymbol{H}$-minimal control problem on the Hamiltonian

$$\boldsymbol{H}(\boldsymbol{x}, p) = p_1 \dot{\boldsymbol{x}}_1 + p_2 \dot{\boldsymbol{x}}_2. \tag{3.19}$$

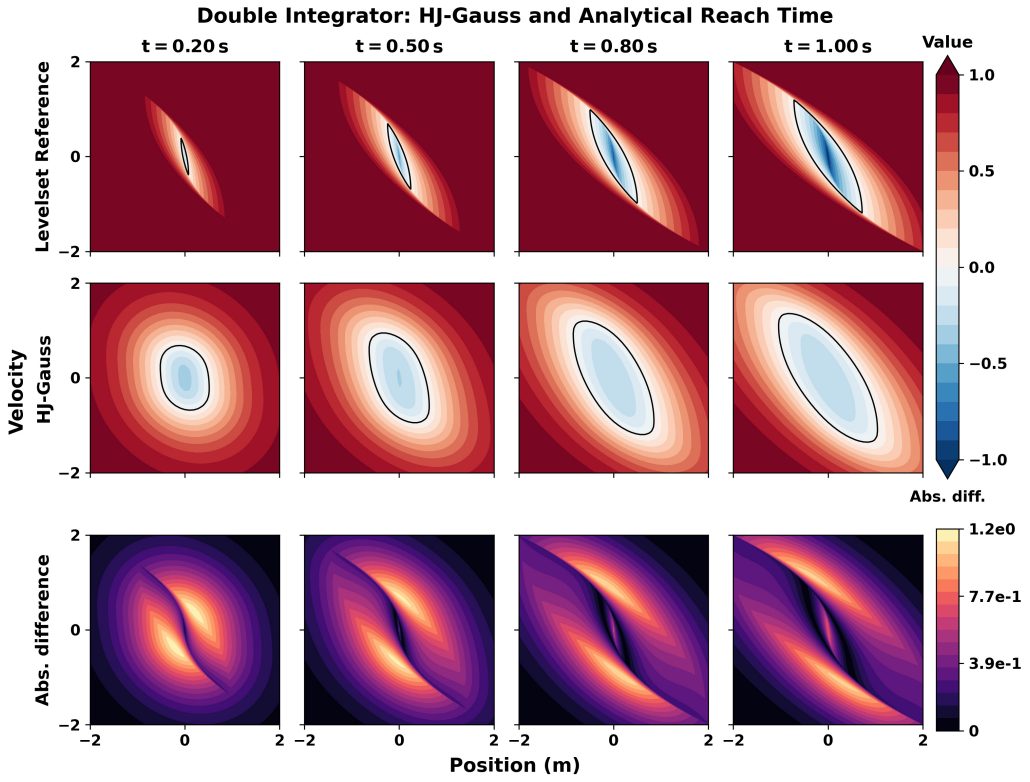


Figure 7: Double-integrator values at four elapsed times; `HJ-Gauss` uses $\delta = 0.08$. Rows show the analytical `Levelset` Reference, `HJ-Gauss`, and their absolute difference. Black contours mark the zero level sets in the first two rows.

The necessary optimality condition stipulates that the minimizing control law be

$$\boldsymbol{u}(t) = -\operatorname{sign}\left(p_2(t)\right) \triangleq \pm 1. \tag{3.20}$$

With co-state equations $\dot{p}_1 = 0$ and $\dot{p}_2 = -p_1$, it can be verified that the state trajectories $(\boldsymbol{\xi}_1, \boldsymbol{\xi}_2)$ obey the relation

$$\boldsymbol{x}_1(t) = \boldsymbol{\xi}_1 + \frac{1}{2}k\left(\boldsymbol{x}_2^2 - \boldsymbol{\xi}_2^2\right), \text{ for } t = k\left(\boldsymbol{x}_2(t) - \boldsymbol{\xi}_2\right). \tag{3.21}$$

The *time to reach the origin* $(0,0)$ from pair (x_1, x_2) on the state trajectories plane *in the shortest possible time* (see [34]) is our approximation problem. The geometry of the analytical time to reach the origin problem (see trajectories plane *in the shortest possible time* (see [34]) is validated against our `HJ-Gauss` method in Figure 7. The bottom row shows the pointwise absolute discrepancy at each integration time. We see that `HJ-Gauss` overapproximates the zero levelset and the pointwise errors between both methods is negligibly zero in the zero levelset.

(d) The Rockets Launch Problem on a Plane Benchmark

The rocket-launch examples of Dreyfus [13] on a plane, translated to a differential game in [34] motivates the next benchmark. We state the important abstractions needed for our problem construction here. Readers interested in the details should consult our previous texts [33, 34]. Its dynamics in relative coordinates is

$$\dot{x}_1 = a\cos\theta + u_e x_1, \quad \dot{x}_2 = a\sin\theta + a + u_p x_1 - g, \quad \dot{\theta} = u_p - u_e. \tag{3.22}$$

where g is the gravitational acceleration, $\theta \in \mathbb{S}^1$ is a relative-heading state, and $u_e, u_p \in [-1, 1]$ are the attitude rate controls for evading and pursuing vehicles respectively, and a is the acceleration in ft/sec^2. For a two-player differential game, $\boldsymbol{P}$ and $\boldsymbol{E}$ share same dynamics so that we freely choose the coordinates of $\boldsymbol{P}$ while $\boldsymbol{E}$'s origin is a distance Φ away from (x, z) at the $x - z$ plane's origin so that the $\boldsymbol{PE}$ vector's inclination measured counterclockwise from the $\boldsymbol{x}$ axis is θ. The evader is maximizing while the pursuer is minimizing the payoff. Capture occurs when $\Phi(x_1, x_2) < \bar{r}$ for a user-specified value r. All capture points are specified by the variational HJ PDE [32] so that the game's Hamiltonian is (see derivation in [34])

$$\boldsymbol{H}(\boldsymbol{x}, p) = -ap_1 \cos\theta - p_2(g - a - a\sin\theta) - \bar{u}|p_1 x + p_3| + \underline{u}|p_2 x + p_3|, \tag{3.23}$$

where we set $a = 1$, $g = 2$, $\Phi = 1.5$, $\delta = 0.08$, and $T = 0.5$. Figure 1 compares `HJ-Gauss`with the `levelsetpy` reference [34].

The initial value function (leftmost inset of Figure 1) is represented as a (closed) dynamic implicit surface over all point sets in the state space (using a signed distance function) for a coordinate-aligned cylinder whose vertical axes runs parallel to the rockets' orientation. Notice the near zero difference in absolute values of the computation using the `levelset`and `HJ-Gauss`schemes across three different representative slices of the value function Reachability [28, 30] enables us to numerically reason about the behavior of these two rockets before motion execution without closed-form geometrical analysis.

(e) Dubins Vehicles Homicidal Chauffeur Problem

Merz's pursuit–evasion game extends the homicidal chauffeur problem to two vehicles with comparable speeds and turn radii [22, 29]. Its value's *boundary of the usable part* [22] constitutes the zero level-set geometry of the BRT benchmark[31]. The experiment described below computes terminal-horizon viscosity values. Let $\boldsymbol{x} = (x_1, x_2, x_3)$ be the pursuer's position and heading relative to an evader centered at the origin. With constant linear speeds $v_e, v_p > 0$ and steering rates $w_e, w_p \in [-w, w]$, in *relative coordinates* the dynamics is (see [32])

$$\dot{x}_1 = -v_e + v_p \cos x_3 + w_e x_2, \quad \dot{x}_2 = v_p \sin x_3 - w_e x_1, \quad \dot{x}_3 = w_p - w_e. \tag{3.24}$$

Set $v_e = v_p = w = 1$ so that for matched headings, identical steering signals preserve the initial separation. The capture target surface is the cylinder $\mathcal{T} = \{\boldsymbol{x} : x_1^2 + x_2^2 \le r^2\}$, with capture radius $r > 0$. In absolute coordinates, the dynamics is $\dot{x}_1 = \cos\theta$, $\dot{x}_2 = \sin\theta$, and $\dot{\theta} = u$, with $|u| \le 1$ and a pose target that includes heading. Write the Hamiltonian as $\boldsymbol{H} = -p_1\cos\theta - p_2\sin\theta + |p_3|$. Capture dynamics is based on the variational HJ equation

$$\ell_t + \max\{0, \boldsymbol{H}(x, D\ell)\} = \frac{\delta}{2}\Delta\ell, \tag{3.25}$$

whose Hamiltonian $\boldsymbol{H}_{\text{phys}} = \max_{w_e} \min_{w_p} p \cdot f$ is

$$\boldsymbol{H}_{\text{phys}}(\boldsymbol{x}, p) = p_1(v_p \cos x_3 - v_e) + p_2(v_p \sin x_3) + w|p_1 x_2 - p_2 x_1 - p_3| - w|p_3|. \tag{3.26}$$

For the normalized bounds $v_e = v_p = w = 1$, the Hamiltonian becomes

$$\boldsymbol{H}_{\text{phys}}(\boldsymbol{x}, p) = p_1(\cos x_3 - 1) + p_2(\sin x_3) + |p_1 x_2 - p_2 x_1 - p_3| - |p_3|. \tag{3.27}$$

For elapsed backward time, the numerical Cauchy problem uses $\boldsymbol{H} = -\boldsymbol{H}_{\text{phys}}$. For unit bounds this is

$$\boldsymbol{H}(\boldsymbol{x}, p) = p_1(1 - \cos x_3) - p_2 \sin x_3 - |p_1 x_2 - p_2 x_1 - p_3| + |p_3|. \tag{3.28}$$

The recomputed terminal-horizon fields use this Hamiltonian without a stopping operator.

The levelset comparison [34] with `HJ-Gauss` for the Dubins vehicles in relative coordinates is shown in Figure 8; Table 2 reports its refinement checks.

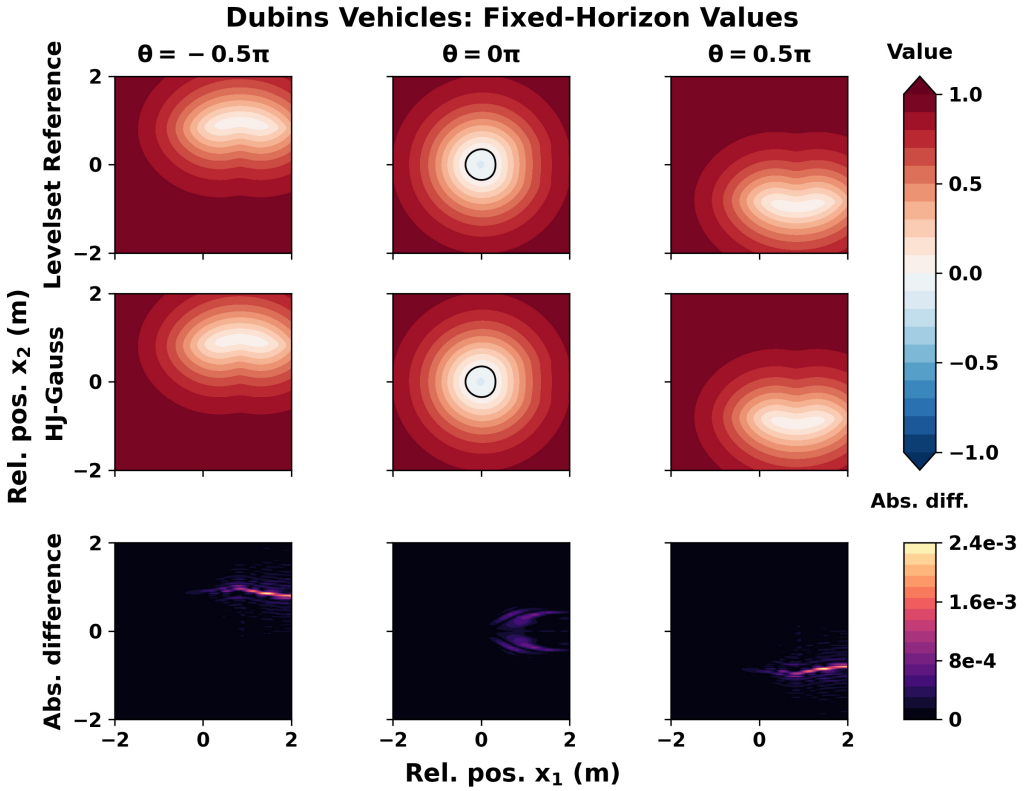


Figure 8: Relative-Dubins terminal-horizon values at $T = 1$ and $\delta = 0.08$. Rows show the `Levelset` Reference [34], `HJ-Gauss`, and their absolute difference. Columns select three headings. Black contours mark the zero level set. The comparison uses the same Hamiltonian, initial condition, and viscosity in both solvers.

(f) Numerical setup notes and hyperparameters

We describe the numerical setup in the codes that implements Algorithm 1. For the rockets and relative-Dubins vehicle cases, `HJ-Gauss` integrates the (2.9) with a second-order exponential Runge–Kutta time stepping on a periodic Fourier grid — a collocation approximation to the transformed equation. A separate fourth-order finite-difference discretization with classical 4th-order Runge–Kutta time stepping evolves the original viscous equation. The bird model uses a 4th-order centered finite differences and classical 4th-order Runge–Kutta integration for `HJ-Gauss`. Its `levelset` reference uses 6th-order centered differences and 4th-order Runge–Kutta integration. The deterministic comparisons test these discretizations.

All figures assign the `levelset`scales $L_* = 1\,\mathrm{m}$ and $T_* = 1\,\mathrm{s}$, so that positions are displayed in meters and elapsed time in seconds. The Starlings example payoff takes the minimum of seven translated radial payoffs, with target centers on a radius-1.15 ring and disk radius 0.6. The baseline spatial box is $[-6, 6]^2$, with periodic heading $[-\pi, \pi)$ for three-dimensional cases. Completion makes the initial condition equal to 1 near each spatial seam while leaving it unchanged on $[-4.5, 4.5]^2$. The observation box is $[-2, 2]^2$ for the single-Dubins, relative-Dubins, and double-integrator cases and $[-3, 3]^2$ for rockets and birds. The prescribed coefficient is

$$\boldsymbol{c}(t, \boldsymbol{x}) = b + 0.04 \cos t \prod_{j=1}^{d} \cos(2\pi \boldsymbol{x}_j / L_j), \qquad b = 0.2, \tag{3.29}$$

Table 1: Maximum sampled changes under time halving, coefficient variation, isolated boundary extension, and combined box sensitivity. Boundary extension fixes the `Levelset` Reference time step, spacing, heading count, Hamiltonian, and initial condition on the common box.

System	Time	Coefficient	Combined box
Rockets	7.19×10^{-6}	1.84×10^{-4}	0.00136
Dubins vehicles	3.97×10^{-5}	2.79×10^{-5}	5.26×10^{-5}
Single Dubins	1.47×10^{-6}	1.02×10^{-5}	8.83×10^{-6}
Double integrator	6.85×10^{-7}	1.73×10^{-6}	4.40×10^{-5}
Bird model	6.23×10^{-7}	7.46×10^{-6}	5.46×10^{-4}

where L_j is the corresponding periodic length. The computation includes its time derivative, gradient, and Laplacian. The sensitivity calculation changes b to 0.35 while keeping the Hamiltonian and initial condition fixed.

Successive refinements can change both spatial resolution and time step; fixed-grid time halving is checked separately. Refinement and solver discrepancies are measured on 65×65 observation points and, in three dimensions, 32 equally spaced headings. These are maximum sampled discrepancies of interpolated fields. Equation residuals use native grid states and spectral spatial derivatives and centered differences at one quarter, one half, and three quarters of the elapsed horizon.

The acceptance thresholds were fixed before the refinement cycle: 0.01 for successive field refinements, 0.02 for the levelset reference [34], 0.002 for time halving and coefficient changes, and 0.005 for each domain-extension test. Sampled original-equation residuals must have maximum at most 0.15 and root-mean-square at most 0.02. The box experiment extends the half-width to 7.5 at fixed spatial spacing and heading resolution; it also changes the time step and coefficient wavelength, so it measures their combined sensitivity. A separate boundary-sensitivity test evolves the original equation on the enlarged box while fixing the time step, spatial spacing, heading count, Hamiltonian, and initial condition on the common domain. This separates boundary sensitivity from changes in the transformation coefficient and time step. Diffusion has infinite propagation speed, so neither finite extension supplies a whole-space boundary-error certificate.

Table 1 separates time-step, coefficient, and domain sensitivity. All comparisons meet their thresholds; these sampled checks do not establish the source iteration's invariant-set hypotheses.

(g) Population-scale evaluation of starling-inspired dynamics

The attacked-agent model has two relative positions, altitude, and heading (Appendix 3(g)). Equal vertical control bounds cancel the vertical Hamiltonian terms; an altitude-independent payoff then reduces the solution to the three-dimensional relative-Dubins equation. Independent enumeration of the four-dimensional controls checks this reduction.

Positions are independently uniform on $[-3, 3]^2$, altitude on $[-1, 1]$, and heading on $[-\pi, \pi)$. Every queried value is retained, together with its levelset reference [34] value.

4. Conclusion

The variable-coefficient Cole–Hopf transformation maps the viscous Hamilton–Jacobi equation to a semilinear heat equation with a complete nonlinear source. Algorithm 1 evaluates its coupled mild value-and-gradient updates by depth-first antithetic Gaussian sampling. It retains the nonlinear source at recursively sampled intermediate states and rejects nonpositive transformed estimates before logarithmic recovery.

The dimension sweep demonstrates the narrow storage advantage of point evaluation: dimensions 8–2048 require no tensor grid, and the counted coordinate payload is linear in

Table 2: Maximum sampled discrepancies for the `HJ-Gauss` computations. Refinement and `Levelset` Reference columns use the common observation queries; the residual column gives the largest sampled residual over native states. The grid column lists the fine spatial and heading counts.

System	Grid	Refinement	Levelset Reference	Residual
Rockets	$160^2 \times 80$	0.007481	0.003223	0.1433
Dubins vehicles	$224^2 \times 112$	0.005717	0.002511	0.1379
Single Dubins	$128^2 \times 64$	0.009297	0.001415	0.02992
Double integrator	128^2	0.006741	5.60×10^{-4}	0.01187
Bird model	$192^2 \times 96$	0.008817	6.44×10^{-4}	0.1295

dimension at fixed sampling and depth. The sample sweep gives a different runtime result. Increasing samples per update from 4 to 48 raises median runtime from 0.0052 to 0.9355 seconds. These measurements do not establish $O(Nn)$ end-to-end runtime or favorable scaling with Picard depth.

Separate deterministic transformed-PDE computations compare rockets, Dubins vehicles, a double integrator, and a starling-inspired relative-state field with independent same-viscosity discretizations. The 100,000-state experiment evaluates one precomputed three-dimensional field at independently sampled four-component queries; it does not solve a coupled flock game. Component, hole, and Euler counts describe thresholded numerical slices rather than continuum topology transitions or population-wide safety certificates.

The released code records finite Picard levels, random seeds, operation counts, fresh-process timing replicates, and memory measurements. Extending a finite collection of point estimates to a continuous reachable set still requires controlled field reconstruction, accumulated sampling-error bounds, verified admissibility, and a stopping rule tied to the contraction hypotheses.

References

1 M. Althoff, G. Frehse, and A. Girard. Set Propagation Techniques for Reachability Analysis. *Annual Review of Control, Robotics, and Autonomous Systems*, 4:369–395, 2021. ISSN 25735144.

2 J.-P. Aubin, A. M. Bayen, and P. Saint-Pierre. *Viability Theory*. Number August 2015. 2011. ISBN 9783642166846.

3 M. Ballerini, N. Cabibbo, R. Candelier, A. Cavagna, E. Cisbani, I. Giardina, V. Lecomte, A. Orlandi, G. Parisi, A. Procaccini, M. Viale, and V. Zdravkovic. interaction Ruling Animal Collective Behavior Depends On Topological Rather Than Metric Distance: Evidence From A Field Study. *Proceedings of the National Academy of Sciences*, 105(4):1232–1237, 2008. ISSN 0027-8424. doi: 10.1073/pnas.0711437105. URL https://www.pnas.org/content/105/4/1232.

4 R. Bellman. *Dynamic programming*. Princeton University Press, 1957. ISBN 0-486-42809-5.

5 S. P. Bhat and D. S. Bernstein. Continuous finite-time stabilization of the translational and rotational double integrators. *IEEE Transactions on automatic control*, 43(5):678–682, 1998.

6 W. Bialek, A. Cavagna, I. Giardina, T. Mora, E. Silvestri, M. Viale, and A. M. Walczak. Statistical mechanics for natural flocks of birds. *Proceedings of the National Academy of Sciences*, 109(13): 4786–4791, 2012.

7 A. Cavagna, A. Cimarelli, I. Giardina, G. Parisi, R. Santagati, F. Stefanini, and M. Viale. Scale-ffee correlations in starling flocks. *Proceedings of the National Academy of Sciences*, 107(26):11865–11870, 2010.

8 M. G. Crandall and P.-L. Lions. Viscosity solutions of hamilton-jacobi equations. *Transactions of the American mathematical society*, 277(1):1–42, 1983.

9 M. G. Crandall and P.-L. Lions. Two Approximations of Solutions of Hamilton-Jacobi Equations. *Mathematics of Computation*, 43(167):1–19, 1984.

10 M. G. Crandall and A. Majda. Monotone Difference Approximations For Scalar Conservation Laws. *Mathematics of Computation*, 34(149):1–21, 1980.

11 M. G. Crandall, L. C. Evans, and P. L. Lions. Some Properties of Viscosity Solutions of Hamilton-Jacobi Equations. *Transactions of the American Mathematical Society*, 282(2):487, 1984. ISSN 00029947.

12 Defense Acquisition University. Verification, 2023. URL https://www.dau.edu/tools/se-brainbook/Pages/Technical%20Processes/verification.aspx. Accessed April 5, 2023.

13 S. E. Dreyfus. Control Problems With Linear Dynamics, Quadratic Criterion, and Linear Terminal Constraints. Technical report, Rand Corp, Santa Monica Calif, 1966.

14 L. Evans and P. E. Souganidis. Differential Games And Representation Formulas For Solutions Of Hamilton-Jacobi-Isaacs Equations. *Indiana Univ. Math. J*, 33(5):773–797, 1984. ISSN 0022-2518.

15 M. Haiken. These birds flock in mesmerizing swarms of thousands—but why is still a mystery., 2021. URL https://tinyurl.com/4973byey. Accessed April 5, 2023.

16 A. Hatcher. *Algebraic Topology*, chapter 2. Cambridge University Press, 2002. Euler characteristic: $\chi(X) = \sum_{n=0}^{\infty}(-1)^n \text{rank}(H_n(X))$.

17 D. Helbing, I. Farkas, and T. Vicsek. Simulating dynamical features of escape panic. *Nature*, 407 (6803):487–490, 2000.

18 D. Henry. *Geometric Theory of Semilinear Parabolic Equations*, volume 840 of *Lecture Notes in Mathematics*. Springer, 1981. doi: 10.1007/BFb0089647.

19 K. L. Hobbs, M. L. Mote, M. C. Abate, S. D. Coogan, and E. M. Feron. Runtime assurance for safety-critical systems. *IEEE Control Systems Magazine*, 43:28–65, 2023.

20 M. Hutzenthaler and T. Kruse. Multilevel Picard Approximations of High-Dimensional Semilinear Parabolic Differential Equations with Gradient-Dependent Nonlinearities. *SIAM Journal on Numerical Analysis*, 58(2):929–961, 2020. doi: 10.1137/17M1157015.

21 R. Isaacs. Differential Games: A Mathematical Theory with Applications to Warfare and Pursuit, Control and Optimization., 1966.

22 R. Isaacs. *Differential Games: A Mathematical Theory with Applications to Warfare and Pursuit, Control and Optimization.* Kreiger, Huntigton, NY, 1999.

23 A. Jadbabaie, J. Lin, and A. S. Morse. Coordination of groups of mobile autonomous agents using nearest neighbor rules. *IEEE Transactions on automatic control*, 48(6):988–1001, 2003.

24 P. Karnakov, S. Litvinov, and P. Koumoutsakos. Computing Foaming Flows Across Scales: From Breaking Waves to Microfluidics. *arXiv preprint arXiv:2103.01513*, 2021.

25 Y. Kim, M.-C. Lai, and C. S. Peskin. Numerical Simulations of Two-Dimensional Foam by the Immersed Boundary Method. *Journal of Computational Physics*, 229(13):5194–5207, 2010.

26 P.-L. Lions. *Generalized solutions of Hamilton-Jacobi equations*, volume 69. London Pitman, 1982.

27 C. Livadas and N. A. Lynch. Formal verification of safety-critical hybrid systems. In *International Workshop on Hybrid Systems: Computation and Control*, pages 253–272. Springer, 1998.

28 J. Lygeros. On reachability and minimum cost optimal control. *Automatica*, 40(6):917–927, 2004.

29 A. Merz. The game of two identical cars. *Journal of Optimization Theory and Applications*, 9(5): 324–343, 1972.

30 I. Mitchell. Games of two identical vehicles. *Dept. Aeronautics and Astronautics, Stanford Univ.*, (July):1–29, 2001.

31 I. Mitchell. A Robust Controlled Backward Reach Tube with (Almost) Analytic Solution for Two Dubins Cars. *EPiC Series in Computing*, 74:242–258, 2020.

32 I. M. Mitchell, A. M. Bayen, and C. J. Tomlin. A Time-Dependent Hamilton-Jacobi Formulation of Reachable Sets for Continuous Dynamic Games. *IEEE Transactions on Automatic Control*, 50 (7):947–957, 2005. ISSN 00189286.

33 L. Molu. The Python LevelSet Toolbox (LevelSetPy). In *2024 IEEE 63rd Conference on Decision and Control (CDC)*, pages 8938–8945, 2024. doi: 10.1109/CDC56724.2024.10886640.

34 L. Molu. LevelSetPy: A GPU-Accelerated Package for Hyperbolic Hamilton-Jacobi Partial Differential Equations' Solubility. *The ACM Transactions on Mathematical Software*, 2025.

35 L. Molu, V. Renganathan, and N. Cho. HJ-Gauss: Monte-Carlo HJ Reachability, 2026. URL https://scriptedonachip.com/downloads/Papers/hjgauss_full.pdf.

36 S. Osher and R. Fedkiw. Level Set Methods and Dynamic Implicit Surfaces. *Applied Mechanics Reviews*, 57(3):B15–B15, 2004. ISSN 0003-6900.
37 S. Osher and J. A. Sethian. Fronts Propagating with Curvature-Dependent Speed: Algorithms based on Hamilton-Jacobi Formulations. *Journal of Computational Physics*, 79(1):12–49, 1988. ISSN 10902716.
38 S. Osher and C.-W. Shu. High-Order Essentially Nonoscillatory Schemes for Hamilton-Jacobi Equations. *SIAM Journal of Numerical Analysis*, 28(4):907–922, 1991. ISSN 15333868.
39 C. J. Tomlin, J. Lygeros, and S. S. Sastry. A game theoretic approach to controller design for hybrid systems. *Proceedings of the IEEE*, 88(7):949–970, 2000.
40 A. van Leeuwenhoek. *The Select Works of Antony Van Leeuwenhoek, Containing His Microscopical Discoveries in Many of the Works of Nature*. 1800. Translated edition.
41 T. Vicsek, A. Czirók, E. Ben-Jacob, I. Cohen, and O. Shochet. Novel type of phase transition in a system of self-driven particles. *Physical review letters*, 75(6):1226, 1995.
42 W. M. Wonham. Linear multivariable control: A geometric approach. *Applications of Mathematics*, 10, 1985.
43 A. Zaitzeff, S. Esedoglu, and K. Garikipati. On the Voronoi Implicit Interface Method. *SIAM Journal on Scientific Computing*, 41(4):A2407–A2429, 2019.

Appendices

A. Transformation and sampling formulas

Retaining both residual components changes the representation formula. Throughout this appendix, the domain is $\mathbb{R}^n$, the diffusion coefficient is $\delta/2$, and $\boldsymbol{c}$ satisfies the bounds in §2(a).

(a) Sampling a fixed source

Monte Carlo sampling estimates a specified source integral. It does not evaluate the unknown source without an approximation of the space–time solution. Conditional on a fixed preceding iterate, write f for its deterministic source. For independent $U_i \sim \mathcal{U}(0,1)$ and $Z_i \sim \mathcal{N}(0, I_n)$, define

$$A_i = \boldsymbol{\omega}_0(\boldsymbol{x} + \sqrt{\delta t} Z_i) + t f(tU_i, \boldsymbol{x} + \sqrt{\delta t(1-U_i)} Z_i), \qquad \widehat{\boldsymbol{\omega}}_N = \frac{1}{N}\sum_{i=1}^{N} A_i. \tag{A.1}$$

The mean is unbiased for one mild update. The two terms may share the Gaussian draw, since independence is needed between the sample pairs, rather than within each pair. Neither $-\boldsymbol{c}^{-1}\log\widehat{\boldsymbol{\omega}}_N$ nor a ratio gradient estimator is generally unbiased. The logarithm requires a positive estimate. Uniform source-time sampling gives a score gradient whose second moment need not be finite near $s = t$. Instead, draw the time from

$$\pi_t(s) = \frac{1}{2\sqrt{t}\sqrt{t-s}}, \quad 0 < s < t, \qquad S = t(1-U^2). \tag{A.2}$$

Using independent copies of S, Z, a gradient sample is

$$B = \frac{Z_0 \boldsymbol{\omega}_0(\boldsymbol{x} + \sqrt{\delta t} Z_0)}{\sqrt{\delta t}} + 2\sqrt{\frac{t}{\delta}}\, Z_1 f(S, \boldsymbol{x} + \sqrt{\delta(t-S)} Z_1). \tag{A.3}$$

Its mean equals (2.21*a*). If $|\boldsymbol{\omega}_0| \le W$ and $|f| \le B_f$, then its second moment obeys

$$\mathbb{E}|B|^2 \le \frac{2nW^2}{\delta t} + \frac{8ntB_f^2}{\delta}. \tag{A.4}$$

The bound follows from $|a+b|^2 \le 2|a|^2 + 2|b|^2$ and $\mathbb{E}|Z|^2 = n$. No bounded-gradient concentration claim follows from a bounded value sample alone.

Endpoint importance sampling remains valid only after including the density ratio for each sampled integral. For example, for a Gaussian displacement h and variance σ^2, the original-to-shifted density ratio is

$$\frac{\varphi_{\boldsymbol{x}}(y)}{\varphi_{\boldsymbol{x}+h}(y)} = \exp\left(\frac{|h|^2 - 2h\cdot(y-\boldsymbol{x})}{2\sigma^2}\right). \tag{A.5}$$

Apply this ratio at each source time with $\sigma^2 = \delta(t-s)$. A tilt of terminal samples alone does not correct the source integral. Importance sampling may change the variance and requires its own moment bounds.

(b) Linear potential solves

A prescribed bounded continuous potential $\boldsymbol{R}$ and bounded initial condition give another representation. For the bounded mild solution, with $X_r = \boldsymbol{x} + \sqrt{\delta}W_r$ for a standard n-dimensional Brownian motion W, the solution of $\boldsymbol{\omega}_t = \delta\Delta\boldsymbol{\omega}/2 + \boldsymbol{R}\boldsymbol{\omega}$ satisfies

$$\boldsymbol{\omega}(t,\boldsymbol{x}) = \mathbb{E}_{\boldsymbol{x}}\left[\boldsymbol{\omega}_0(X_t)\exp\left\{\int_0^t \boldsymbol{R}(t-r,X_r)\,dr\right\}\right]. \tag{A.6}$$

Freezing $\boldsymbol{R}[\boldsymbol{c},\boldsymbol{v}^{(k)},D\boldsymbol{v}^{(k)}]$ produces a linear potential equation, whose positive solution defines a possible iteration. For the full equation, however, $\boldsymbol{R}$ depends on the unknown value and gradient. The formula then remains implicit. Its potential iteration is distinct from the additive source iteration analyzed below, and requires a separate convergence proof. A stable log-sum-exp evaluation applies to positive path weights after time discretization; it does not remove the path quadrature error.

B. Convergence, residual certification, and robustness

This appendix gives conditional convergence and a residual certificate for the full transformed equation.

(a) Convergence of source iteration

Gradient dependence requires a norm that controls both values and first derivatives. Let $\mathrm{BUC}^1(\mathbb{R}^n)$ contain continuously differentiable functions with bounded, uniformly continuous values and gradients. On $X_\tau = C([0,\tau];\mathrm{BUC}^1)$, use

$$\|z\|_{X_\tau} = \sup_{s\le\tau}\|z(s)\|_\infty + \sup_{s\le\tau}\|Dz(s)\|_\infty. \tag{B.1}$$

For $\boldsymbol{g}\in\mathrm{BUC}^1$ and bounded coefficient derivatives, the transformed initial condition belongs to the same space. Define the mild update by

$$(\mathcal{M}z)(t) = S_t\boldsymbol{\omega}_0 + \int_0^t S_{t-s}F(s,\cdot,z(s),Dz(s))\,ds. \tag{B.2}$$

Theorem B.1 (Conditional finite-time convergence)**.** *Let $\mathcal{A}\subset X_\tau$ be nonempty, closed, and invariant under $\mathcal{M}$. Suppose all its members satisfy $a\le z\le b$ and $|Dz|\le Q$, with $a>0$. Assume that the source is bounded and uniformly continuous on the relevant coefficient and state ranges, and satisfies*

$$|F(t,\boldsymbol{x},z,q) - F(t,\boldsymbol{x},\widetilde{z},\widetilde{q})| \le L\big(|z-\widetilde{z}| + |q-\widetilde{q}|\big). \tag{B.3}$$

If $\theta = L(\tau + 2\sqrt{n\tau/\delta}) < 1$, then $\mathcal{M}$ has a unique fixed point in $\mathcal{A}$. Its iterates satisfy

$$\|z^{(K)} - \boldsymbol{\omega}\|_{X_\tau} \le \theta^K\|z^{(0)} - \boldsymbol{\omega}\|_{X_\tau}, \qquad \|z^{(K)} - \boldsymbol{\omega}\|_{X_\tau} \le \frac{\theta}{1-\theta}\|z^{(K)} - z^{(K-1)}\|_{X_\tau} \tag{B.4}$$

for $K \geq 1$. Here finite-time convergence means convergence as $K \to \infty$ on the finite elapsed-time interval $[0, \tau]$; it does not assert termination after finitely many updates. If the fixed point is classical, its inverse solves the original viscous equation.

Proof of Theorems 2.2 and B.1. The space $\mathrm{BUC}^1(\mathbb{R}^n)$ is Banach under the sum of its uniform value and gradient norms. Hence X_τ is Banach, and its closed nonempty subset $\mathcal{A}$ is complete. Invariance makes $\mathcal{M}$ a self-map on this set.

For $z, \widetilde{z} \in \mathcal{A}$, write

$$d(s, \boldsymbol{x}) = F(s, \boldsymbol{x}, z(s, \boldsymbol{x}), Dz(s, \boldsymbol{x})) - F(s, \boldsymbol{x}, \widetilde{z}(s, \boldsymbol{x}), D\widetilde{z}(s, \boldsymbol{x})).$$

The source Lipschitz bound gives $\|d(s)\|_\infty \leq L\|z - \widetilde{z}\|_{X_\tau}$ for every $s \in [0, \tau]$. The common initial-condition terms cancel in $\mathcal{M}z - \mathcal{M}\widetilde{z}$. Since the heat semigroup contracts the uniform norm,

$$\|\mathcal{M}z(t) - \mathcal{M}\widetilde{z}(t)\|_\infty \leq \int_0^t \|d(s)\|_\infty \, ds \leq Lt\|z - \widetilde{z}\|_{X_\tau}. \tag{B.5}$$

Gaussian differentiation, followed by $\mathbb{E}|Z| \leq \sqrt{n}$ for $Z \sim \mathcal{N}(0, I_n)$, yields

$$\|DS_r h\|_\infty \leq \frac{\mathbb{E}|Z|}{\sqrt{\delta r}}\|h\|_\infty \leq \sqrt{\frac{n}{\delta r}}\|h\|_\infty, \qquad r > 0. \tag{B.6}$$

The singularity at $r = 0$ is integrable. Differentiating the source integral therefore gives

$$\|D\mathcal{M}z(t) - D\mathcal{M}\widetilde{z}(t)\|_\infty \leq L\int_0^t \sqrt{\frac{n}{\delta(t-s)}}\, ds \, \|z - \widetilde{z}\|_{X_\tau} = 2L\sqrt{\frac{nt}{\delta}}\|z - \widetilde{z}\|_{X_\tau}. \tag{B.7}$$

Taking the two time suprema and adding proves

$$\|\mathcal{M}z - \mathcal{M}\widetilde{z}\|_{X_\tau} \leq L\left(\tau + 2\sqrt{n\tau/\delta}\right)\|z - \widetilde{z}\|_{X_\tau} = \theta\|z - \widetilde{z}\|_{X_\tau}.$$

The assumption $\theta < 1$ and Banach's fixed-point theorem give a unique fixed point $\boldsymbol{\omega} \in \mathcal{A}$. For each starting iterate $z^{(0)} \in \mathcal{A}$, invariance keeps every subsequent iterate in $\mathcal{A}$. Repeated contraction against the fixed point gives the first estimate in (B.4).

For the posterior estimate, contraction between successive iterates gives

$$\|z^{(K+j+1)} - z^{(K+j)}\|_{X_\tau} \leq \theta^{j+1}\|z^{(K)} - z^{(K-1)}\|_{X_\tau}, \qquad j \geq 0.$$

Sum these differences from $j = 0$ to infinity and use convergence to the fixed point. The geometric series has sum $\theta/(1-\theta)$, proving the second estimate. When $\theta = 0$, the map is constant on $\mathcal{A}$ and one update reaches its fixed point.

Finally, $\boldsymbol{\omega} \geq a > 0$ permits logarithmic inversion. If the fixed point is classical, its mild equation implies the full transformed differential equation. Lemma 2.1 then shows that its inverse solves the original viscous equation. □

For bounded coefficient derivatives, $\boldsymbol{c} \geq \boldsymbol{c}_- > 0$, and $z \geq a > 0$, the inverse maps have bounded derivatives on bounded (z, q) ranges. A Hamiltonian uniformly bounded and locally Lipschitz in p on those ranges, uniformly over $(t, \boldsymbol{x})$, therefore supplies a finite source constant L. No convexity assumption is needed for this contraction argument.

The global complexity theorem in [20, Corollary 4.8] does not follow from this local argument. Their setting imposes a globally Lipschitz source in the value and gradient variables, together with strong smoothness and derivative-growth bounds for the complexity result. The transformed source need not meet those assumptions. Even with constant $\boldsymbol{c}$ and the linear Hamiltonian $\boldsymbol{H}(p) = ap$ in one dimension, it becomes

$$F(z, q) = -aq - \frac{\delta}{2z}q^2. \tag{B.8}$$

This source is singular as z approaches zero and is quadratic in q. Using a globally Lipschitz extension to analyze the original equation would require a proof that the solution remains inside the bounded positive region on which the sources agree. Transferring their complexity result

would also require verifying their dimension-dependent constants and quadrature regularity. The present contraction theorem establishes neither requirement.

Invariance is a substantive hypothesis. One sufficient construction starts with $b_0(t) = S_t\boldsymbol{\omega}_0$ and an X_τ ball of radius r about b_0. Choose $r < \inf \boldsymbol{\omega}_0$ to preserve positivity. If $|F| \le B_F$ throughout that ball and

$$B_F\left(\tau + 2\sqrt{n\tau/\delta}\right) \le r, \tag{B.9}$$

then the mild map preserves the ball. Under the preceding uniform-in-state Hamiltonian and coefficient assumptions, bounded initial values and gradients give finite source ranges for sufficiently small τ. The case of a constant $\boldsymbol{c} = 1/\delta$ and $\boldsymbol{H} = |p|^2/2$ has $F = 0$ and $\theta = 0$, so the hypotheses include a nonempty solvable case. Extension to a longer horizon requires renewed bounds and invariance on each interval; a global contraction is not asserted.

If $0 < \theta < 1$ and the initial error is bounded by $E_0 > 0$, a sufficient iteration count is

$$K \ge \max\left\{0, \left\lceil \frac{\log(E_0/\varepsilon)}{\log(1/\theta)} \right\rceil \right\}. \tag{B.10}$$

When $\theta = 0$, one update reaches the fixed point. A small change between noisy iterates alone does not certify convergence.

(b) Sampling and accumulated implementation error

For the fixed-source value samples in (A.1), suppose $a_0 \le \boldsymbol{\omega}_0 \le b_0$ and $|f| \le B_f$. Each sample lies in an interval of width $D = b_0 - a_0 + 2tB_f$. At M fixed queries, Hoeffding's inequality and a union bound give

$$\max_{1\le m\le M} |\widehat{\boldsymbol{\omega}}_N(t_m, \boldsymbol{x}_m) - (\mathcal{M}z)(t_m, \boldsymbol{x}_m)| \le D_* \sqrt{\frac{\log(2M/\eta)}{2N}} \tag{B.11}$$

with probability at least $1 - \eta$, using a common width bound D_*. Independence between different queries is unnecessary. The preceding iterate must be fixed conditional on the past, and the new samples must have the stated conditional law. A union bound over K adaptive iterations replaces M by MK.

(i) Proof of finite-query concentration

Proof of Theorem 2.1. Let $\mathcal{F}_{k-1}$ contain the sampling history before update k, for $k = 1, \dots, K$. Conditional on this history, the preceding iterate z_{k-1} and its source f_{k-1} are fixed. Write $(t_m, \boldsymbol{x}_m)$ for query m, where $m = 1, \dots, M$, and let $A_{k,m,i}$ be its ith sample, for $i = 1, \dots, N$.

Independence of the uniform time and Gaussian variables in the sampling law, together with boundedness of the source, permits integration in either order. The conditional sample mean is therefore

$$\mu_{k,m} := \mathbb{E}[A_{k,m,i} \mid \mathcal{F}_{k-1}] = S_{t_m}\boldsymbol{\omega}_0(\boldsymbol{x}_m) + \int_0^{t_m} S_{t_m - s} f_{k-1}(s)(\boldsymbol{x}_m)\, ds = (\mathcal{M}z_{k-1})(t_m, \boldsymbol{x}_m). \tag{B.12}$$

This is the prescribed mild update, rather than the unknown converged solution. Define its Monte Carlo estimate by $\widehat{\mu}_{k,m} = N^{-1}\sum_{i=1}^{N} A_{k,m,i}$.

Each sample lies between $a_0 - t_m B_F$ and $b_0 + t_m B_F$. Its interval width is at most the common deterministic bound D_*. If $D_* = 0$, all samples equal their conditional means almost surely, and the theorem follows with $h_N = 0$. Suppose henceforth that $D_* > 0$. Conditional independence within the batch and Hoeffding's inequality give, for every $h > 0$,

$$\mathbb{P}\left(|\widehat{\mu}_{k,m} - \mu_{k,m}| > h \mid \mathcal{F}_{k-1}\right) \le 2\exp\left(-\frac{2Nh^2}{D_*^2}\right). \tag{B.13}$$

Taking expectations removes the conditioning and preserves this upper bound. A union bound over the at most MK batches now yields

$$\mathbb{P}\left(\max_{k,m} |\widehat{\mu}_{k,m} - \mu_{k,m}| > h\right) \leq 2MK \exp\left(-\frac{2Nh^2}{D_*^2}\right). \tag{B.14}$$

Choosing $h = h_N$ from (2.24) makes the right-hand side equal to η. If the algorithm stops before K updates, its executed batches form a subset of those covered by the same bound. No independence between queries or updates is used; only the stated conditional independence within each batch is required. □

(ii) Proof of sample count and logarithmic recovery

Proof of Corollary 2.3. The sample-count condition (2.25) implies $h_N \leq \varepsilon_\omega$ by direct substitution into (2.24). Theorem 2.1 therefore bounds all covered transformed-value sampling errors by ε_ω on an event of probability at least $1 - \eta$.

On that same event, consider a target update $\mu \geq a > 0$ and its estimate $\widehat{\mu}$, with $|\widehat{\mu} - \mu| \leq h_N < a$. Then $\widehat{\mu} \geq a - h_N > 0$, so both logarithms are defined. Every point between μ and $\widehat{\mu}$ is at least $a - h_N$. Integrating the derivative of the logarithm along this interval gives

$$|\log \widehat{\mu} - \log \mu| \leq \frac{|\widehat{\mu} - \mu|}{a - h_N} \leq \frac{h_N}{a - h_N}.$$

At the query point, logarithmic inversion multiplies this difference by $1/\boldsymbol{c}(t, \boldsymbol{x}) \leq 1/\boldsymbol{c}_-$. Hence

$$\left| -\frac{\log \widehat{\mu}}{\boldsymbol{c}(t, \boldsymbol{x})} + \frac{\log \mu}{\boldsymbol{c}(t, \boldsymbol{x})} \right| \leq \frac{h_N}{\boldsymbol{c}_-(a - h_N)}. \tag{B.15}$$

The inequality holds simultaneously at every covered query satisfying the lower-bound hypothesis, without another failure-probability allocation. It bounds recovery of the prescribed update and makes no assertion of unbiased logarithmic recovery or error relative to the converged solution. □

(iii) Accumulated implementation error

The prototype instead uses the source-time density π_t for both value and gradient samples. Its value-source weight is $1/\pi_t(S) \leq 2t$, so the corresponding interval width is $b_0 - a_0 + 4tB_f$. The same concentration argument applies with that larger width.

This is a finite-query value bound. A continuum X_τ guarantee also requires gradient concentration and control of reconstruction between queries. Gaussian score samples are unbounded, so the same Hoeffding argument does not establish that guarantee. The second-moment bound in (A.4) supports variance estimates, but does not by itself control a continuum supremum.

If a numerical implementation separately establishes X_τ errors ϵ_k in each mild update and keeps its iterates in the admissible set, then contraction gives

$$\|\widehat{z}^{(K)} - \boldsymbol{\omega}\|_{X_\tau} \leq \theta^K E_0 + \sum_{j=0}^{K-1} \theta^{K-1-j} \epsilon_j. \tag{B.16}$$

The proof iterates the inequality $e_{k+1} \leq \theta e_k + \epsilon_k$. Thus sampling, source quadrature, and reconstruction errors accumulate through the solve. They do not enter only as a final independent sample average. If both transformed values are at least a, logarithmic inversion yields

$$\|\widehat{\boldsymbol{v}}^{(K)} - \boldsymbol{v}^\delta\|_\infty \leq \frac{\theta^K E_0 + \sum_{j=0}^{K-1} \theta^{K-1-j} \epsilon_j}{\boldsymbol{c}_- a}. \tag{B.17}$$

This conditional bound does not certify the accompanying numerical prototype, which uses time stepping rather than whole-interval Picard iteration.

(c) A posteriori certificate for the original equation

A residual of the full transformed equation measures the residual of the original equation directly. Assume comparison for the original parabolic equation on the whole space. Let U be a bounded classical approximation with bounded residual, and define

$$w = e^{-\boldsymbol{c}U}, \qquad d = w_t - \frac{\delta}{2}\Delta w - F(t, \boldsymbol{x}, w, Dw). \tag{B.18}$$

Theorem B.2 (Residual certificate)**.** *Under these hypotheses, the viscous solution with initial condition* $\boldsymbol{g}$ *obeys*

$$\|U(t) - \boldsymbol{v}^{\delta}(t)\|_{\infty} \le \|U(0) - \boldsymbol{g}\|_{\infty} + \int_0^t \left\| \frac{d(s)}{\boldsymbol{c}(s)w(s)} \right\|_{\infty} ds. \tag{B.19}$$

Proof. Lemma 2.1 gives $d = -\boldsymbol{c}wr_U$, with $r_U = U_t + \boldsymbol{H}(t, \boldsymbol{x}, DU) - \delta\Delta U/2$. Let $e_0 = \|U(0) - \boldsymbol{g}\|_{\infty}$ and $E(t) = e_0 + \int_0^t \|r_U(s)\|_{\infty} ds$. Since the Hamiltonian does not depend on the value, $U - E$ is a subsolution and $U + E$ is a supersolution. Their initial conditions bracket $\boldsymbol{g}$. Comparison gives $U - E \le \boldsymbol{v}^{\delta} \le U + E$, proving the bound. □

The source must include both residual components when computing d. If a solver omits $\boldsymbol{R}_{\mathrm{der}}$, then its transformed residual contains the missing term $-w\boldsymbol{R}_{\mathrm{der}}$. The certificate charges that omission. Residuals evaluated only at sampled states are diagnostic; rigorous use requires a bound between those states and through time, including derivative approximation error.

(d) Viscosity, model perturbations, and level sets

Vanishing viscosity remains a separate approximation. Suppose the Hamiltonian and initial condition satisfy a comparison and approximation theorem giving $\|\boldsymbol{v}^{\delta} - \boldsymbol{v}^{0}\|_{\infty} \le C_{\mathrm{vis}}\sqrt{\delta}$ on the chosen horizon. This rate is an additional hypothesis, rather than a consequence of the transformation. Under the same uniform norm and domain, Theorem B.2 then gives

$$\|U(t) - \boldsymbol{v}^{0}(t)\|_{\infty} \le \|U(0) - \boldsymbol{g}\|_{\infty} + \int_0^t \left\| \frac{d(s)}{\boldsymbol{c}(s)w(s)} \right\|_{\infty} ds + C_{\mathrm{vis}}\sqrt{\delta}. \tag{B.20}$$

No universal choice $\delta \sim N^{-1/3}$ follows: source, stability, and sampling constants can depend on δ. Their dependence must be established before optimizing a combined rate.

Proposition B.1 (Hamiltonian and initial condition perturbations)**.** *Suppose two equations with the same diffusion admit comparison, their initial conditions differ by at most* ϵ_g*, and* $|\boldsymbol{H}_1(t, \boldsymbol{x}, p) - \boldsymbol{H}_2(t, \boldsymbol{x}, p)| \le \epsilon_H$ *throughout the required co-state range. For classical solutions, it suffices that this range contains* $D\boldsymbol{v}_2$*. Then*

$$\|\boldsymbol{v}_1(t) - \boldsymbol{v}_2(t)\|_{\infty} \le \epsilon_g + t\epsilon_H. \tag{B.21}$$

Proof. The functions $\boldsymbol{v}_2 - \epsilon_g - t\epsilon_H$ and $\boldsymbol{v}_2 + \epsilon_g + t\epsilon_H$ are a subsolution and a supersolution for the first equation. Comparison proves the claim. For viscosity solutions, impose the Hamiltonian bound on the test-gradient range required by comparison, or uniformly over all p. □

Changing a prescribed positive transformation coefficient does not perturb the Hamiltonian when both source components are retained. Thus a smooth positive regularization that preserves the stated differentiability changes coordinates; it creates no Hamiltonian modeling error by itself. Hard clipping can destroy that differentiability and falls outside these classical results. If the Hamiltonian is deliberately replaced, the perturbation bound applies to that separate change. The matching choice $\boldsymbol{c} = 2\boldsymbol{H}(t, \boldsymbol{x}, D\boldsymbol{v})/(\delta|D\boldsymbol{v}|^2)$ can vanish, change sign, or become singular. At $\boldsymbol{c} = 0$, the exponential transform loses the value information. Recomputing this ratio also requires regularity beyond the prescribed-coefficient hypotheses. `HJ-Gauss` permits a positive constant coefficient even in these cases.

If a uniform bound $|U - \boldsymbol{v}^0| \le E$ has been established on a specified domain, then

$$\{U < -E\} \subseteq \{\boldsymbol{v}^0 < 0\}, \qquad \{U > E\} \subseteq \{\boldsymbol{v}^0 > 0\}. \tag{B.22}$$

Proof of Corollary 2.4. Fix a point in the specified set and suppress its arguments. The assumed error bound is equivalent to

$$U - E \le \boldsymbol{v}^0 \le U + E. \tag{B.23}$$

If $U > E$, then the lower bound gives $\boldsymbol{v}^0 \ge U - E > 0$. Thus the point lies outside the target-reaching class $\{\boldsymbol{v}^0 \le 0\}$. If $U < -E$, then the upper bound gives $\boldsymbol{v}^0 \le U + E < 0$, placing the point in the strict negative sublevel set. Since the argument applies at every point covered by the error bound, it proves both inclusions in (B.22).

The states assigned neither strict sign by this rule satisfy $-E \le U \le E$, which is the abstention band $|U| \le E$. At its boundaries, the corresponding estimate in (B.23) is nonstrict, so the corollary makes no strict sign assertion there.

If the assumed error bound holds simultaneously on an event of probability at least $1 - \eta$, both sign implications hold simultaneously on that same event. No additional union bound is needed. The domain of the sign certificate is the domain on which that error bound has been established; a bound at finitely many queries supplies no assertion between those queries. □

Interpreting either sign as safety also requires the corresponding game's value convention and, for reach–avoid problems, a correctly solved obstacle problem.